

# Cloud Native, Capacity, Performance and Cost Optimization Tools and Techniques

CMG Workshop

November 2013

Adrian Cockcroft

@adrianco @NetflixOSS

<http://www.linkedin.com/in/adriancockcroft>

# Presentation vs. Workshop

- Presentation
  - Short duration, focused subject
  - One presenter to many anonymous audience
  - A few questions at the end
- Workshop
  - Time to explore in and around the subject
  - Tutor gets to know the audience
  - Discussion, rat-holes, “bring out your dead”

# Attendee Introductions

- Who are you, where do you work
- Why are you here today, what do you need
- “Bring out your dead”
  - Do you have a specific problem or question?
  - One sentence elevator pitch
- What instrument do you play?

# Content

---

**Cloud Native**

---

**Migration Path**

---

**Service and API Architectures**

---

**Storage Architecture**

---

**Operations and Tools**

---

**Cost Optimization**

---



More?

# Cloud Native

What is it?

Why?

# Strive for perfection

Perfect code

Perfect hardware

Perfectly operated



But perfection takes too long...

Compromises...

Time to market vs. Quality

Utopia remains out of reach

# Where time to market wins big

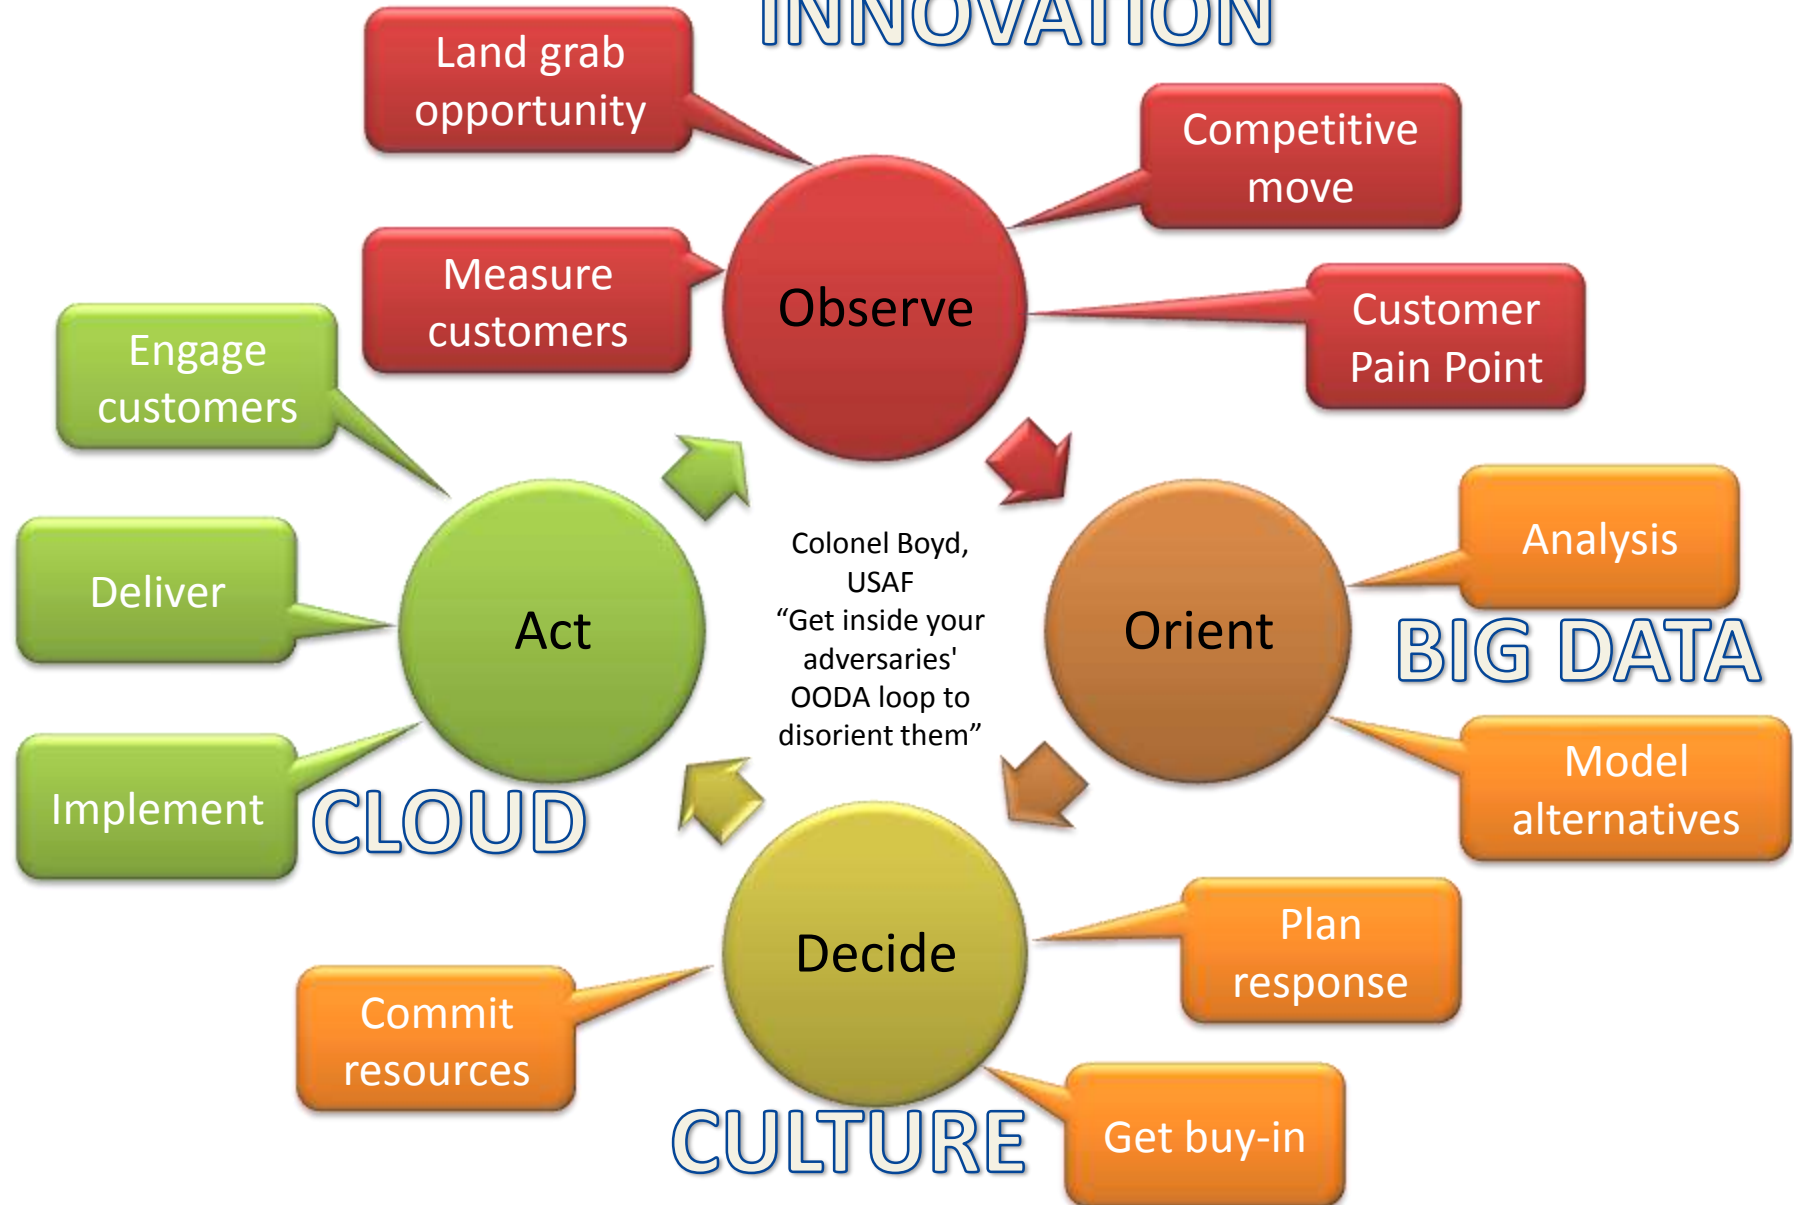
Making a land-grab

Disrupting competitors (OODA)

Anything delivered as web services



# INNOVATION



# How Soon?

Product features in days instead of months

Deployment in minutes instead of weeks

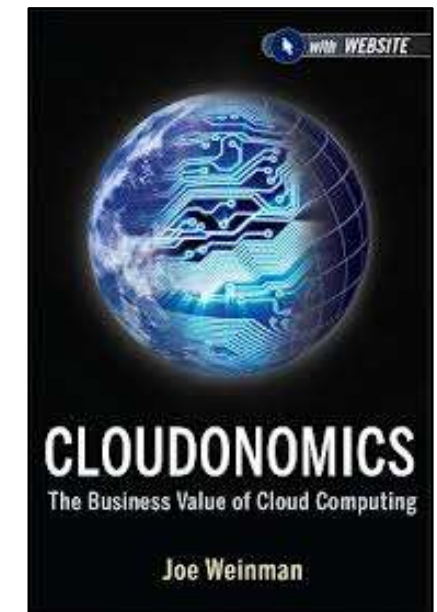
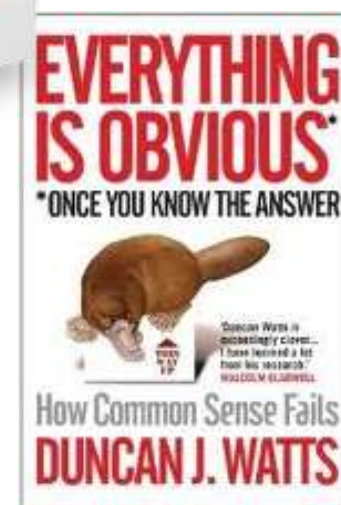
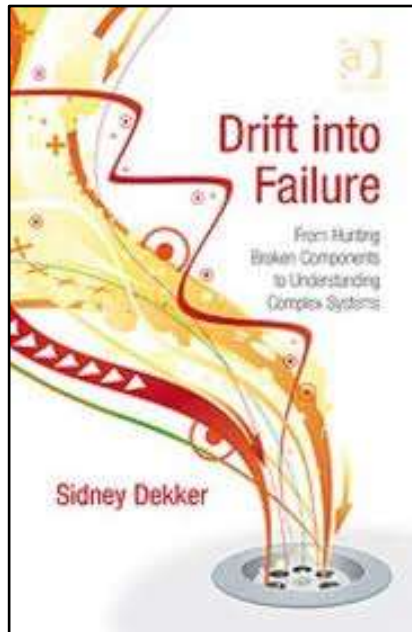
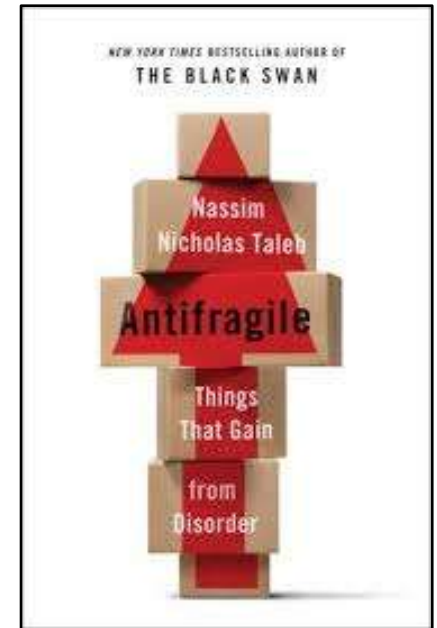
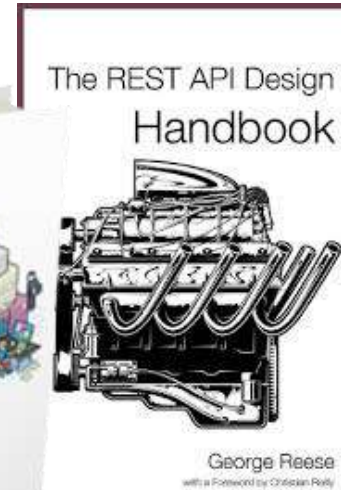
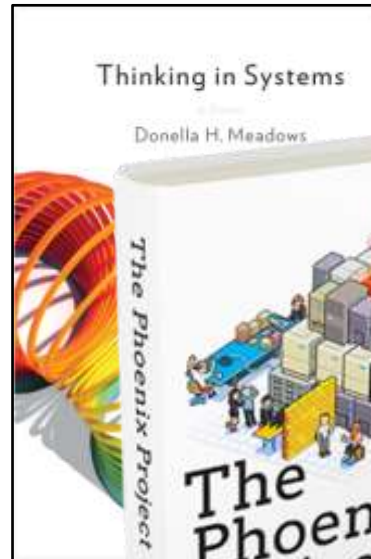
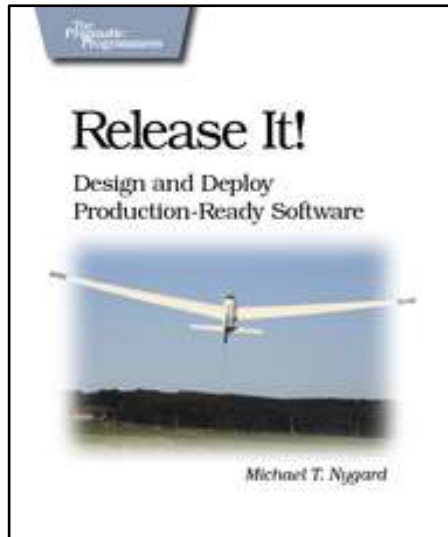
Incident response in seconds instead of hours

# Cloud Native

## A new engineering challenge

Construct a highly agile and highly available service from ephemeral and assumed broken components

# Inspiration



# How to get to Cloud Native

Freedom and Responsibility for Developers

Decentralize and Automate Ops Activities

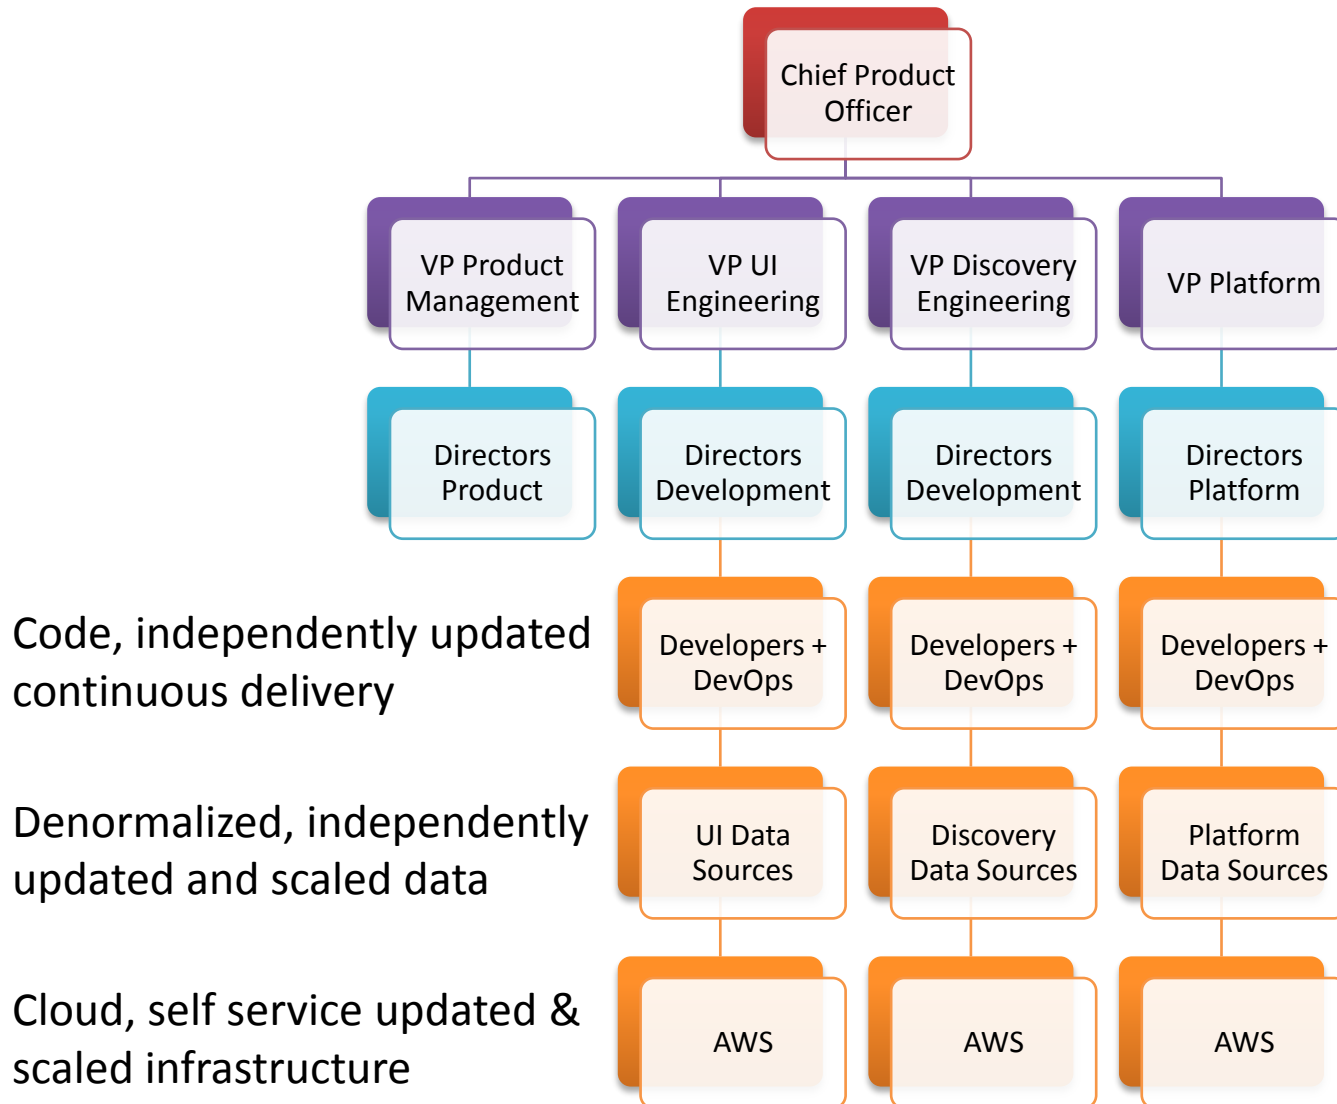
Integrate DevOps into the Business Organization

**Re-Org!**

# Four Transitions

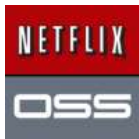
- Management: Integrated Roles in a Single Organization
  - Business, Development, Operations -> BusDevOps
- Developers: Denormalized Data – NoSQL
  - Decentralized, scalable, available, polyglot
- Responsibility from Ops to Dev: Continuous Delivery
  - Decentralized small daily production updates
- Responsibility from Ops to Dev: Agile Infrastructure - Cloud
  - Hardware in minutes, provisioned directly by developers

# Netflix BusDevOps Organization



# Decentralized Deployment





# Asgard Developer Portal

<http://techblog.netflix.com/2012/06/asgard-web-based-cloud-management-and.html>

**asgardprod/us-east-1/cluster/show/obiwan** **This cluster contains two ASGs**

**ASGARD prod** us-east-1 CMC

Home App AMI Cluster ELB EC2 SDB SNS SQS RDS Task

### Manage Cluster of Sequential Auto Scaling Groups

Recommended next step: Switch traffic to the preferred group, then delete legacy group

**obiwan-v063**

Launch and Terminate are disabled

Resize to 9 min / 12 max

Delete Disable Enable

9 instances grouped by state

Count	State	Build	ELB	Disc
+ 9	InService	580	OUT_OF_SERVICE	

**No traffic on old version**

**obiwan-v064**

Resize to 9 min / 12 max

Delete Disable Enable

9 instances grouped by state

Count	State	Build	ELB	Disc
+ 9	InService	583	UP	

**Live traffic on new version**

Create Next Group: **obiwan-v065** [+ Advanced Options](#)

AMI Image ID: 179123456789/obiwan-41.2-141730! [Filter](#) [+ Show more AMIs](#)

Instance Type: m1.large \$230.400/mo [Filter](#)

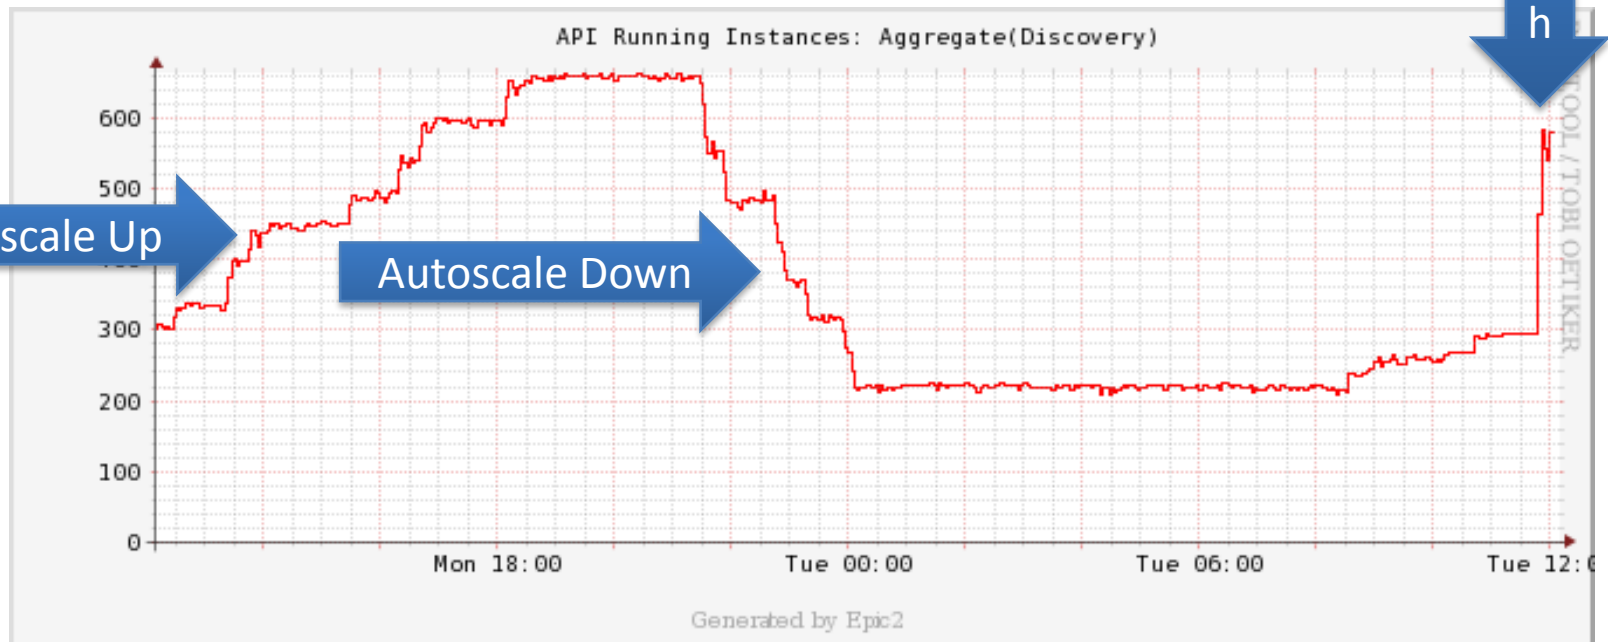
Instance Counts: Min: 9 Desired: 9 Max: 12

After launch: ☒ Wait for Discovery health check pass

[+ Create Next Group obiwan-v065](#)

# Ephemeral Instances

- Largest services are autoscaled
- Average lifetime of an instance is 36 hours



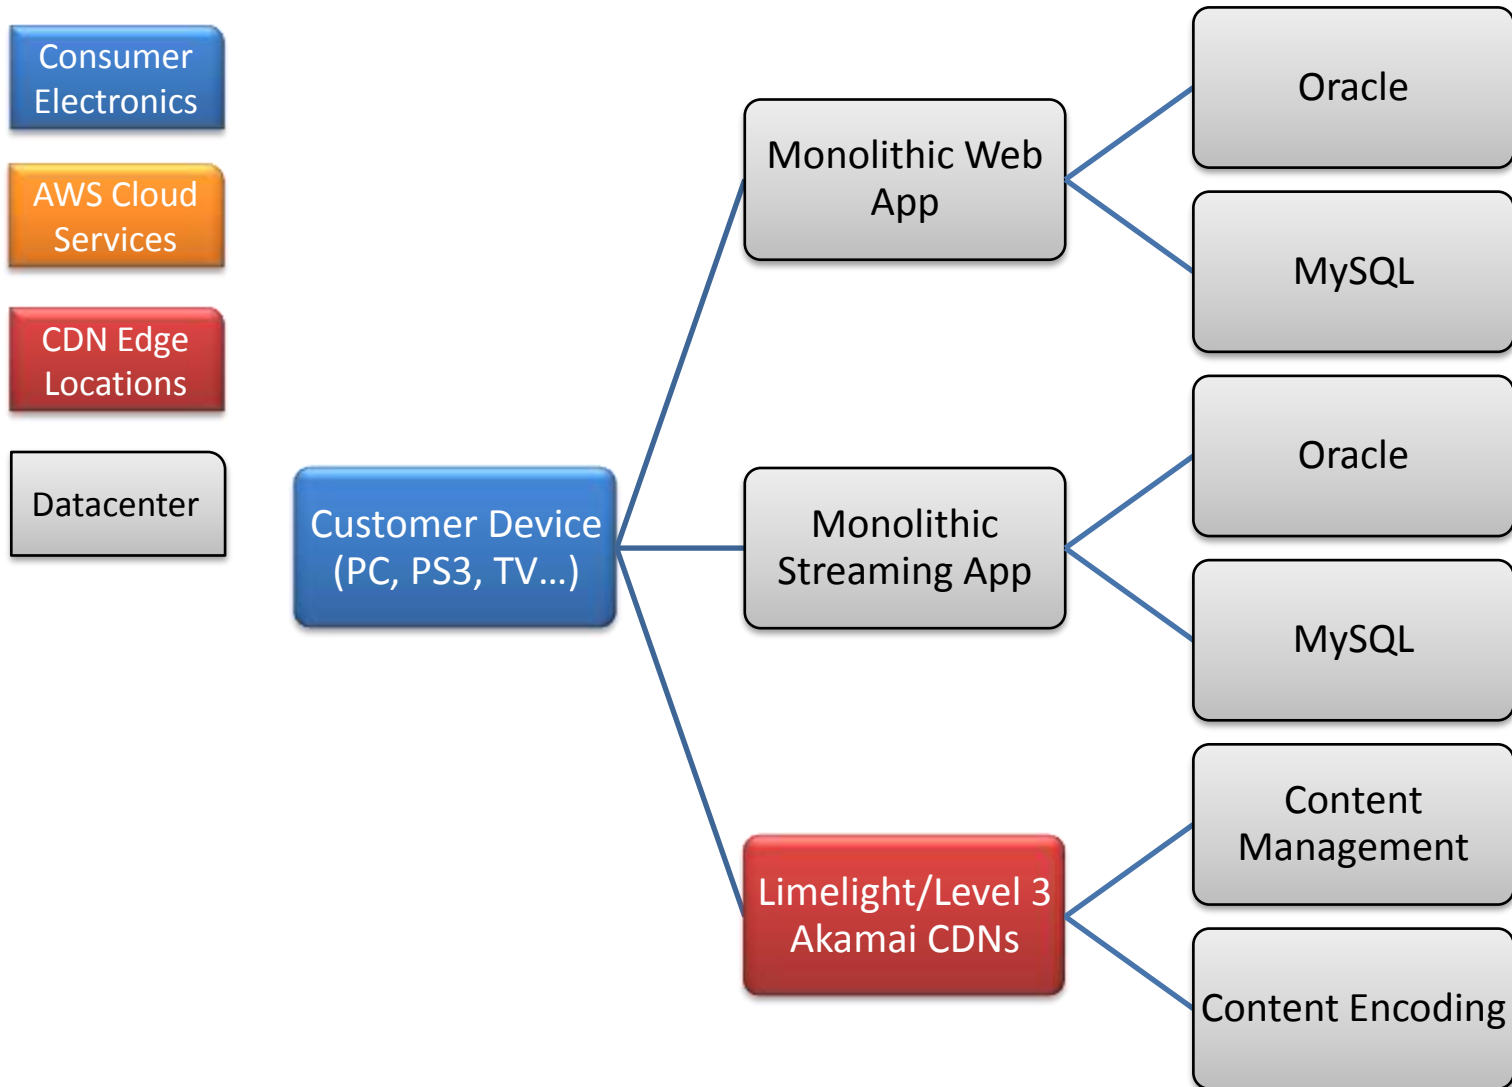
# Netflix Member Web Site Home Page

Personalization Driven – How Does It Work?

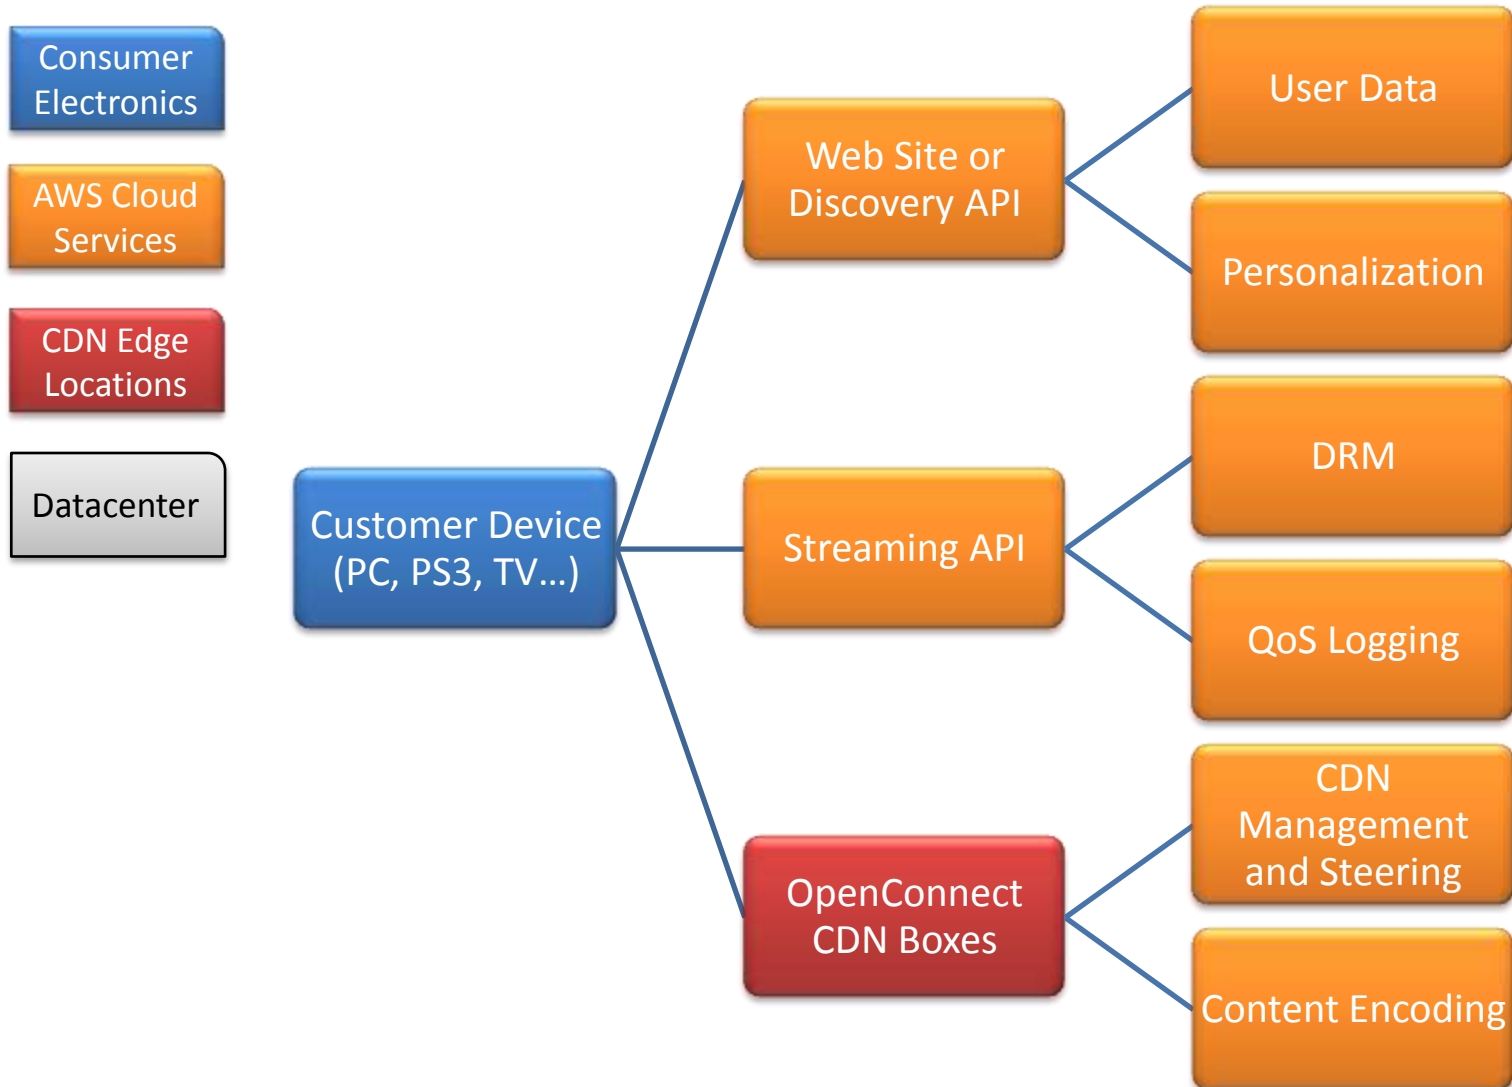
The screenshot displays the Netflix home page for a member named Adrian Cockcroft. The interface is divided into several sections:

- Header:** The Netflix logo is on the left, and the member's name "Adrian Cockcroft" with a dropdown arrow and a link to "Your Account & Help" is on the right.
- Navigation Bar:** A row of tabs includes "Watch Instantly", "Just for Kids", "Browse DVDs", "Your Queue", and "Taste Profile". A search bar on the right contains the text "Movies, TV shows, actors, directors, genres" and a magnifying glass icon.
- Genre Links:** Below the navigation bar are links for "Genres", "New Arrivals", and "Instantly to your TV".
- Recently Watched:** A section on the left showing a poster for "JOHN MAYALL & THE JOHN MAYALL BAND LIVE! JOHN CLAPTON, GARY BARBER, MICK TAYLOR JOHN BIRTHDAY CONCERT".
- Top 10 for Adrian:** A row of six posters: "DEFYING DISEASE TED TALKS", "SAM KINISON Breaking the Rules", "ANCIENT INVENTIONS OF WAR, SEX AND CITY LIFE", "ROBOTIC MACHINATIONS TED TALKS", "Ira Mitchell and Webb look", and "Bartleby".
- Friends' Favorites:** A section titled "Based on these friends:" with two small profile pictures. Below are six posters: "Breaking Bad", "LOST IN TRANSLATION", "THE TERMINATOR", "Audrey", "THE HUNT FOR RED OCTOBER", and "GOOD WILL HUNTING".

# How Netflix Used to Work



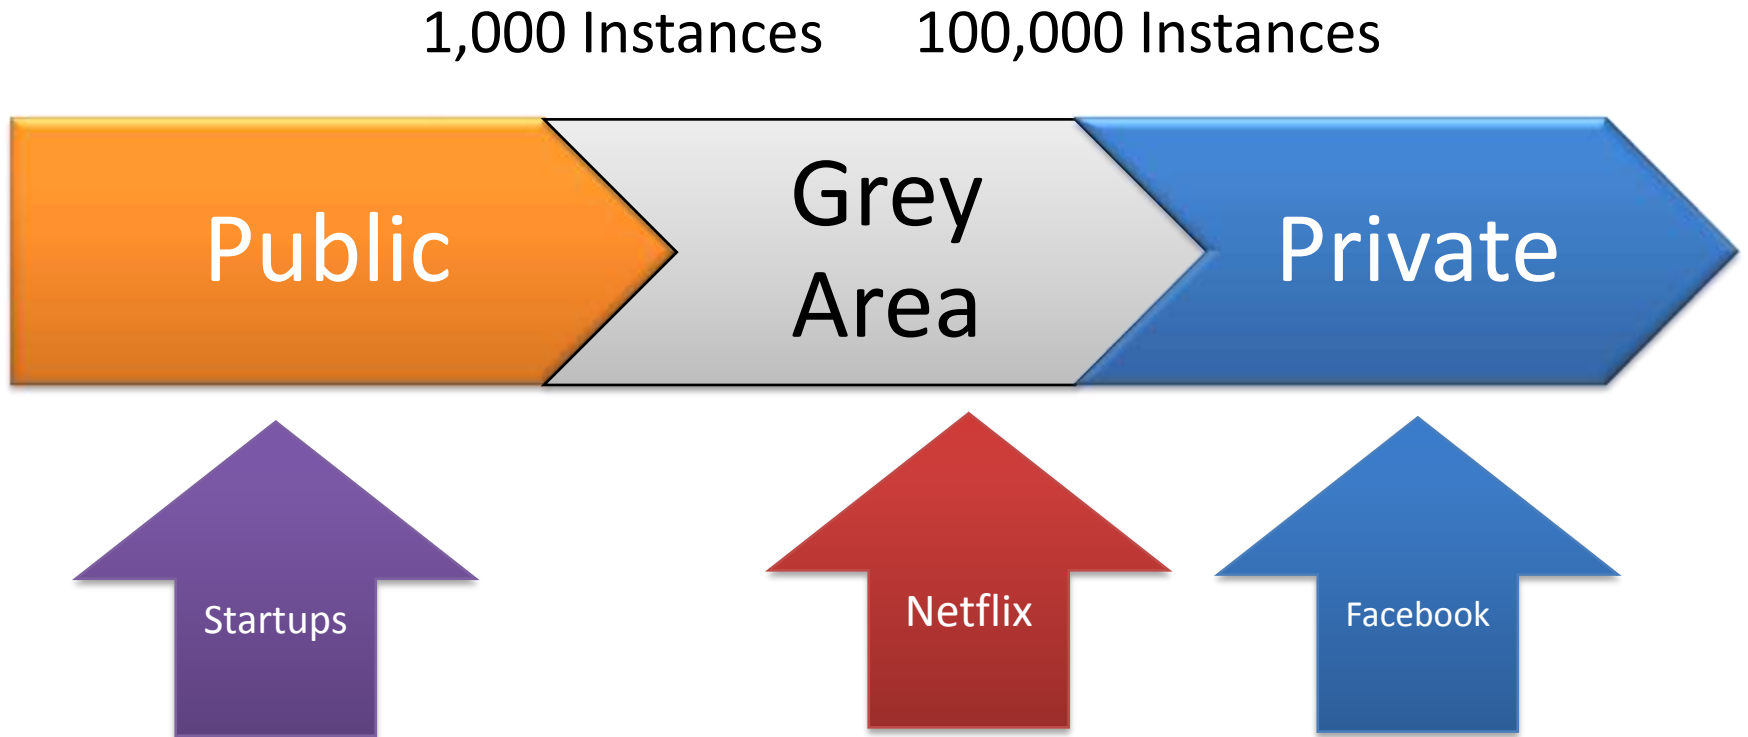
# How Netflix Streaming Works Today



# The DIY Question

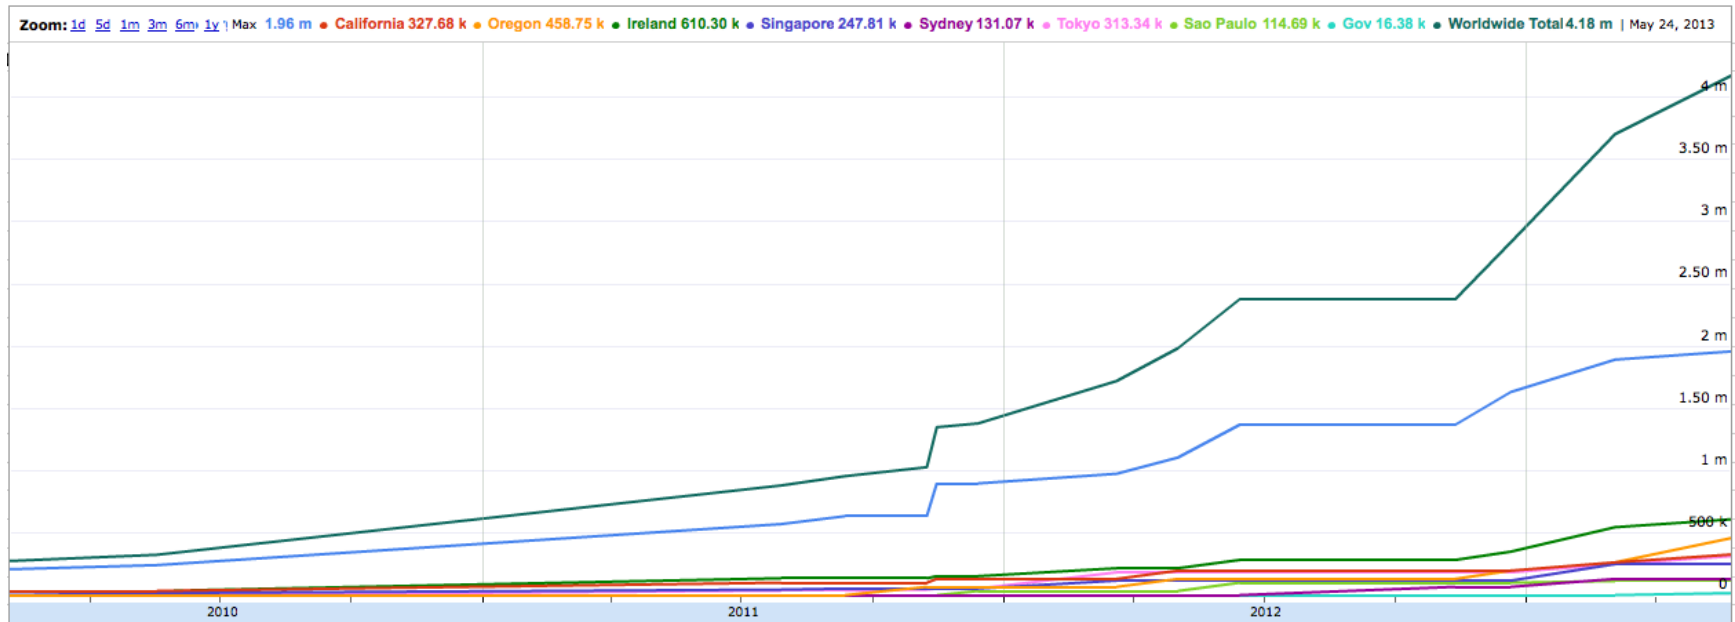
Why doesn't Netflix build and run its own cloud?

# Fitting Into Public Scale



# How big is Public?

AWS Maximum Possible Instance Count 4.2 Million – May 2013  
Growth >10x in Three Years, >2x Per Annum - <http://bit.ly/awsiprange>



AWS upper bound estimate based on the number of public IP Addresses  
Every provisioned instance gets a public IP by default (some VPC don't)



# The Alternative Supplier Question

What if there is no clear leader for a feature, or AWS doesn't have what we need?

# Things We Don't Use AWS For

SaaS Applications – Pagerduty, Appdynamics

Content Delivery Service

DNS Service

Nov  
2012  
Streaming  
Bandwidth

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	36.8%	Netflix	33.0%	Netflix	28.8%
2	HTTP	9.83%	YouTube	14.8%	YouTube	13.1%
3	Skype	4.76%	HTTP	12.0%	HTTP	11.7%
4	Netflix	4.51%	BitTorrent	5.89%	BitTorrent	10.3%
5	SSL	3.73%	iTunes	3.92%	iTunes	3.43%
6	YouTube	2.70%	MPEG	2.22%	SSL	2.23%
7	PPStream	1.65%	Flash Video	2.21%	MPEG	2.05%
8	Facebook	1.62%	SSL	1.97%	Flash Video	2.01%
9	Apple PhotoStream	1.46%	Amazon Video	1.75%	Facebook	1.50%
10	Dropbox	1.17%	Facebook	1.48%	RTMP	1.41%
	Top 10	68.24%	Top 10	79.01%	Top 10	76.54%



Table 3 - Top 10 Peak Period Applications (North America, Fixed Access)

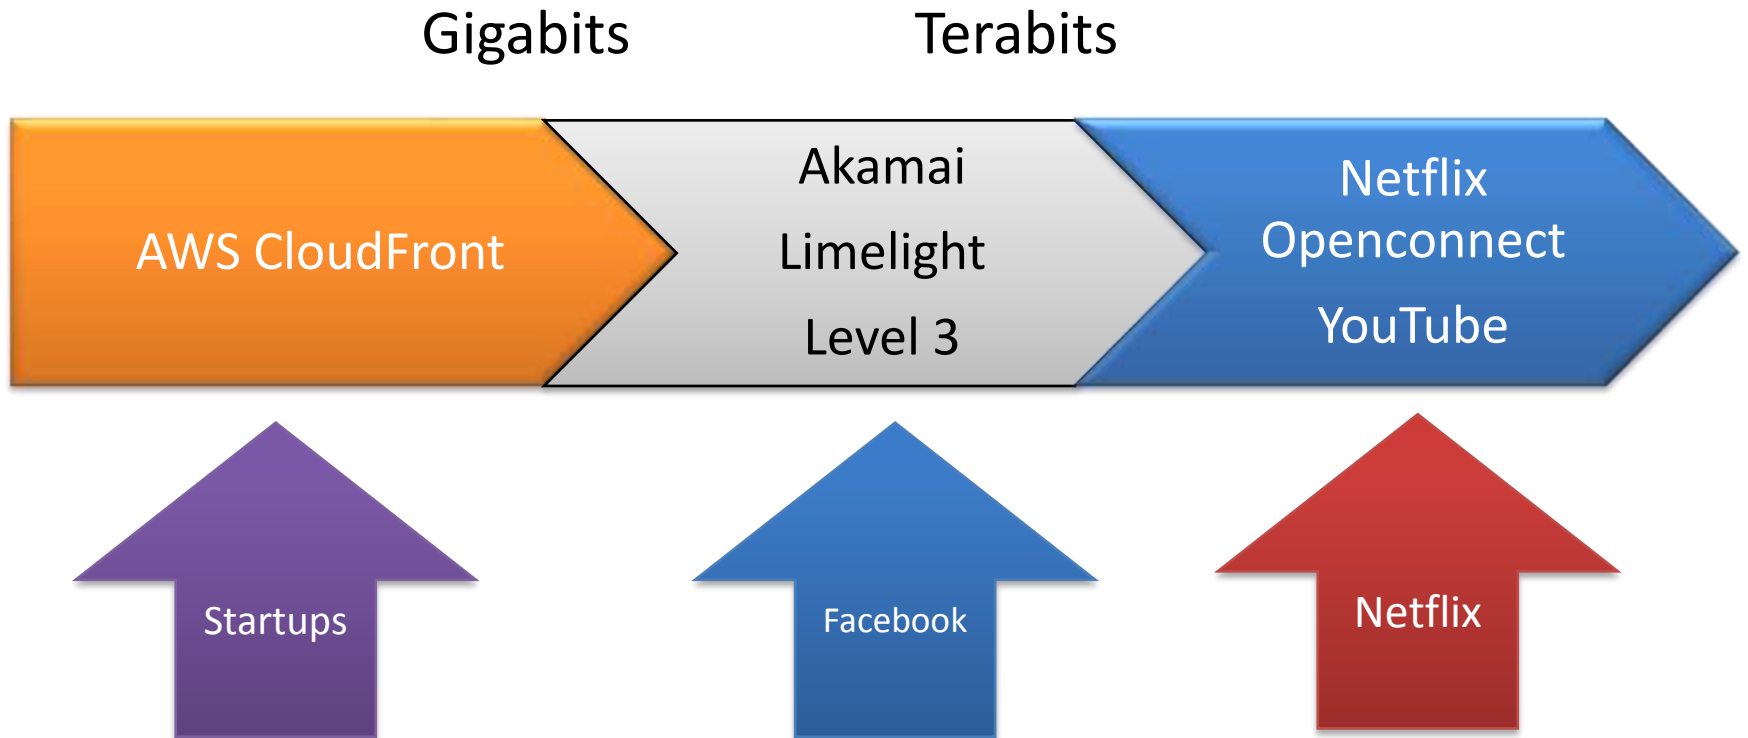
March  
2013

Mean  
Bandwidth  
+39% 6mo

Rank	Upstream		Downstream		Aggregate	
	Application	Share	Application	Share	Application	Share
1	BitTorrent	34.81%	Netflix	32.25%	Netflix	28.88%
2	HTTP	7.53%	YouTube	17.11%	YouTube	15.43%
3	SSL	5.81%	HTTP	11.11%	HTTP	10.66%
4	Netflix	5.38%	BitTorrent	5.57%	BitTorrent	9.23%
5	Skype	4.88%	MPEG	2.58%	SSL	2.39%
6	YouTube	3.71%	Hulu	2.41%	MPEG	2.30%
7	Facebook	1.71%	iTunes	1.90%	Hulu	2.16%
8	Apple Photostream	1.34%	SSL	1.89%	iTunes	1.71%
9	Dropbox	1.21%	Flash Video	1.72%	Flash Video	1.53%
10	Carbonite	0.99%	Facebook	1.48%	Facebook	1.52%
	Top 10	67.38%		78.03%		75.82%



# CDN Scale



# Content Delivery Service

Open Source Hardware Design + FreeBSD, bird, nginx  
see [openconnect.netflix.com](http://openconnect.netflix.com)

NETFLIX

## Open Connect

[Overview](#)[FAQ](#)[Peering Information](#)

### > Hardware Design

[Software Design](#)[Deployment Guide](#)[ISP Inquiry](#)

## Open Connect Appliance Hardware

### Objectives

When designing the Open Connect Appliance Hardware, we focused on these fundamental design goals:

- Very high storage density without sacrificing space and power efficiency. Our target was fitting 100 terabytes into a 4u chassis that is less than 2' deep.
- High throughput: 10 Gbps throughput via an optical network connection.
- Very low field maintenance: the appliance must tolerate a variety of hardware failures including hard drives, network optics, and power supply units.
- Simple racking and installation. Front mounted power and network ports are the only things to connect at install time.



Open Connect Appliances are servers based on commodity PC components (similar to the model used by all large scale content delivery networks). We were influenced by the excellent write-ups from the [Backblaze](#) team, and use a custom chassis due to a lack of ready made options for a compact unit.

To achieve over 100 TB of storage, spinning hard drives provide the highest affordable density, in particular 36 3TB SATA units. The hard drives are not hot swappable, as we wish to avoid the operational burden of field service. For lower power utilization and simpler sourcing we select commodity units from two vendors and use software to manage failure modes and avoid field replacement. Dead drives reduce the total storage available for the system, but don't take it offline. We also add 1 TB of flash storage (2 solid state drives) for system files, logs and popular content. To augment the motherboard attached controller, we use two 16 port LSI SAS controller cards that connect directly to the SATA drives. This avoids I/O bottlenecks of SATA multipliers or SAS expanders, and also reduces system complexity.

From a compute point of view, the system has modest requirements moving bits from the storage to network packets on the interface. To reduce the power usage and hence also cooling requirement (which in turn reduces vibration from case fans) we use a single low power 4 core Intel Sandy Bridge CPU on a small form factor [Supermicro](#) mATX board with the full 32 GB of RAM installed.

We use redundant, hot swappable power supply units that have interchangeable AC and DC options for maximum installation flexibility. [Zipay](#) reversed the fan rotation of the units to allow mounting at the front of the case, and thus allow network and power connects to be positioned here.

The network card has two 10 Gbps modules, which can power a variety of SR and LR optic modules, for installation flexibility and scalable interconnection.

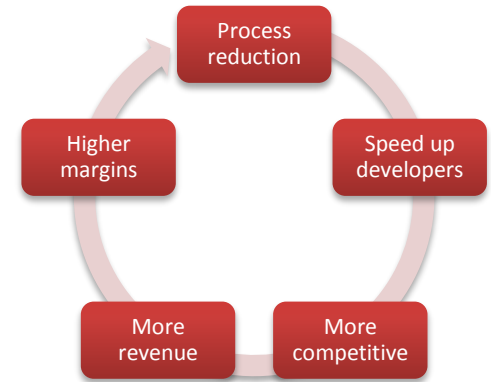
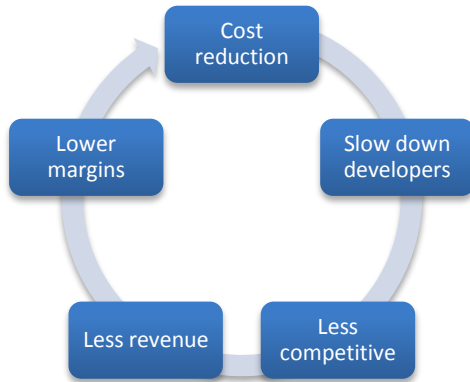


# DNS Service

AWS Route53 is missing too many features (for now)

Multiple vendor strategy Dyn, Ultra, Route53

Abstracted (broken) DNS APIs with Denominator



# What Changed?

Get out of the way of innovation

Best of breed, by the hour

Choices based on scale

# Availability Questions

Is it running yet?

How many places is it running in?

How far apart are those places?



# The STRANGE WORLD of the FUTURE

*STRANDED without video!  
No way to fill their empty hours!  
They were victims of...*

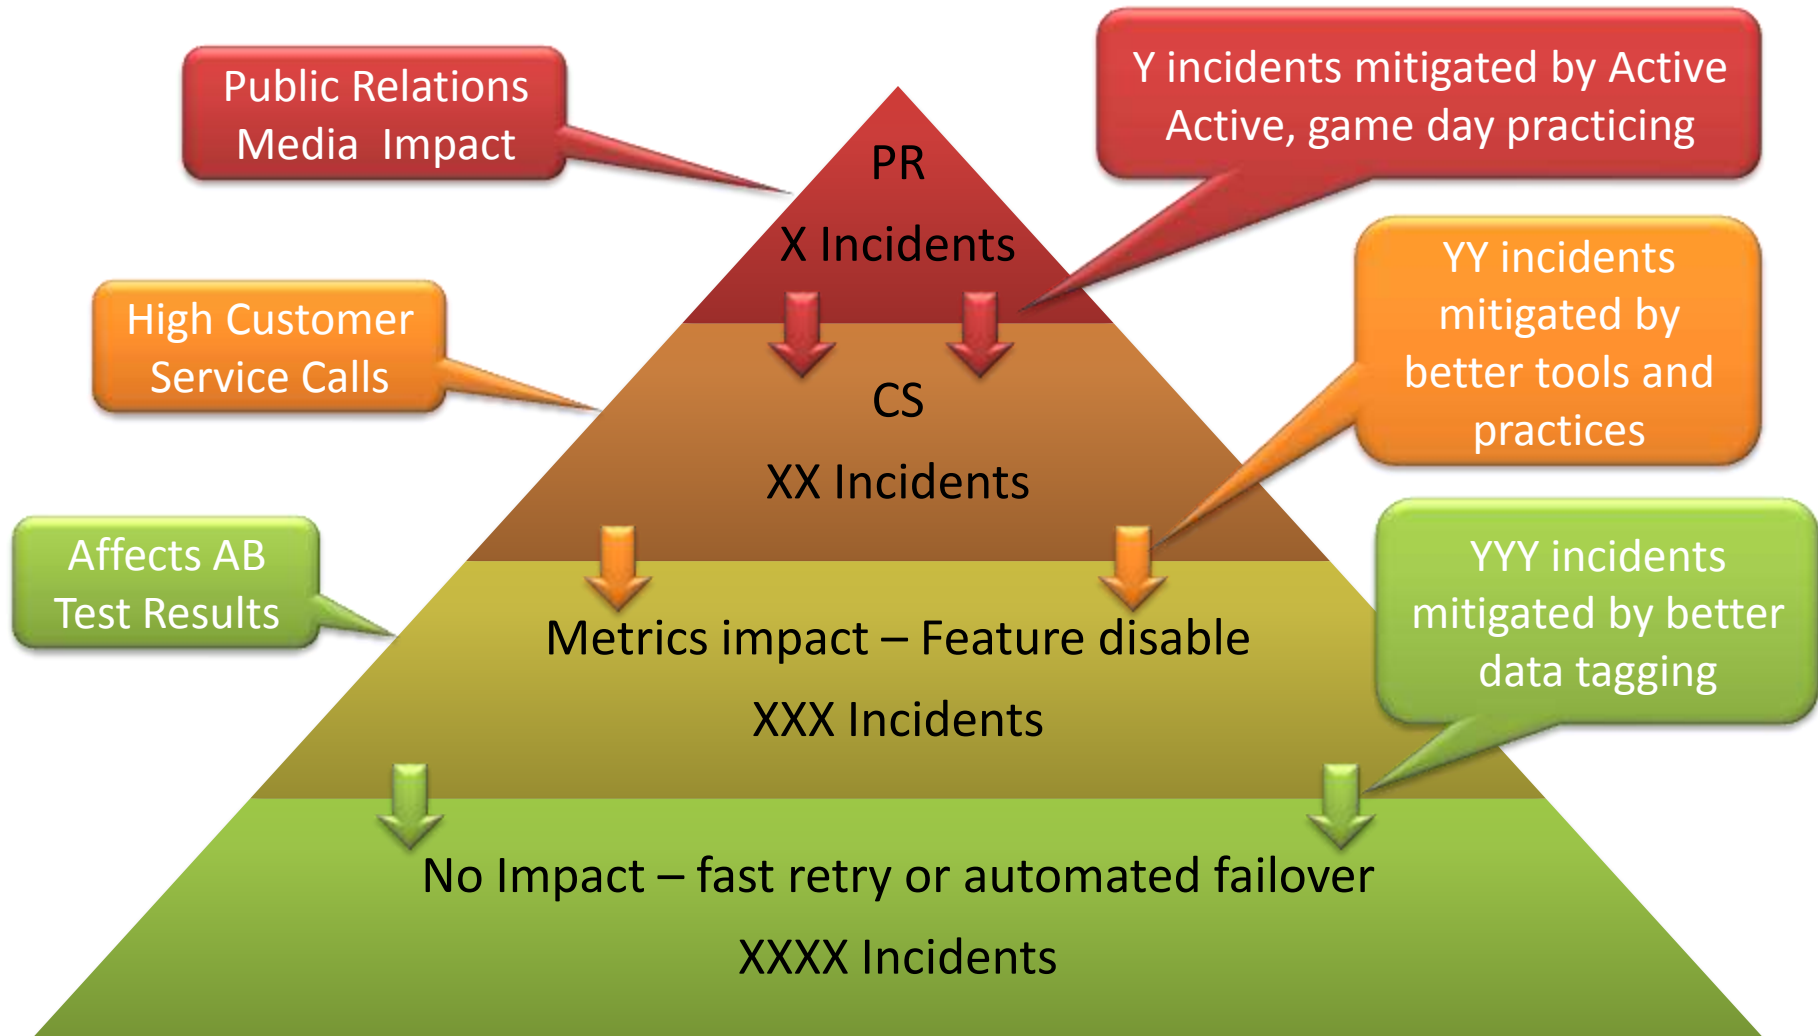
## THE CLOUD OF BROKEN STREAMS

CREATED WITH PULP-O-MIZER COVER MAKER

# Netflix Outages

- Running very fast with scissors
  - Mostly self inflicted – bugs, mistakes from pace of change
  - Some caused by AWS bugs and mistakes
- Incident Life-cycle Management by Platform Team
  - No runbooks, no operational changes by the SREs
  - Tools to identify what broke and call the right developer
- Next step is multi-region active/active
  - Investigating and building in stages during 2013
  - Could have prevented some of our 2012 outages

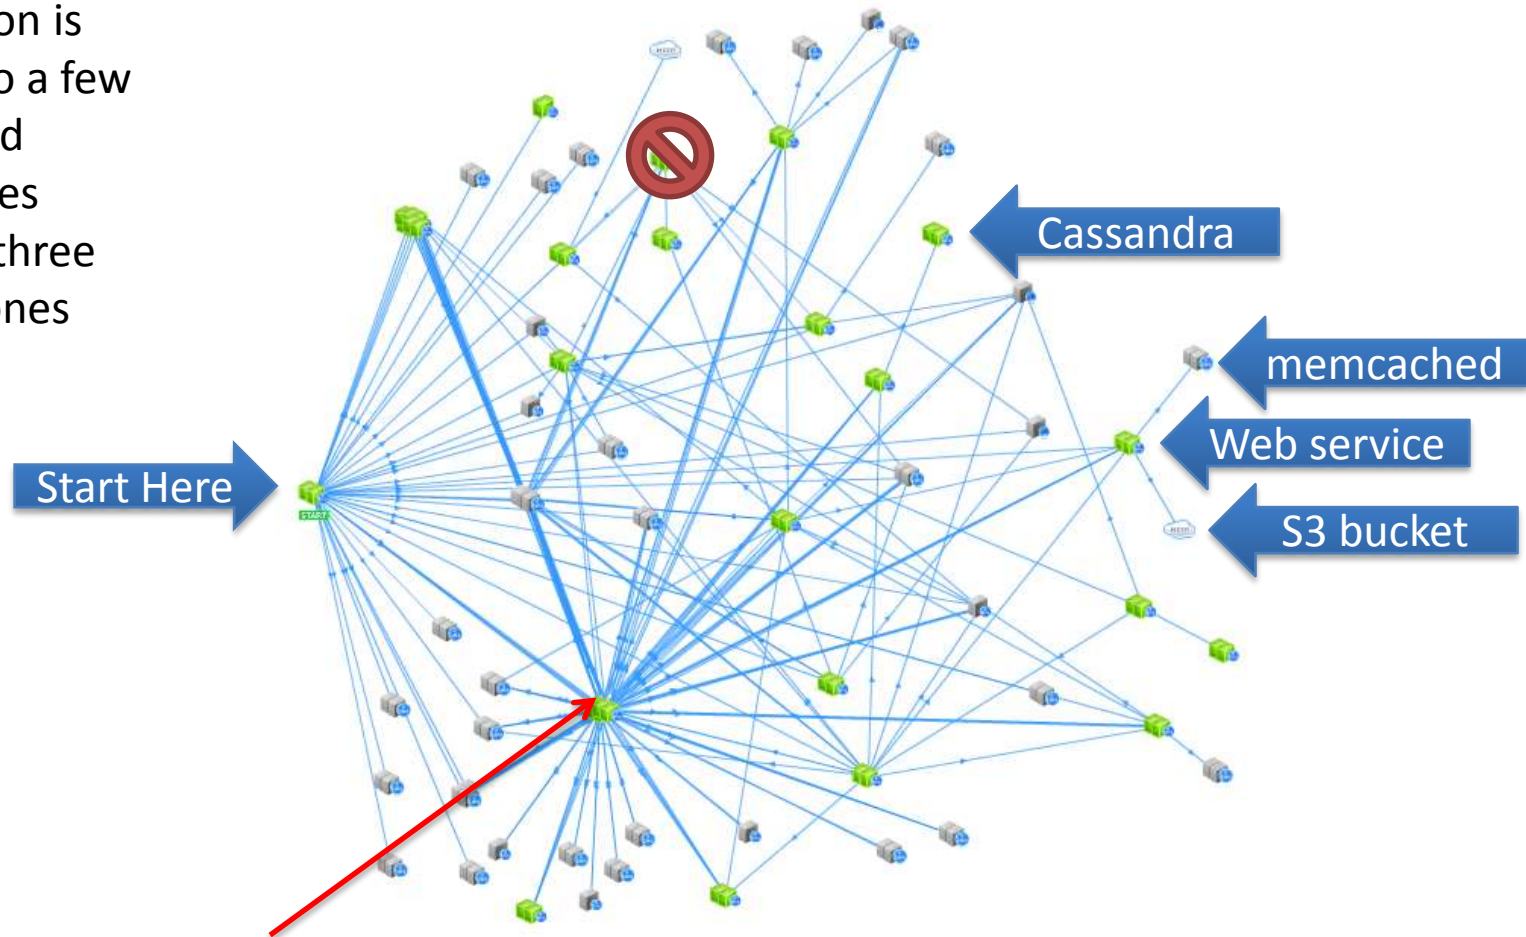
# Incidents – Impact and Mitigation



# Real Web Server Dependencies Flow

(Netflix Home page business transaction as seen by AppDynamics)

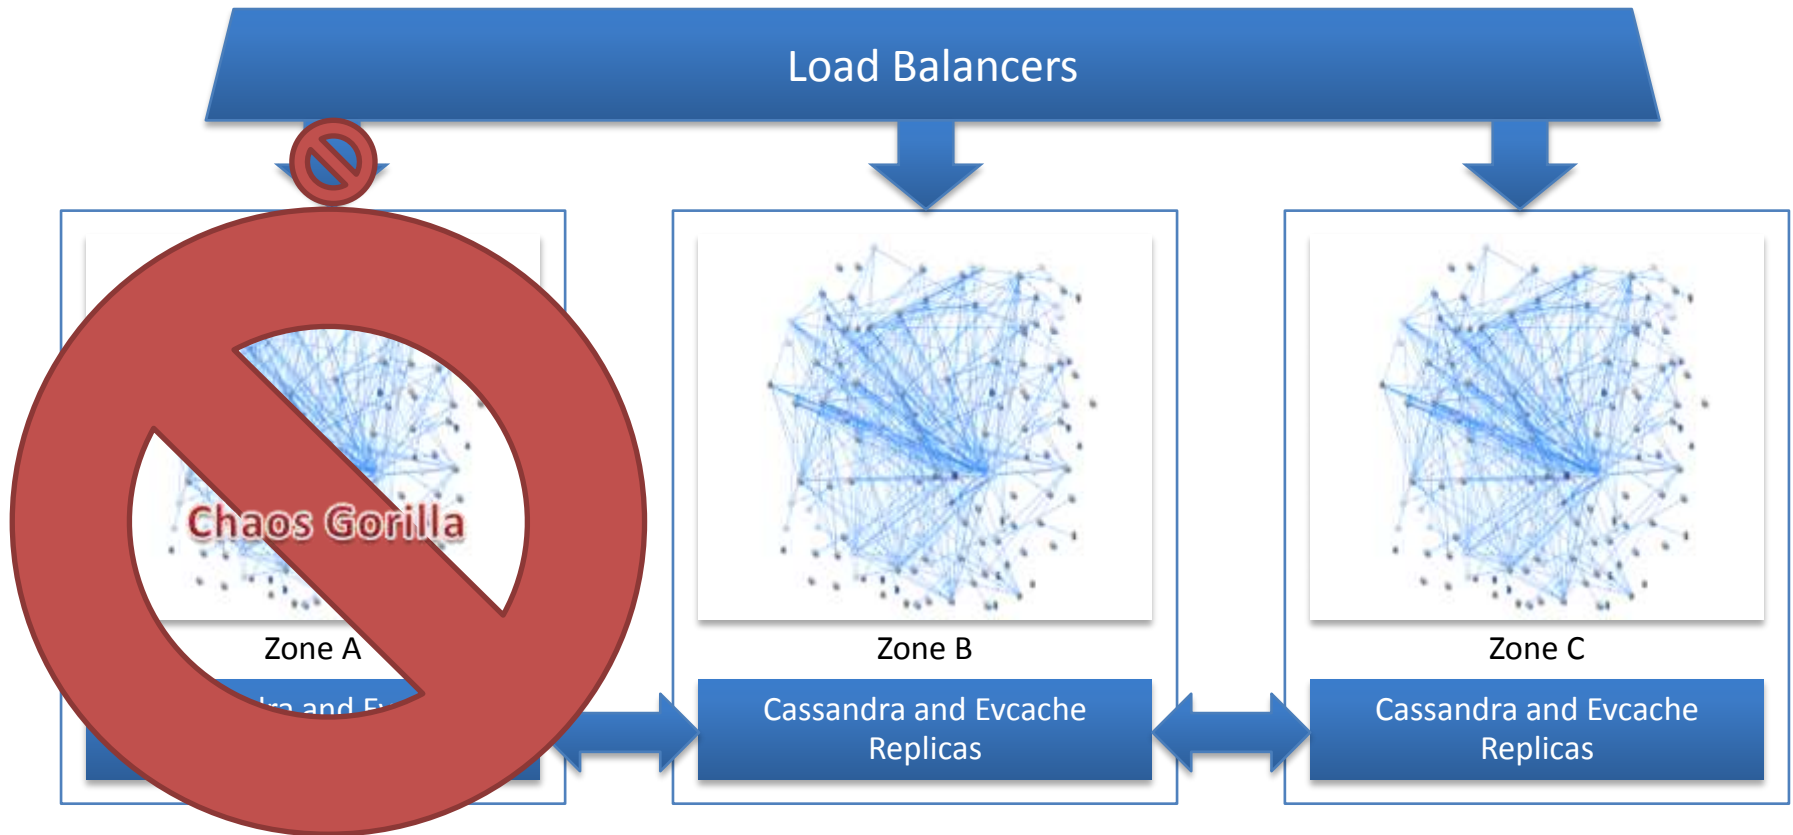
Each icon is  
three to a few  
hundred  
instances  
across three  
AWS zones



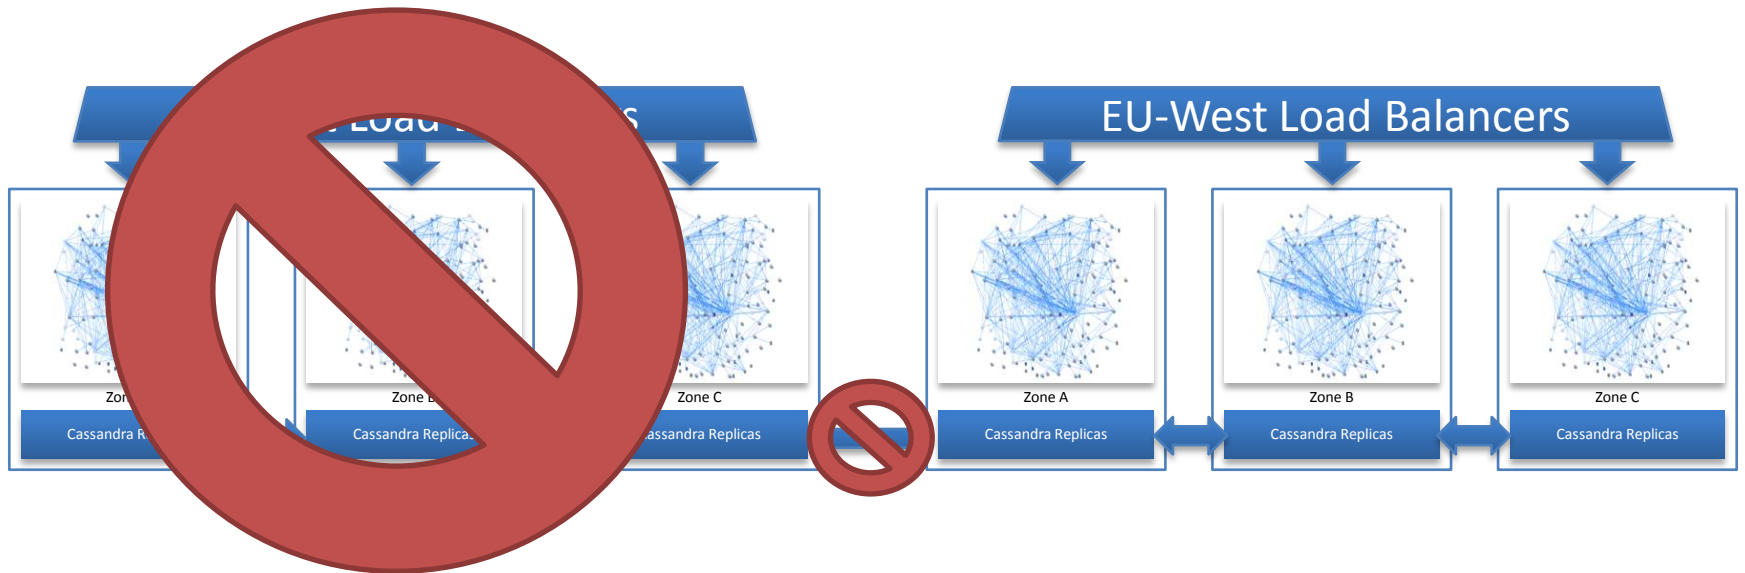
Personalization movie group choosers  
(for US, Canada and Latam)

# Three Balanced Availability Zones

Test with Chaos Gorilla



# Isolated Regions



More?

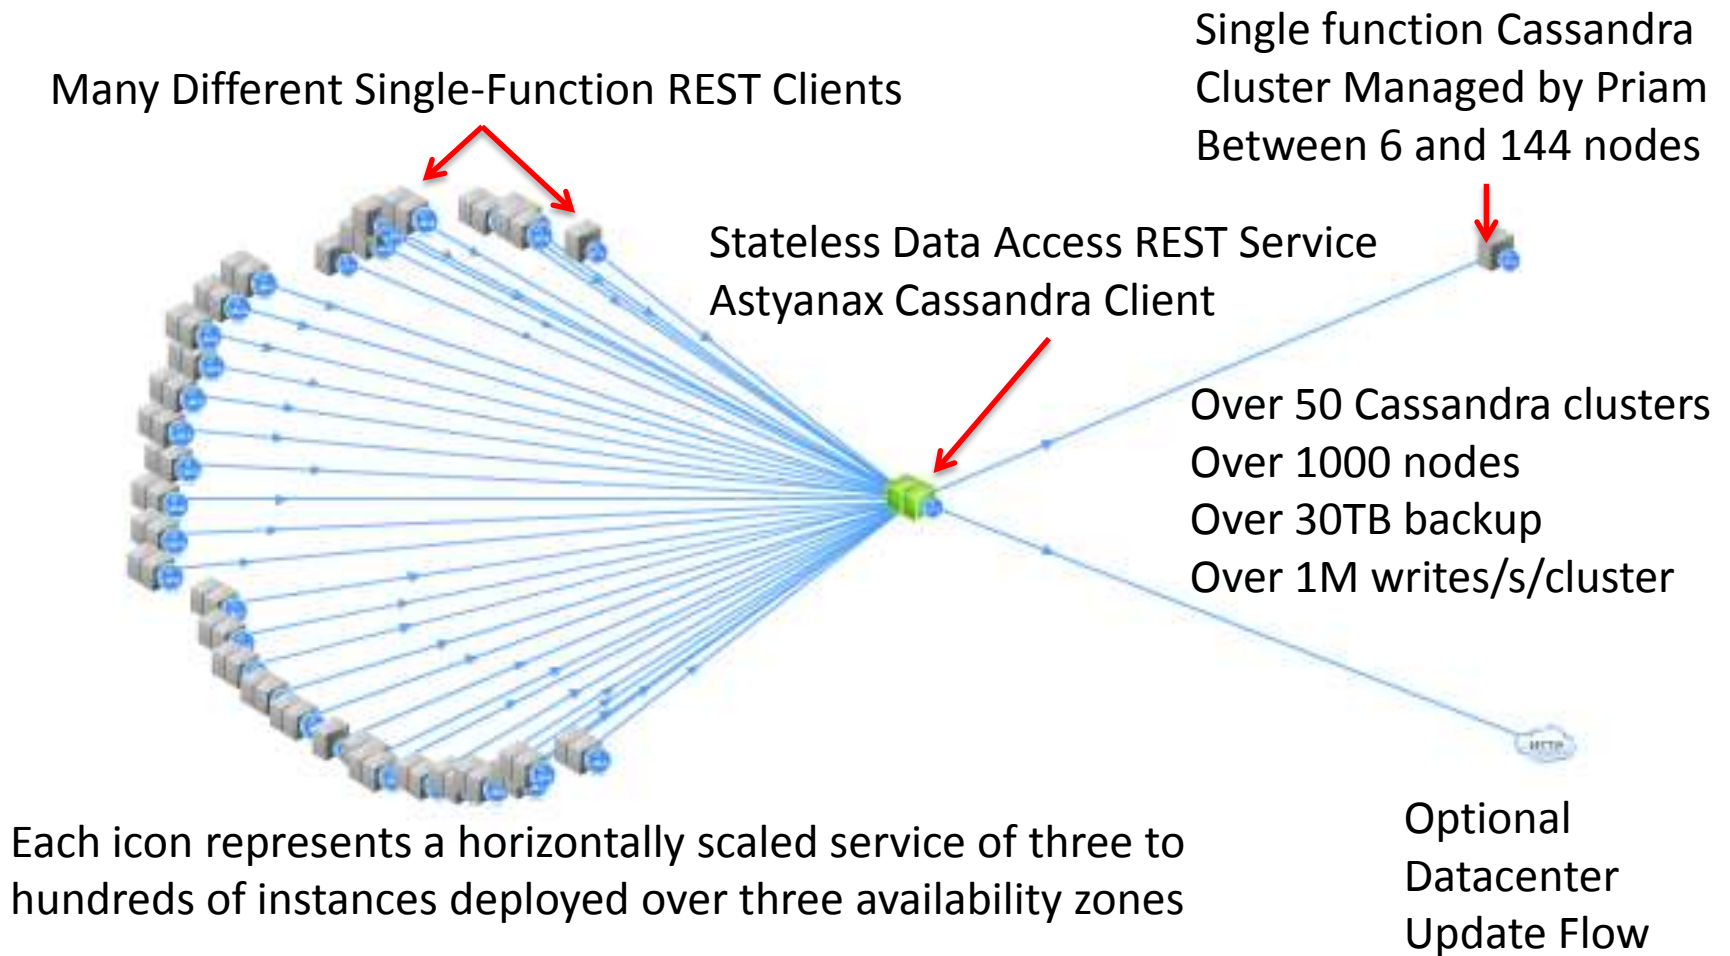
# Highly Available NoSQL Storage

A highly scalable, available and durable deployment pattern based on Apache Cassandra



# Single Function Micro-Service Pattern

One keyspace, replaces a single table or materialized view





# Stateless Micro-Service Architecture

## Linux Base AMI (CentOS or Ubuntu)

Optional  
Apache  
frontend,  
memcached,  
non-java apps

Monitoring  
Log rotation  
to S3  
AppDynamics  
machineagent  
Epic/Atlas

### Java (JDK 6 or 7)

AppDynamics  
appagent  
monitoring

GC and thread  
dump logging

### Tomcat

Application war file, base  
servlet, platform, client  
interface jars, Astyanax

Healthcheck, status  
servlets, JMX interface,  
Servo autoscale

# Cassandra Instance Architecture

Linux Base AMI (CentOS or Ubuntu)

Tomcat and  
Priam on JDK  
Healthcheck,  
Status

Monitoring  
AppDynamics  
machineagent  
Epic/Atlas

Java (JDK 7)

AppDynamics  
appagent  
monitoring

GC and thread  
dump logging

Cassandra Server

Local Ephemeral Disk Space – 2TB of SSD or 1.6TB disk  
holding Commit log and SSTables

# Apache Cassandra

- Scalable and Stable in large deployments
  - No additional license cost for large scale!
  - Optimized for “OLTP” vs. Hbase optimized for “DSS”
- Available during Partition (AP from CAP)
  - Hinted handoff repairs most transient issues
  - Read-repair and periodic repair keep it clean
- Quorum and Client Generated Timestamp
  - Read after write consistency with 2 of 3 copies
  - Latest version includes Paxos for stronger transactions

# Astyanax Cassandra Client for Java

Available at <http://github.com/netflix>

- Features
  - Abstraction of connection pool from RPC protocol
  - Fluent Style API
  - Operation retry with backoff
  - Token aware
  - Batch manager
  - Many useful recipes
  - New: Entity Mapper based on JPA annotations



# C\* Astyanax Recipes

- Distributed row lock (without needing zookeeper)
- Multi-region row lock
- Uniqueness constraint
- Multi-row uniqueness constraint
- Chunked and multi-threaded large file storage
- Reverse index search
- All rows query
- Durable message queue
- Contributed: High cardinality reverse index

# Astyanax - Cassandra Write Data Flows

Single Region, Multiple Availability Zone, Token Aware

1. Client Writes to local coordinator
2. Coordinator writes to other zones
3. Nodes return ack
4. Data written to internal commit log disks (no more than 10 seconds later)



If a node goes offline, hinted handoff completes the write when the node comes back up.

Requests can choose to wait for one node, a quorum, or all nodes to ack the write

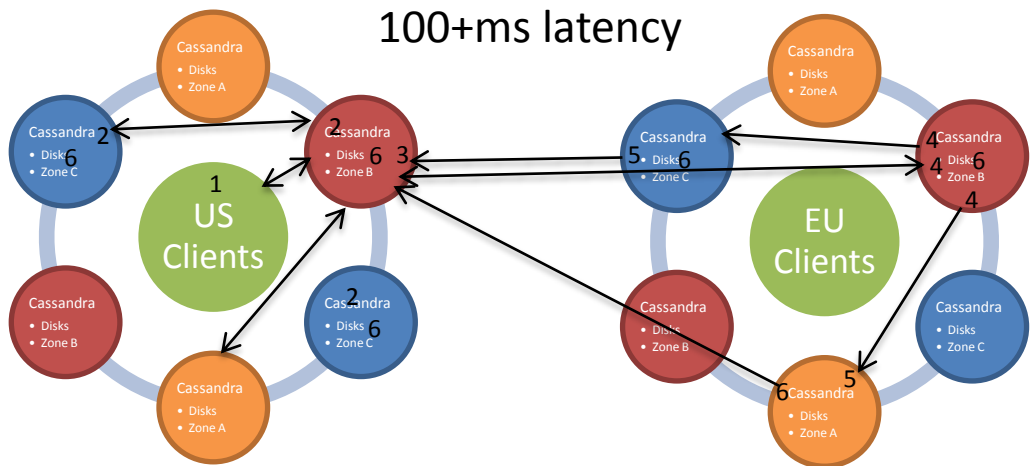
SSTable disk writes and compactions occur asynchronously

# Data Flows for Multi-Region Writes

## Token Aware, Consistency Level = Local Quorum

1. Client writes to local replicas
2. Local write acks returned to Client which continues when 2 of 3 local nodes are committed
3. Local coordinator writes to remote coordinator.
4. When data arrives, remote coordinator node acks and copies to other remote zones
5. Remote nodes ack to local coordinator
6. Data flushed to internal commit log disks (no more than 10 seconds later)

If a node or region goes offline, hinted handoff completes the write when the node comes back up. Nightly global compare and repair jobs ensure everything stays consistent.



# Cassandra at Scale

Benchmarking to Retire Risk



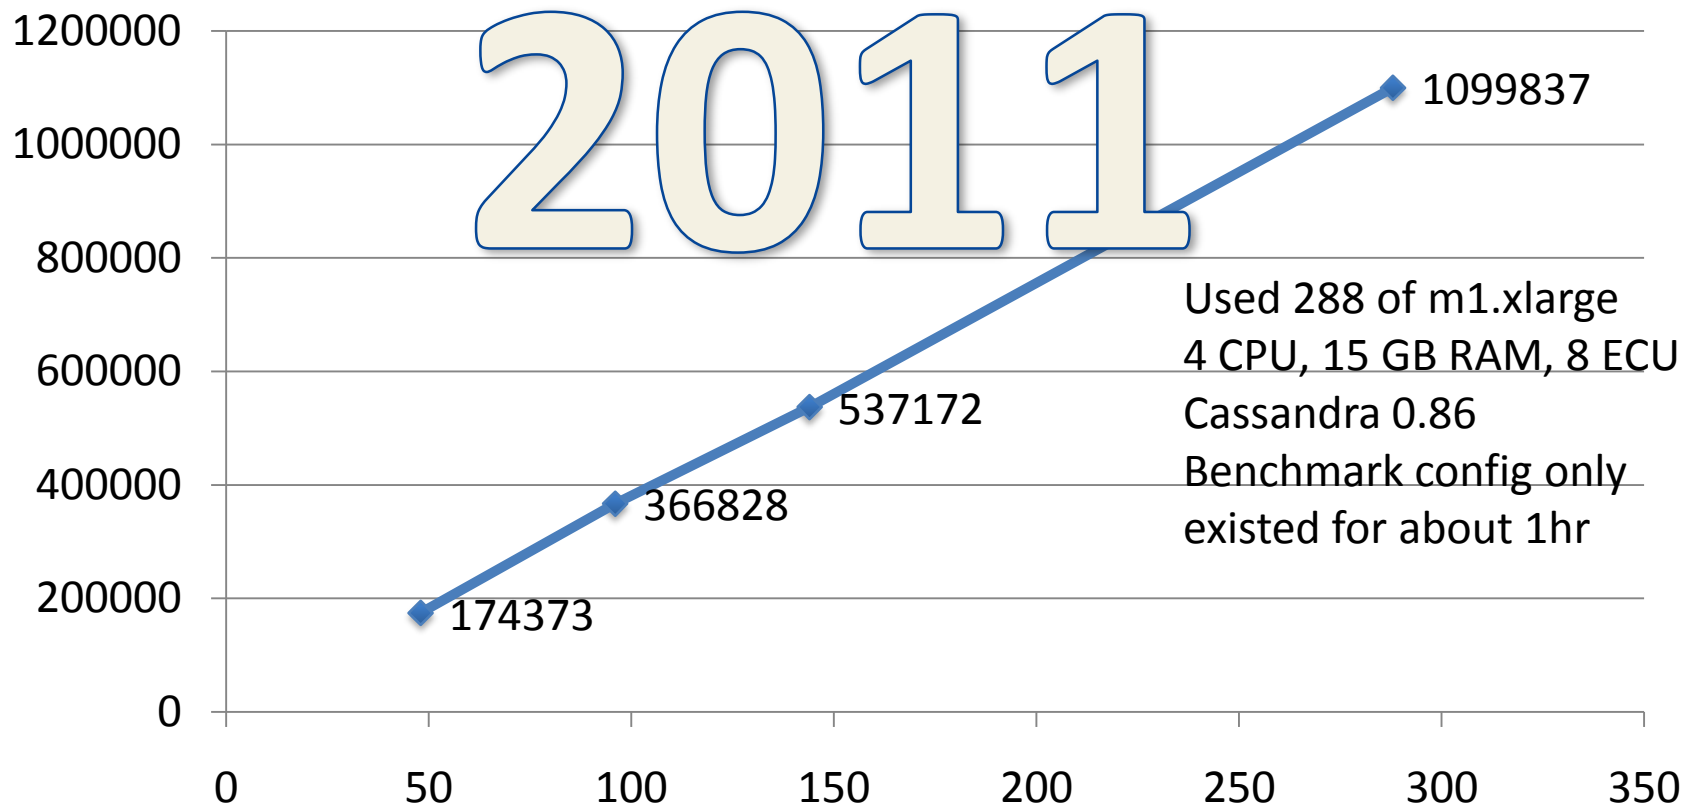
More?



# Scalability from 48 to 288 nodes on AWS

<http://techblog.netflix.com/2011/11/benchmarking-cassandra-scalability-on.html>

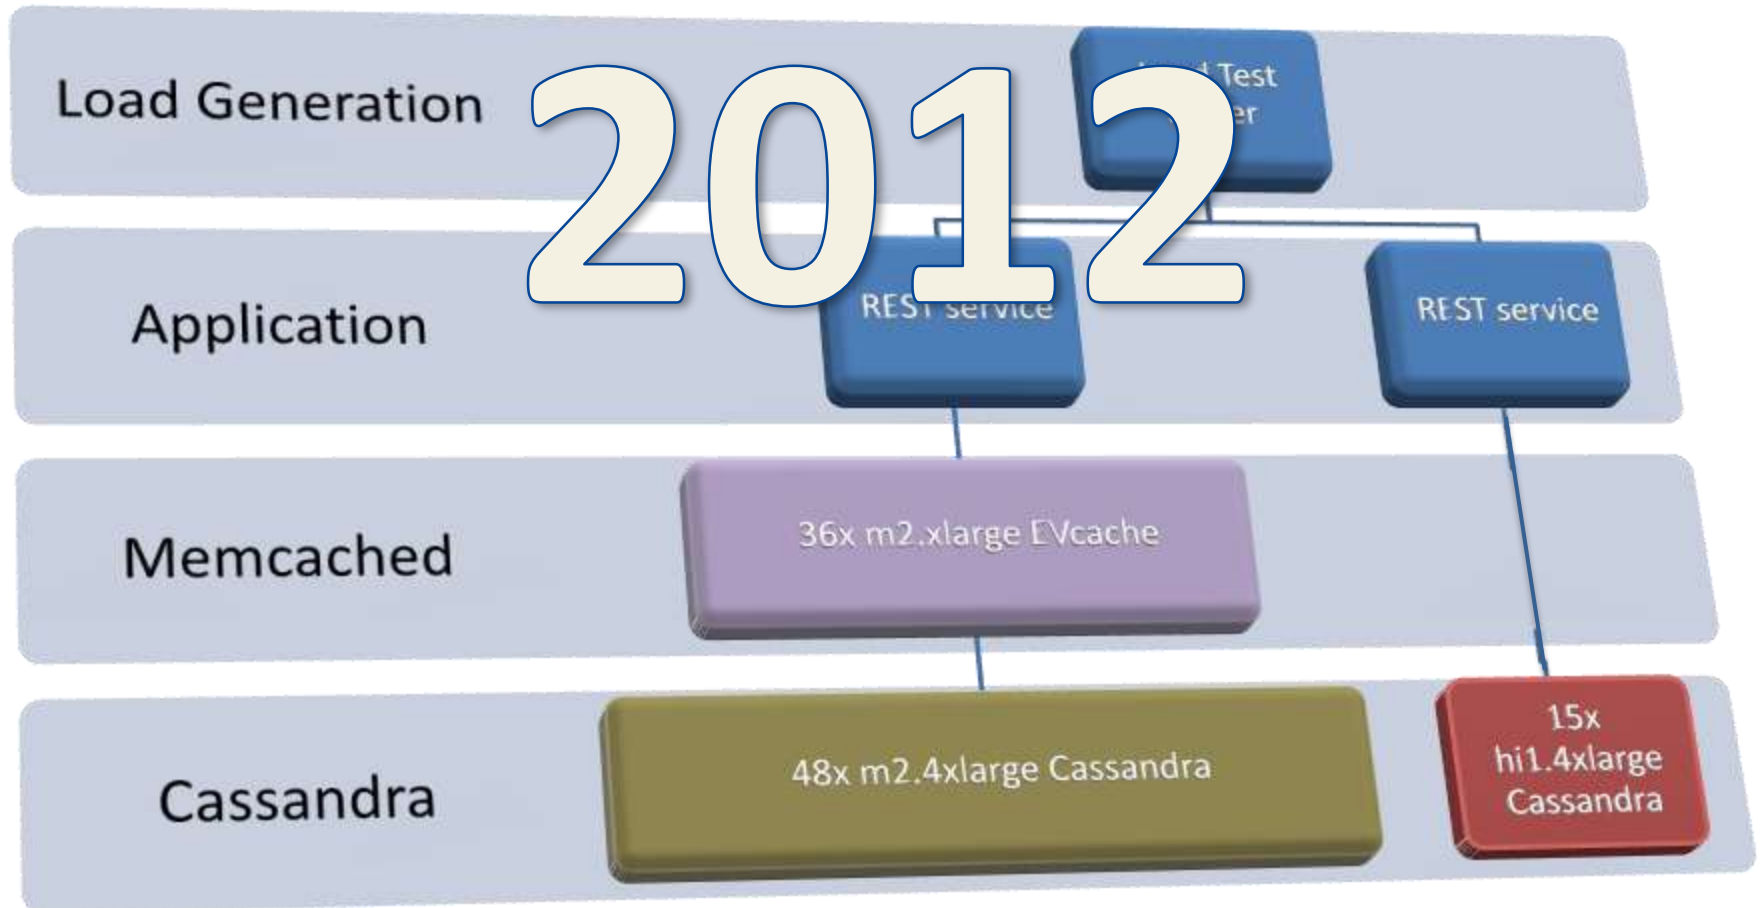
Client Writes/s by node count – Replication Factor = 3



# Cassandra Disk vs. SSD Benchmark

Same Throughput, Lower Latency, Half Cost

<http://techblog.netflix.com/2012/07/benchmarking-high-performance-io-with.html>



# 2013 - Cross Region Use Cases

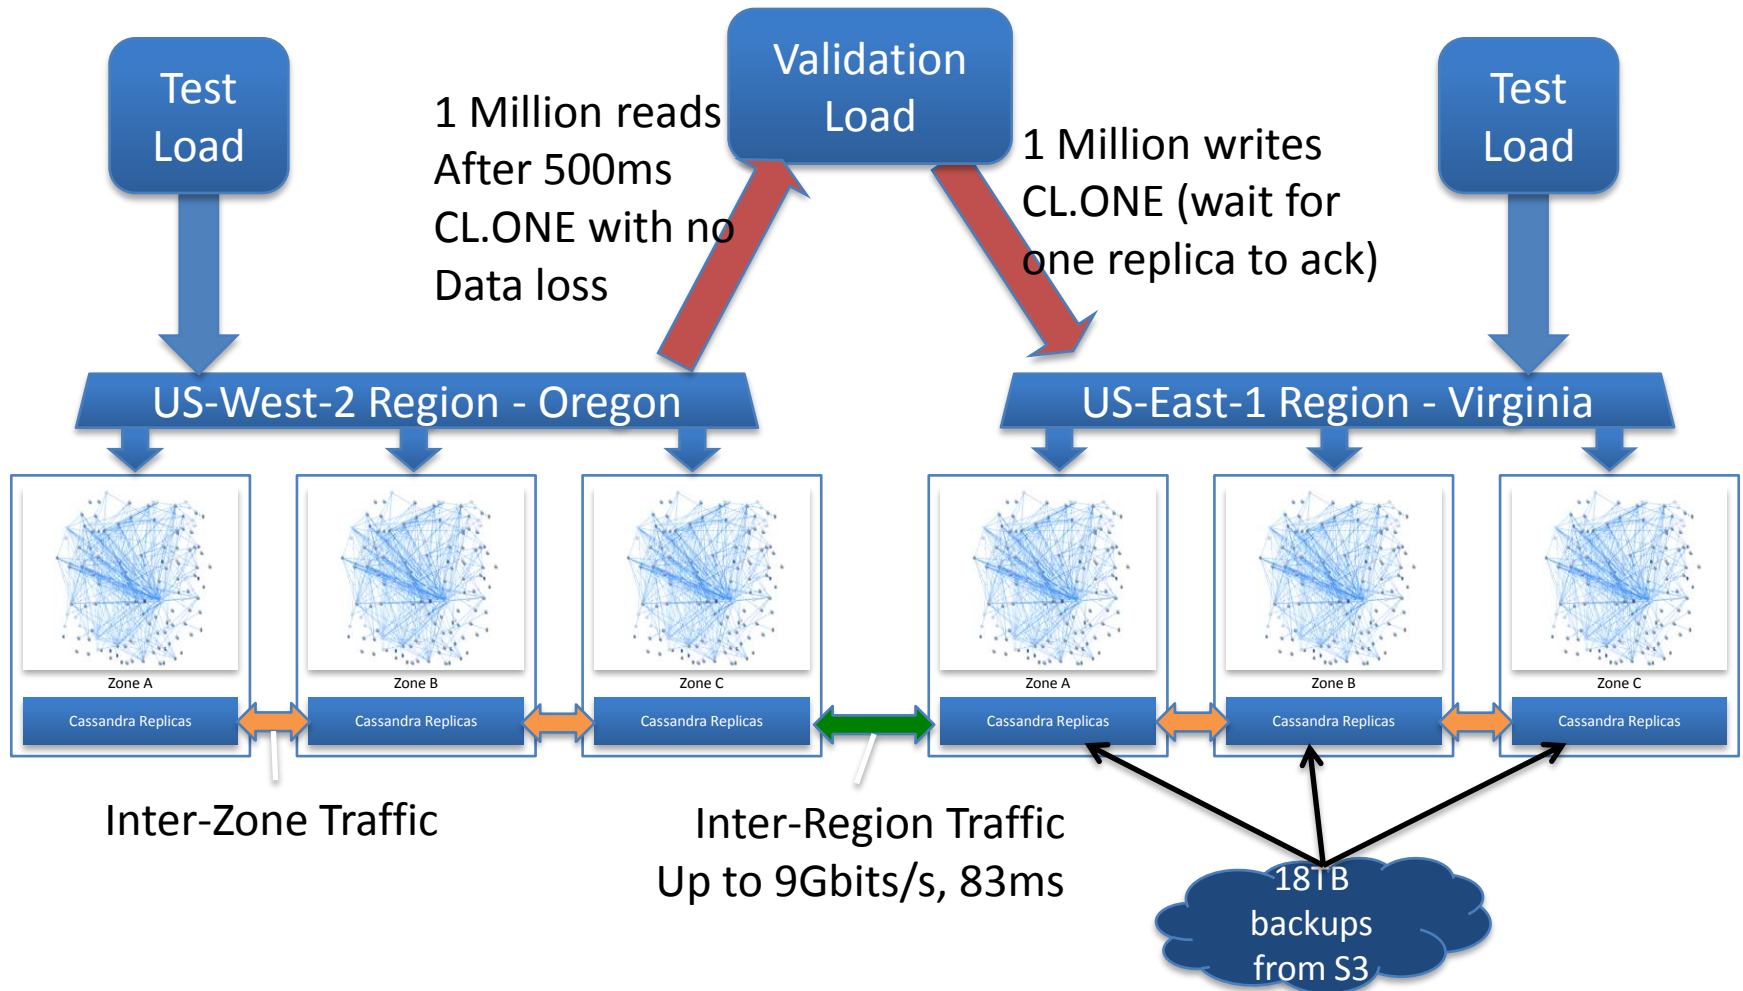
- Geographic Isolation
  - US to Europe replication of subscriber data
  - Read intensive, low update rate
  - Production use since late 2011
- Redundancy for regional failover
  - US East to US West replication of everything
  - Includes write intensive data, high update rate
  - Testing now

# Benchmarking Global Cassandra

Write intensive test of cross region replication capacity

16 x hi1.4xlarge SSD nodes per zone = 96 total

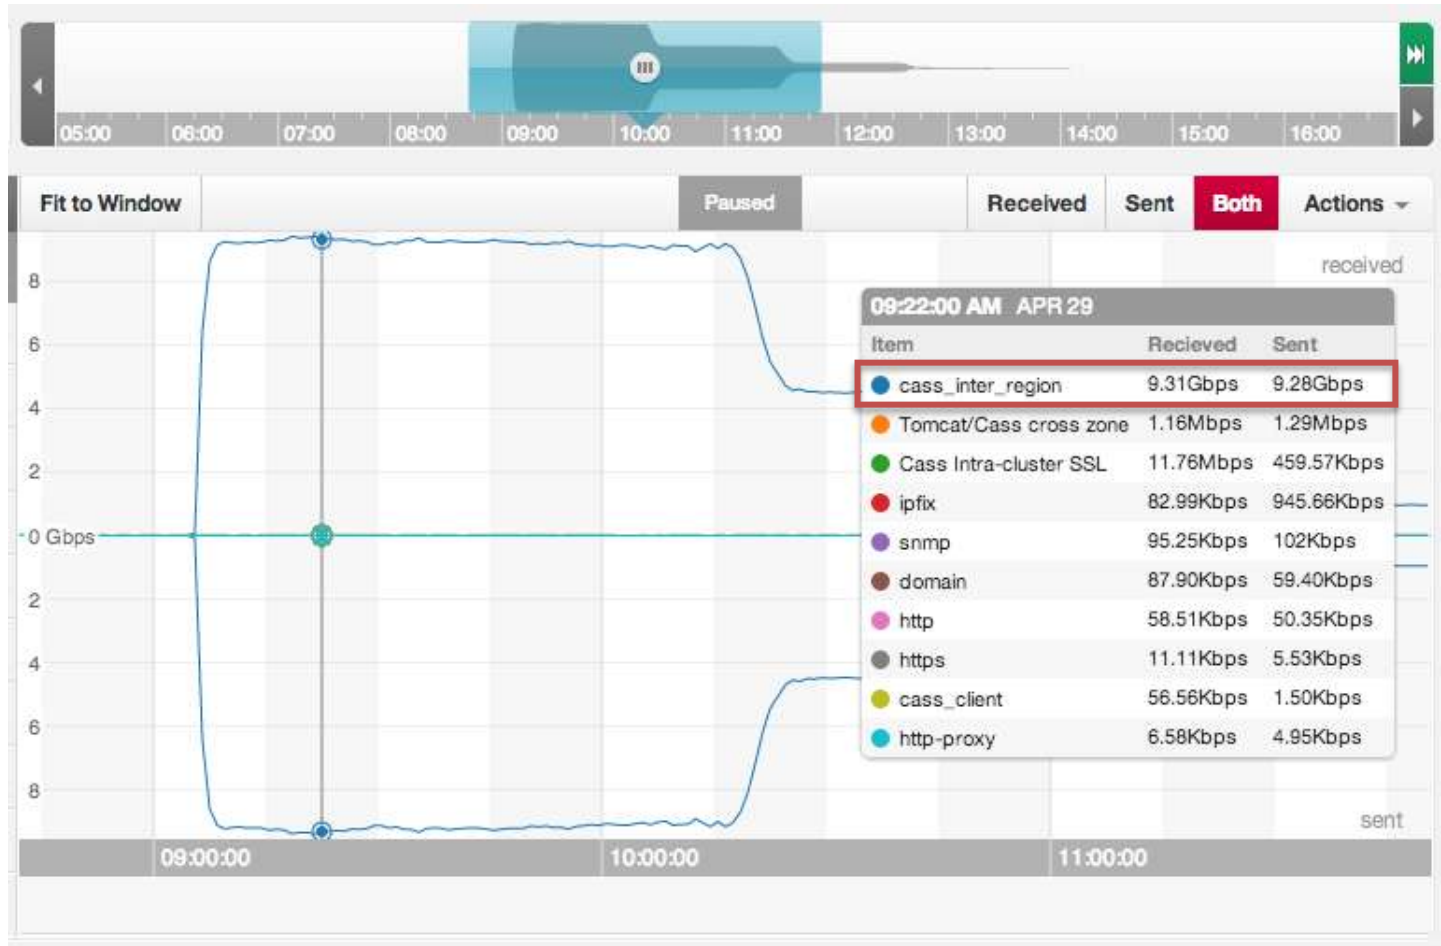
192 TB of SSD in six locations up and running Cassandra in 20 minutes



# Copying 18TB from East to West

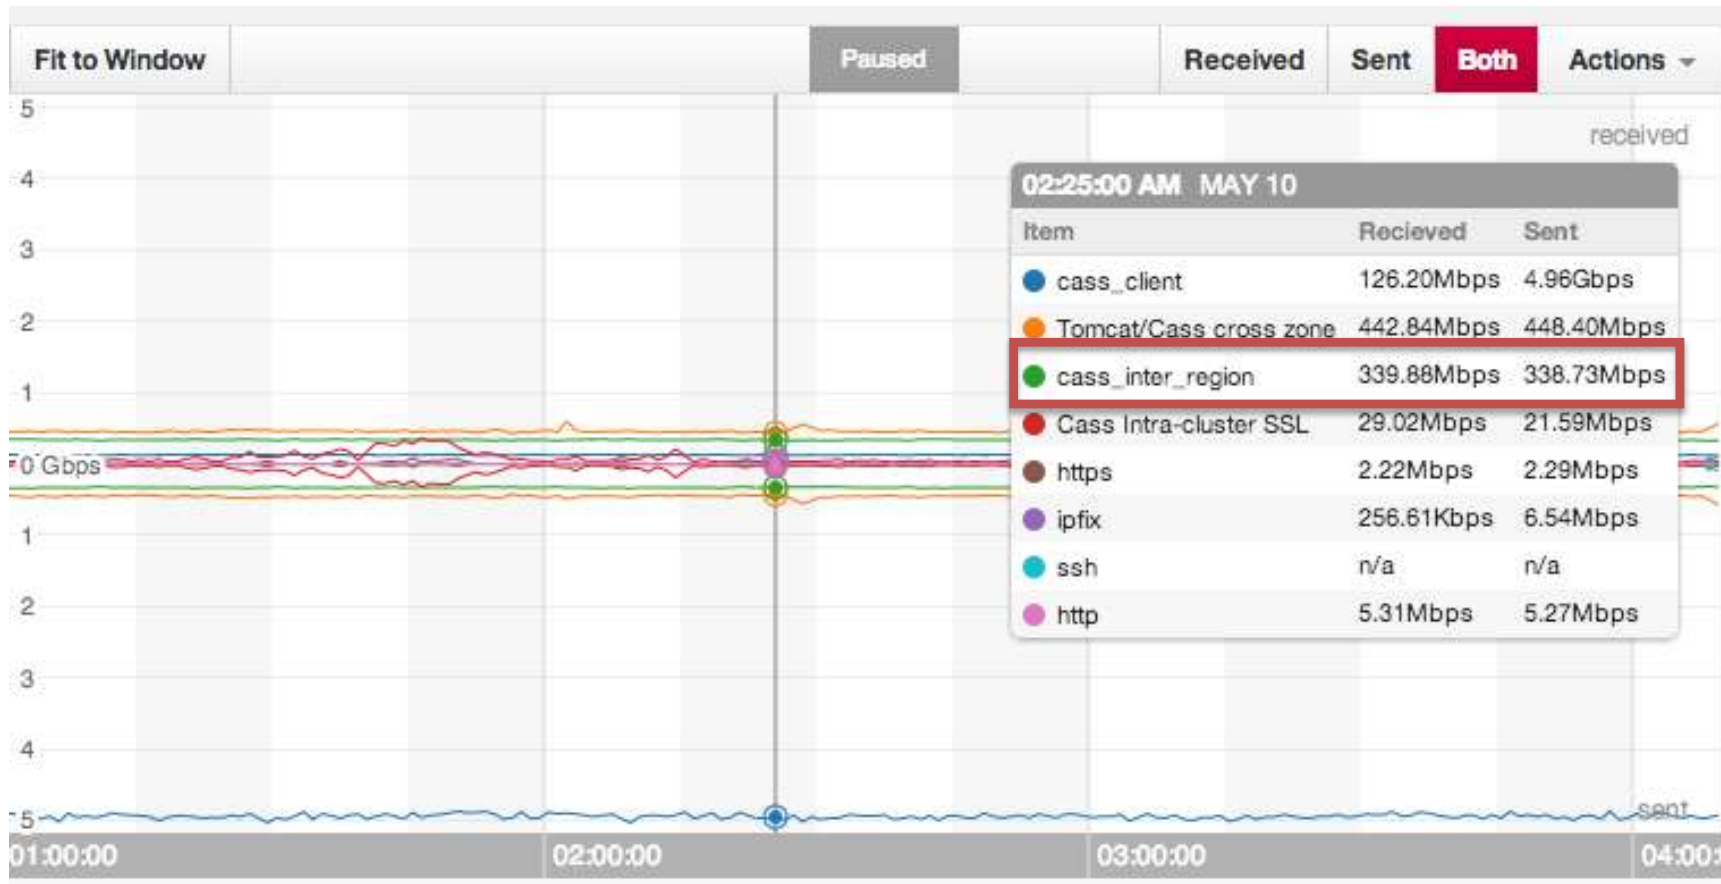
Cassandra bootstrap 9.3 Gbit/s single threaded 48 nodes to 48 nodes

Thanks to boundary.com for these network analysis plots



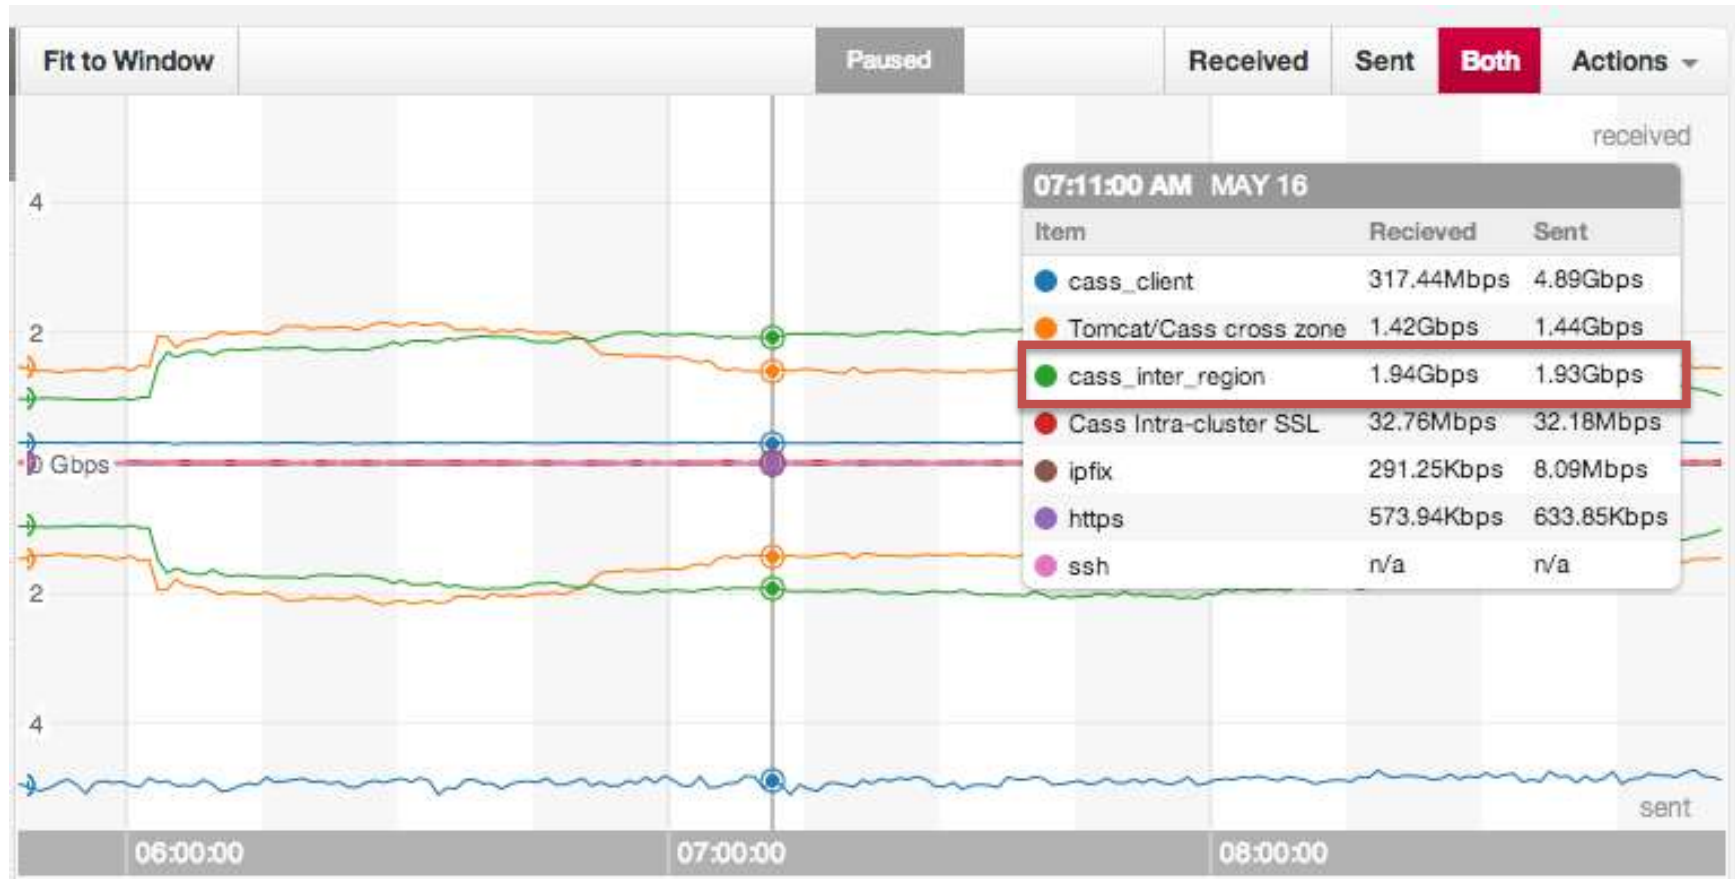
# Inter Region Traffic Test

Verified at desired capacity, no problems, 339 MB/s, 83ms latency

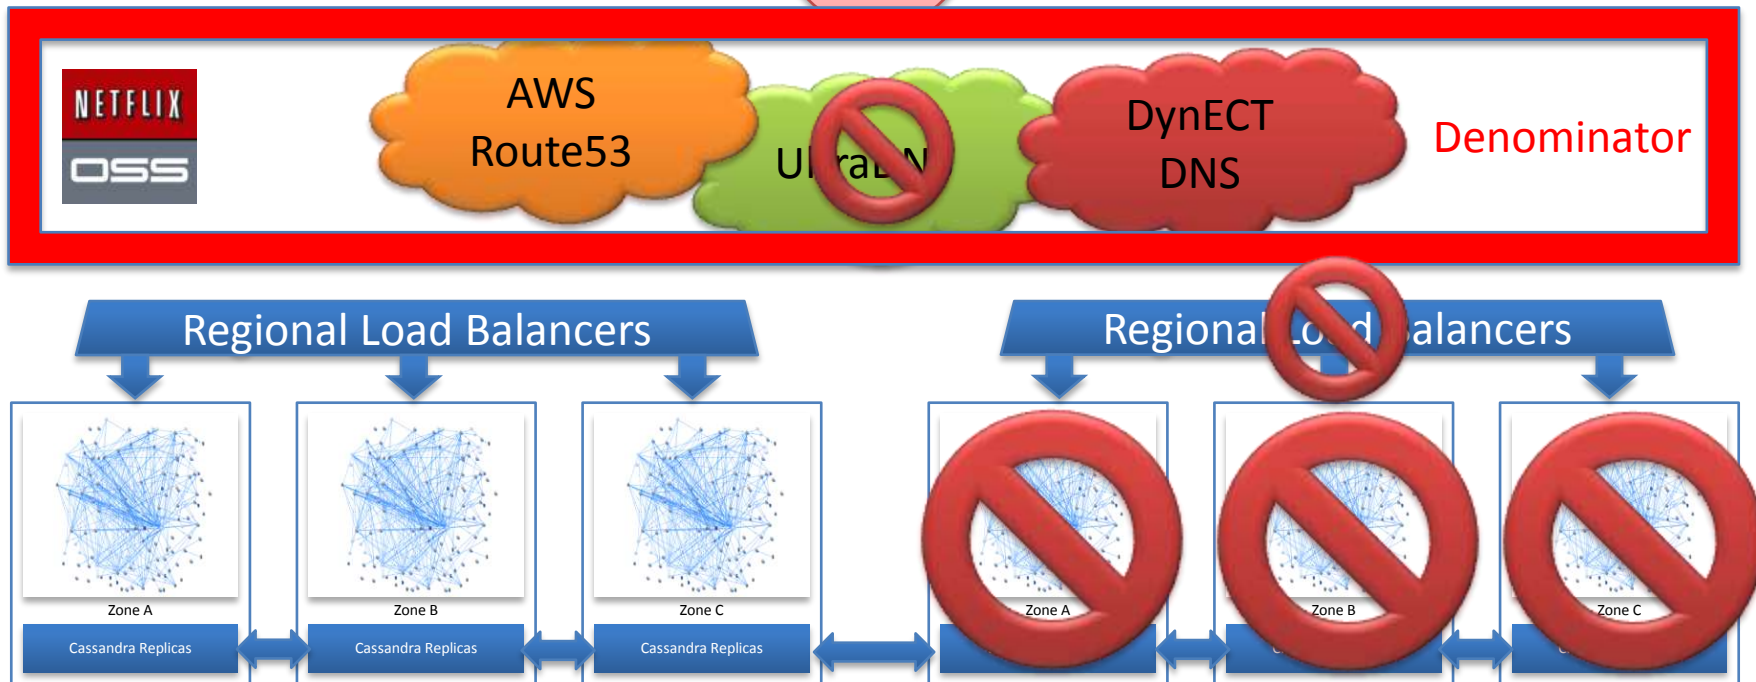


# Ramp Up Load Until It Breaks!

Unmodified tuning, dropping client data at 1.93GB/s inter region traffic  
Spare CPU, IOPS, Network, just need some Cassandra tuning for more



# Managing Multi-Region Availability



Denominator – manage traffic via multiple DNS providers with Java code  
2013 Timeline - Concept Jan, Code Feb, OSS March, Production use May



# Failure Modes and Effects

Failure Mode	Probability	Current Mitigation Plan
Application Failure	High	Automatic degraded response
AWS Region Failure	Low	Active-Active multi-region deployment
AWS Zone Failure	Medium	Continue to run on 2 out of 3 zones
Datacenter Failure	Medium	Migrate more functions to cloud
Data store failure	Low	Restore from S3 backups
S3 failure	Low	Restore from remote archive

Until we got really good at mitigating high and medium probability failures, the ROI for mitigating regional failures didn't make sense. Getting there...

# Application Resilience

Run what you wrote

Rapid detection

Rapid Response

# Chaos Monkey

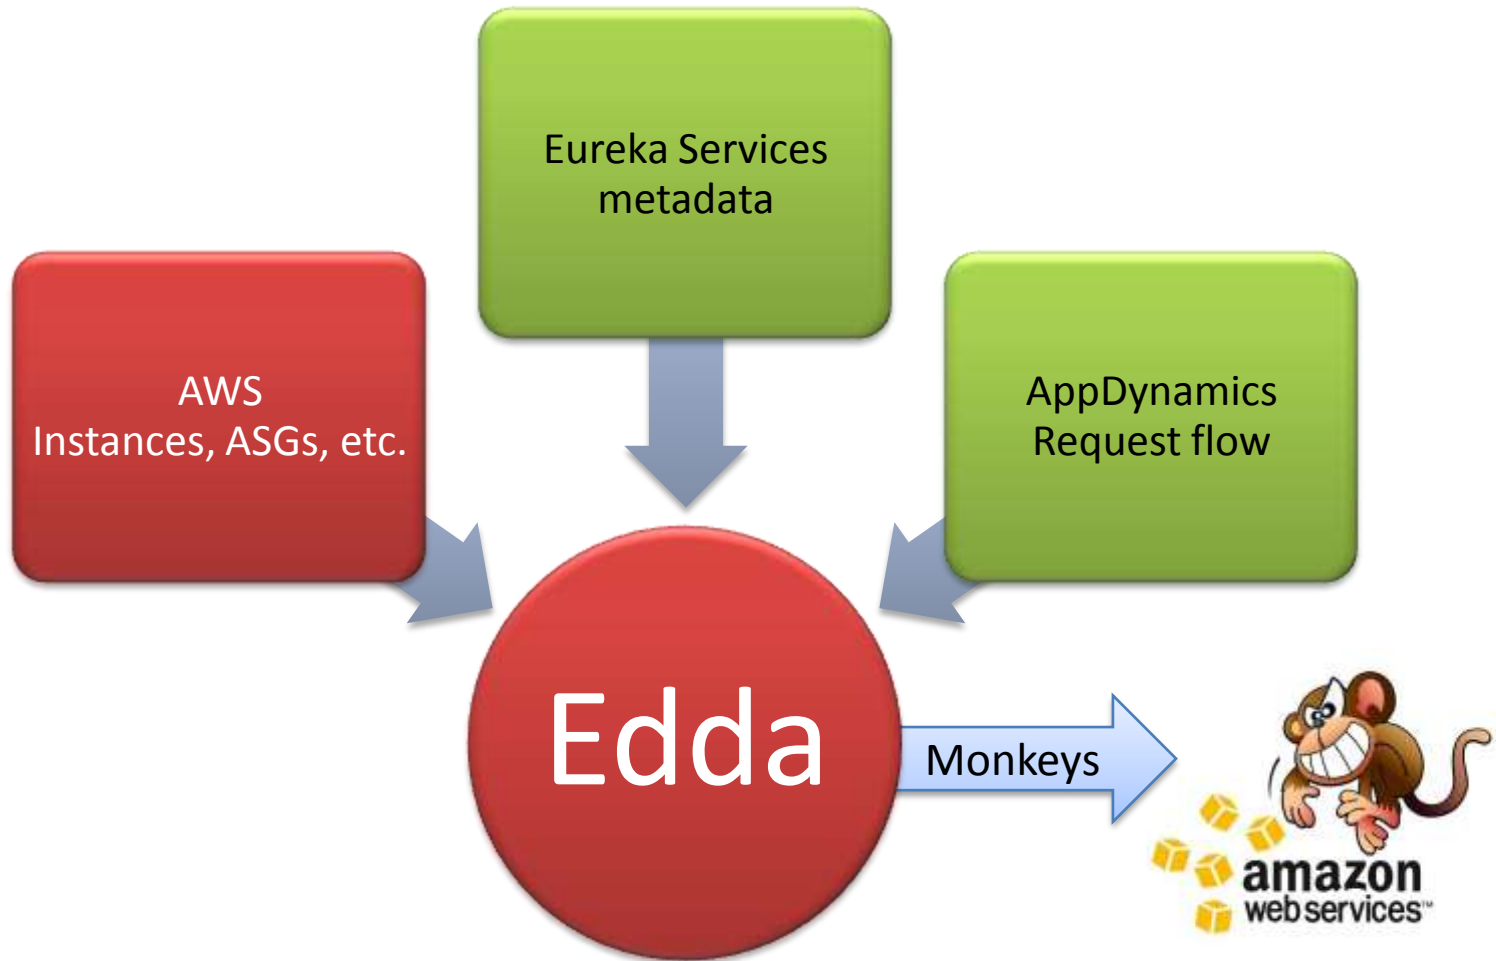
<http://techblog.netflix.com/2012/07/chaos-monkey-released-into-wild.html>

- Computers (Datacenter or AWS) randomly die
  - Fact of life, but too infrequent to test resiliency
- Test to make sure systems are resilient
  - Kill individual instances without customer impact
- Latency Monkey (coming soon)
  - Inject extra latency and error return codes



# Edda – Configuration History

<http://techblog.netflix.com/2012/11/edda-learn-stories-of-your-cloud.html>



# Edda Query Examples

Find any instances that have ever had a specific public IP address

```
$ curl "http://edda/api/v2/view/instances;publicIpAddress=1.2.3.4;_since=0"
["i-0123456789","i-012345678a","i-012345678b"]
```

Show the most recent change to a security group

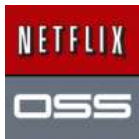
```
$ curl "http://edda/api/v2/aws/securityGroups/sg-0123456789;_diff;_all;_limit=2"
--- /api/v2/aws.securityGroups/sg-0123456789;_pp;_at=1351040779810
+++ /api/v2/aws.securityGroups/sg-0123456789;_pp;_at=1351044093504
@@ -1,33 +1,33 @@
{
...
  "ipRanges" : [
    "10.10.1.1/32",
    "10.10.1.2/32",
+   "10.10.1.3/32",
-   "10.10.1.4/32"
...
  ]
}
```

# Cloud Native Big Data

Size the cluster to the data

Size the cluster to the questions

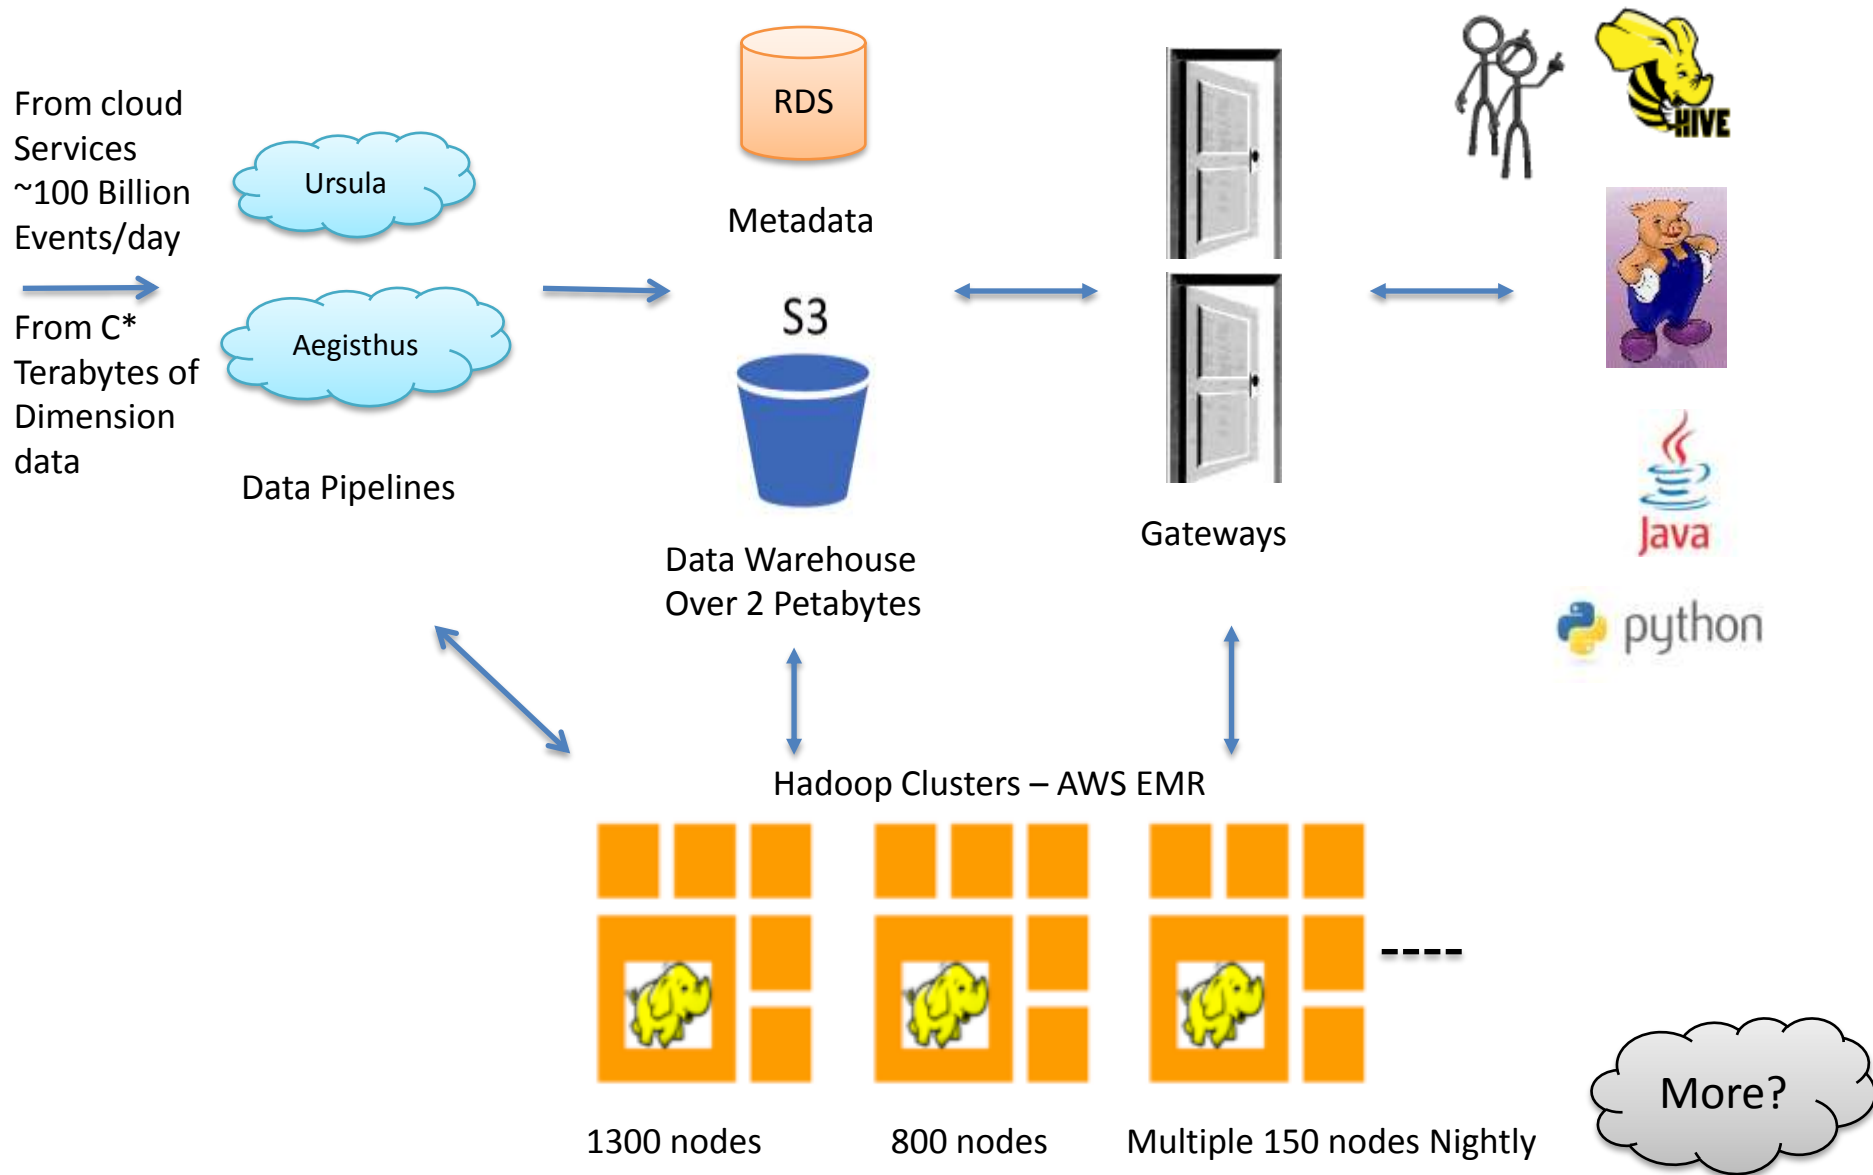
Never wait for space or answers



# Netflix Dataoven

From cloud  
Services  
~100 Billion  
Events/day

From C\*  
Terabytes of  
Dimension  
data



# Cloud Native Development Patterns

Master copies of data are cloud resident

Dynamically provisioned micro-services

Services are distributed and ephemeral



# Datacenter to Cloud Transition Goals

- **Faster**
  - **Lower latency** than the equivalent datacenter web pages and API calls
  - Measured as mean and 99<sup>th</sup> percentile
  - For both first hit (e.g. home page) and in-session hits for the same user
- **Scalable**
  - **Avoid needing any more datacenter capacity** as subscriber count increases
  - No central vertically scaled databases
  - Leverage AWS elastic capacity effectively
- **Available**
  - Substantially **higher robustness and availability** than datacenter services
  - Leverage multiple AWS availability zones
  - No scheduled down time, no central database schema to change
- **Productive**
  - Optimize **agility** of a large development team with automation and tools
  - Leave behind complex tangled datacenter code base (~8 year old architecture)
  - Enforce clean layered interfaces and re-usable components

# Datacenter Anti-Patterns

What do we currently do in the datacenter that prevents us from meeting our goals?

# Rewrite from Scratch

Not everything is cloud specific

Pay down technical debt

Robust patterns

# Netflix Datacenter vs. Cloud Arch

## Anti-Architecture

Central SQL Database

Distributed Key/Value NoSQL

Sticky In-Memory Session

Shared Memcached Session

Chatty Protocols

Latency Tolerant Protocols

Tangled Service Interfaces

Layered Service Interfaces

Instrumented Code

Instrumented Service Patterns

Fat Complex Objects

Lightweight Serializable Objects

Components as Jar Files

Components as Services

More?

# Cloud Security

Fine grain security rather than perimeter

Leveraging AWS Scale to resist DDOS attacks

Automated attack surface monitoring and testing

[http://www.slideshare.net/jason\\_chan/resilience-and-security-scale-lessons-learned](http://www.slideshare.net/jason_chan/resilience-and-security-scale-lessons-learned)

# Security Architecture

- Instance Level Security baked into base AMI
  - Login: ssh only allowed via portal (not between instances)
  - Each app type runs as its own userid app{test|prod}
- AWS Security, Identity and Access Management
  - Each app has its own security group (firewall ports)
  - Fine grain user roles and resource ACLs
- Key Management
  - AWS Keys dynamically provisioned, easy updates
  - High grade app specific key management using HSM



More?

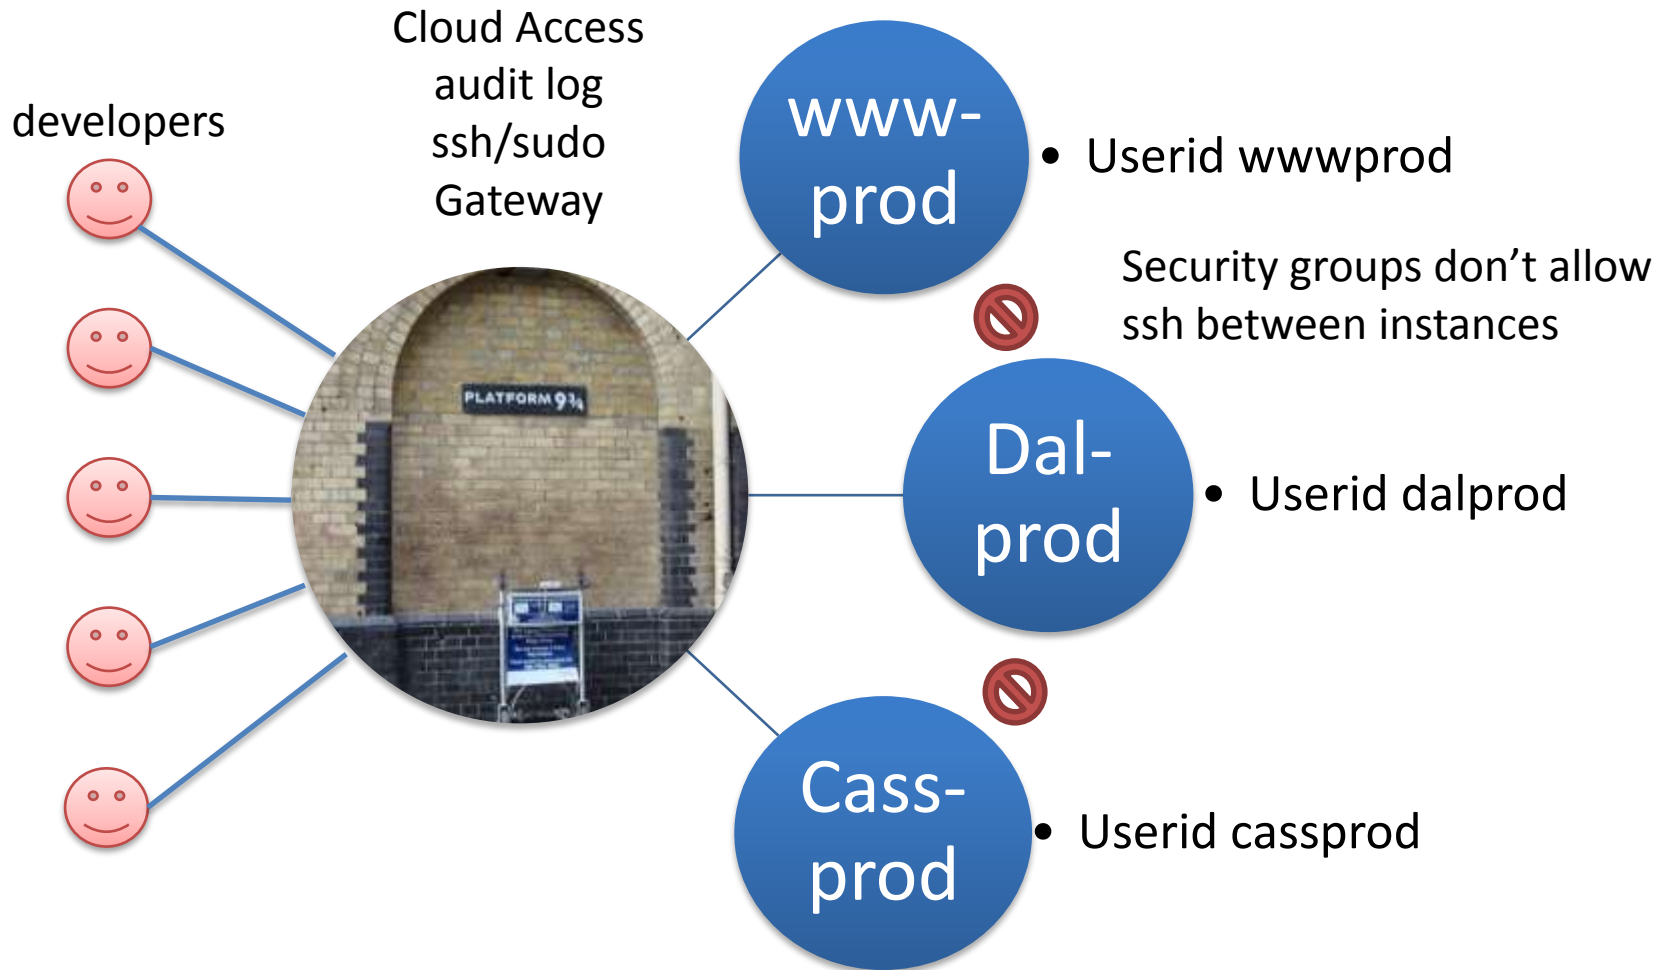
# AWS Accounts

# Accounts Isolate Concerns

- paastest – for development and testing
  - Fully functional deployment of all services
  - Developer tagged “stacks” for separation
- paasprod – for production
  - Autoscale groups only, isolated instances are terminated
  - Alert routing, backups enabled by default
- paasaudit – for sensitive services
  - To support SOX, PCI, etc.
  - Extra access controls, auditing
- paasarchive – for disaster recovery
  - Long term archive of backups
  - Different region, perhaps different vendor



# Cloud Access Control



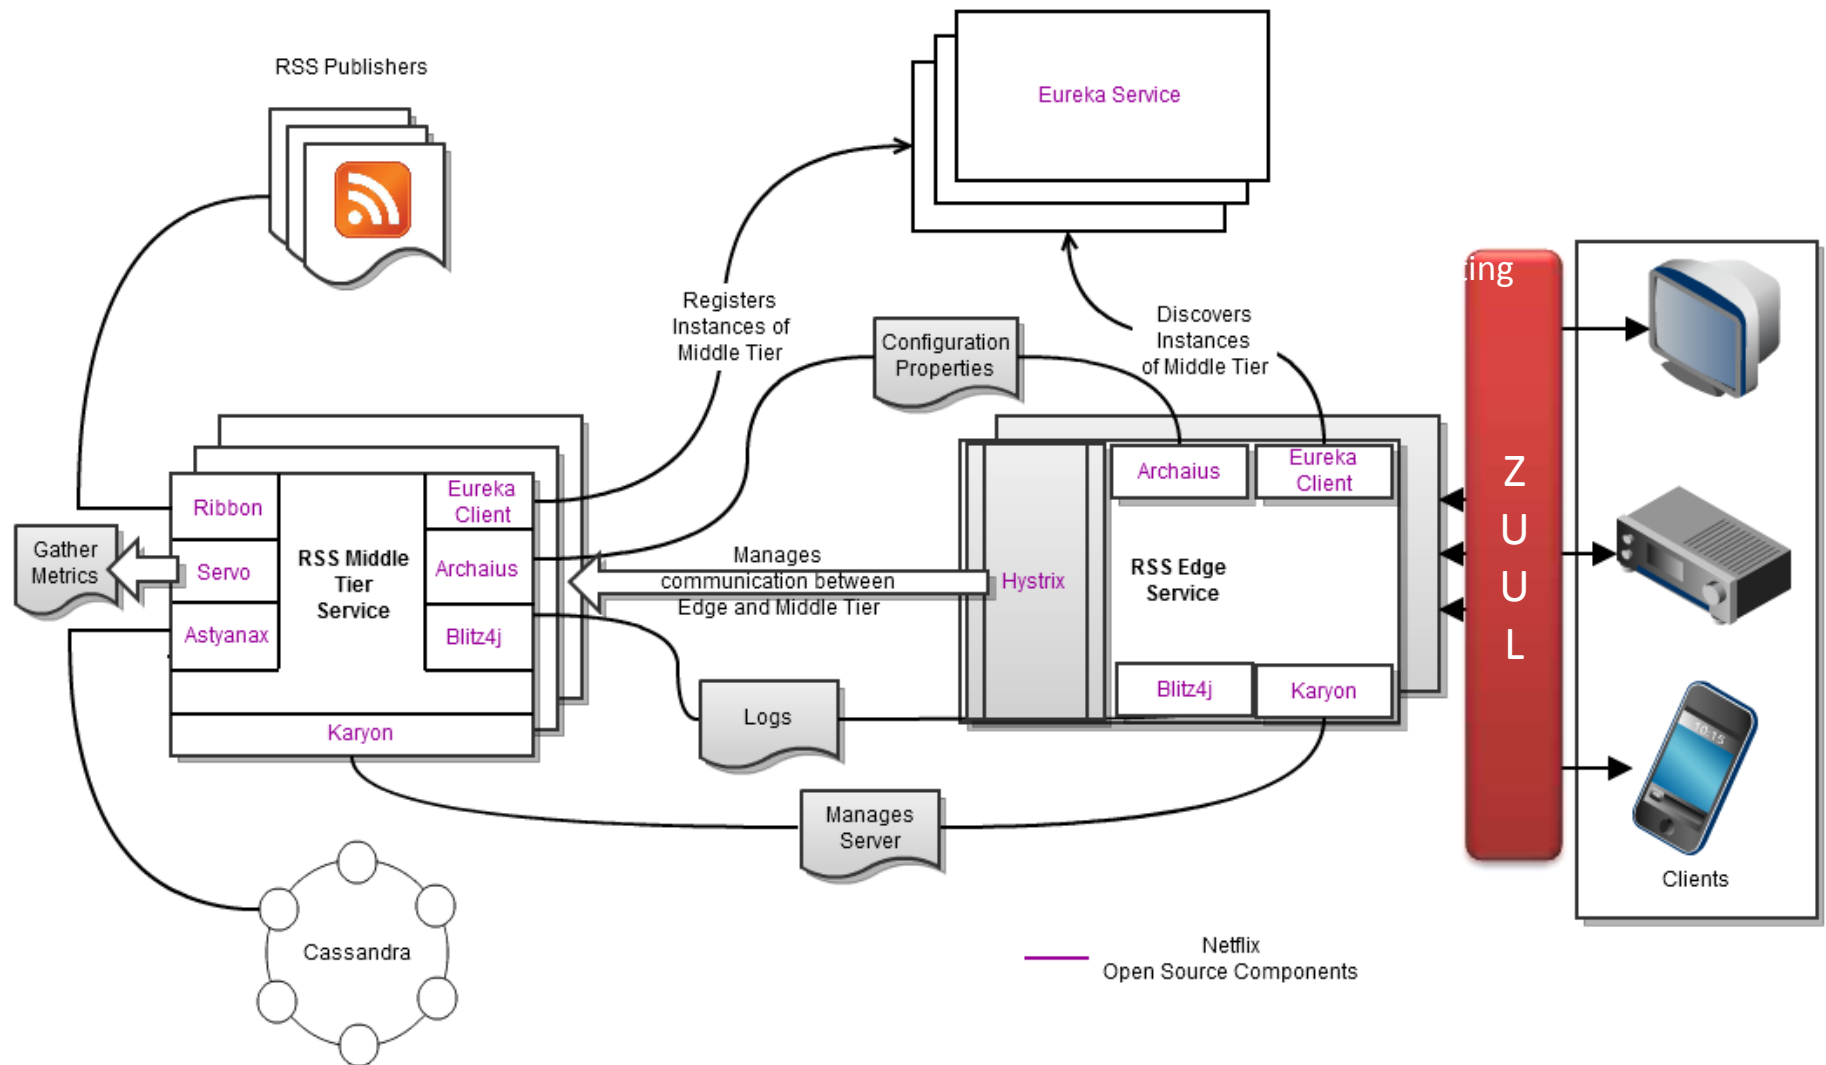


Our perspiration...

A Cloud Native Open Source Platform

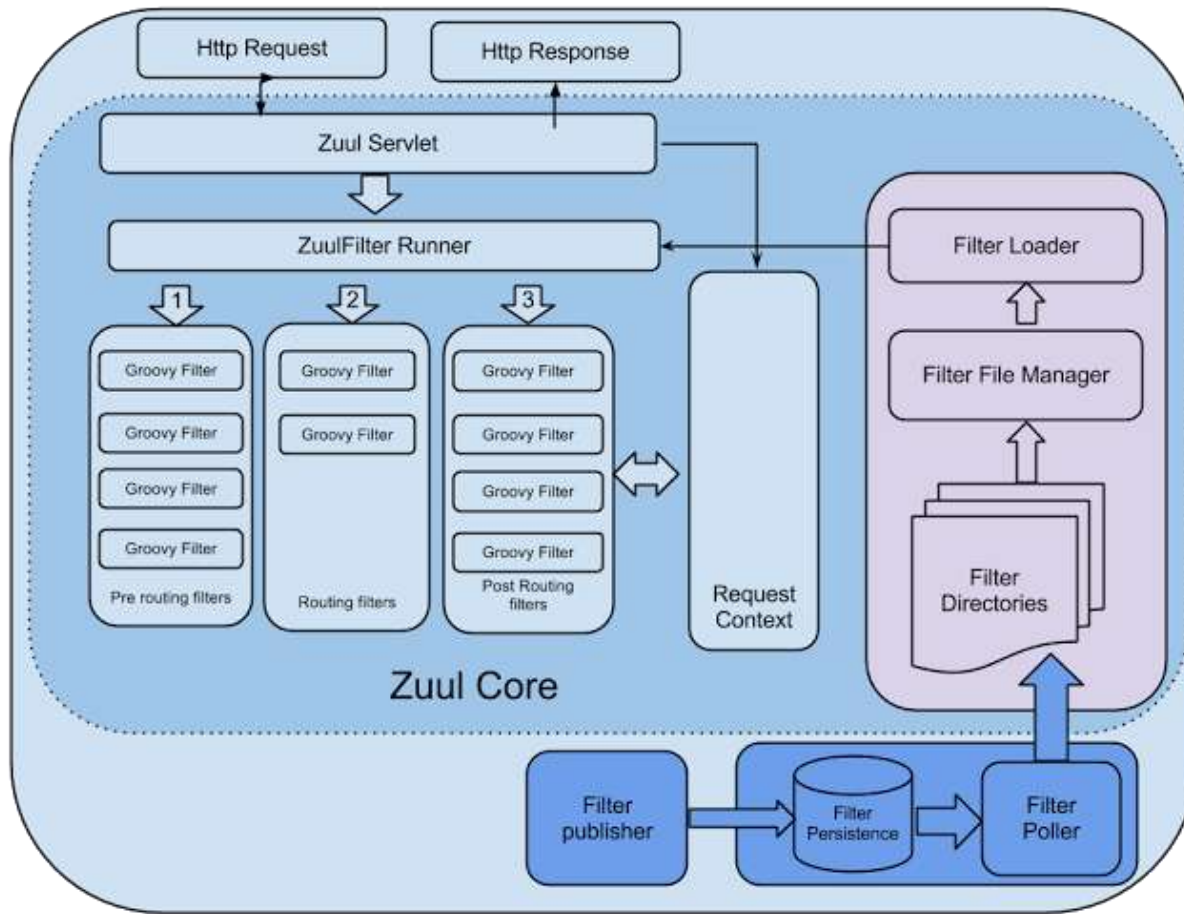
See [netflix.github.com](https://netflix.github.com)

# Example Application – RSS Reader



# Zuul Architecture

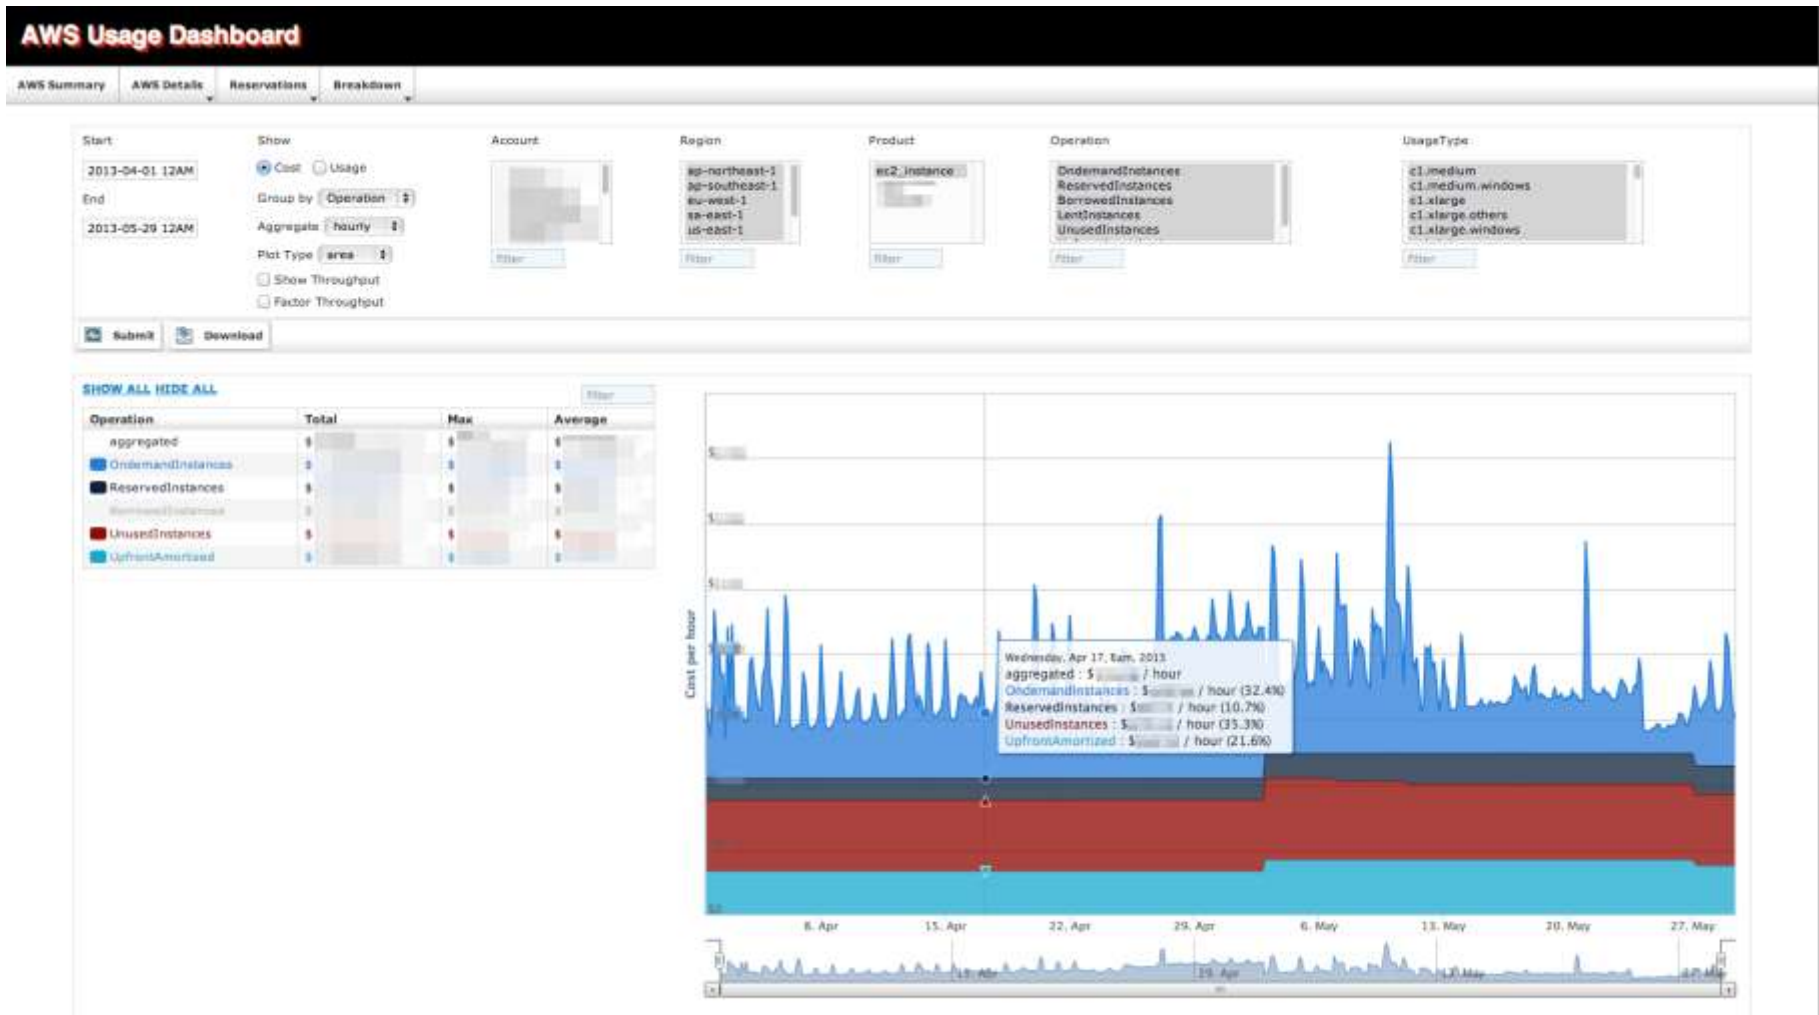
<http://techblog.netflix.com/2013/06/announcing-zuul-edge-service-in-cloud.html>



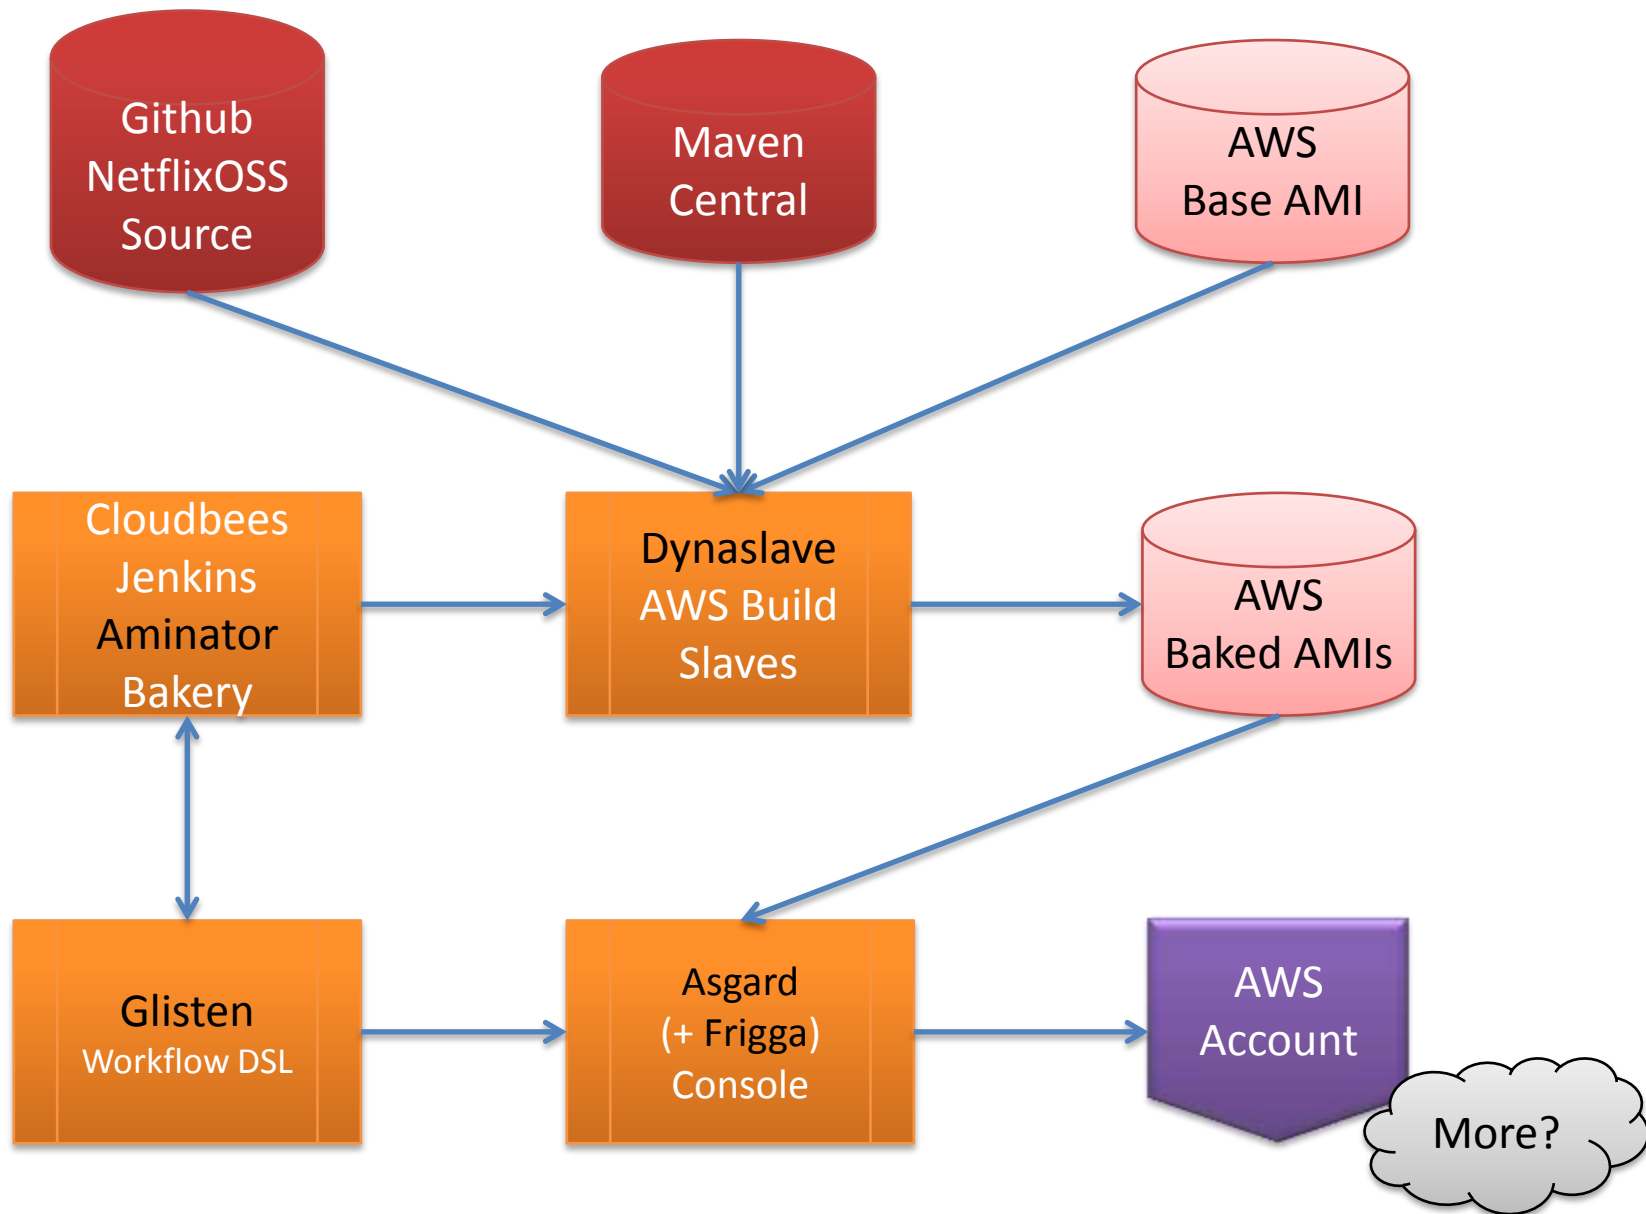


# Ice – AWS Usage Tracking

<http://techblog.netflix.com/2013/06/announcing-ice-cloud-spend-and-usage.html>



# NetflixOSS Continuous Build and Deployment



# NetflixOSS Services Scope

## AWS Account

Asgard Console

Archaius  
Config Service

Cross region Priam C\*

Pytheas  
Dashboards

Atlas  
Monitoring

Genie, Lipstick  
Hadoop Services

Ice – AWS Usage  
Cost Monitoring

## Multiple AWS Regions

Eureka Registry

Exhibitor  
Zookeeper

Edda History

Simian Army

Zuul Traffic Mgr

## 3 AWS Zones

Application Clusters  
Autoscale Groups  
Instances

Priam  
Cassandra  
Persistent Storage

Evcache  
Memcached  
Ephemeral Storage

More?

# NetflixOSS Instance Libraries

## Initialization

- Baked AMI – Tomcat, Apache, your code
- Governor – Guice based dependency injection
- Archaius – dynamic configuration properties client
- Eureka - service registration client

## Service Requests

- Karyon - Base Server for inbound requests
- RxJava – Reactive pattern
- Hystrix/Turbine – dependencies and real-time status
- Ribbon and Feign - REST Clients for outbound calls

## Data Access

- Astyanax – Cassandra client and pattern library
- Evcache – Zone aware Memcached client
- Curator – Zookeeper patterns
- Denominator – DNS routing abstraction

## Logging

- Blitz4j – non-blocking logging
- Servo – metrics export for autoscaling
- Atlas – high volume instrumentation

More?



# NetflixOSS Testing and Automation

## Test Tools

- CassJmeter – Load testing for Cassandra
- Circus Monkey – Test account reservation rebalancing

## Maintenance

- Janitor Monkey – Cleans up unused resources
- Efficiency Monkey
- Doctor Monkey
- Howler Monkey – Complains about AWS limits

## Availability

- Chaos Monkey – Kills Instances
- Chaos Gorilla – Kills Availability Zones
- Chaos Kong – Kills Regions
- Latency Monkey – Latency and error injection

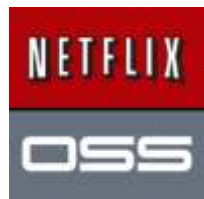
## Security

- Conformity Monkey – architectural pattern warnings
- Security Monkey – security group and S3 bucket permissions

More?

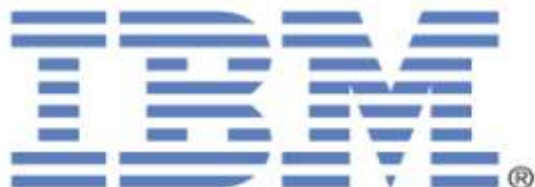
# Vendor Driven Portability

Interest in using NetflixOSS for Enterprise Private Clouds



“It’s done when it runs Asgard”  
Functionally complete  
Demonstrated March  
Released June in V3.3

IBM Example application “Acme Air”  
Based on NetflixOSS running on AWS  
Ported to IBM Softlayer with Rightscale



Vendor and end user interest  
Openstack “Heat” getting there  
Paypal C3 Console based on Asgard

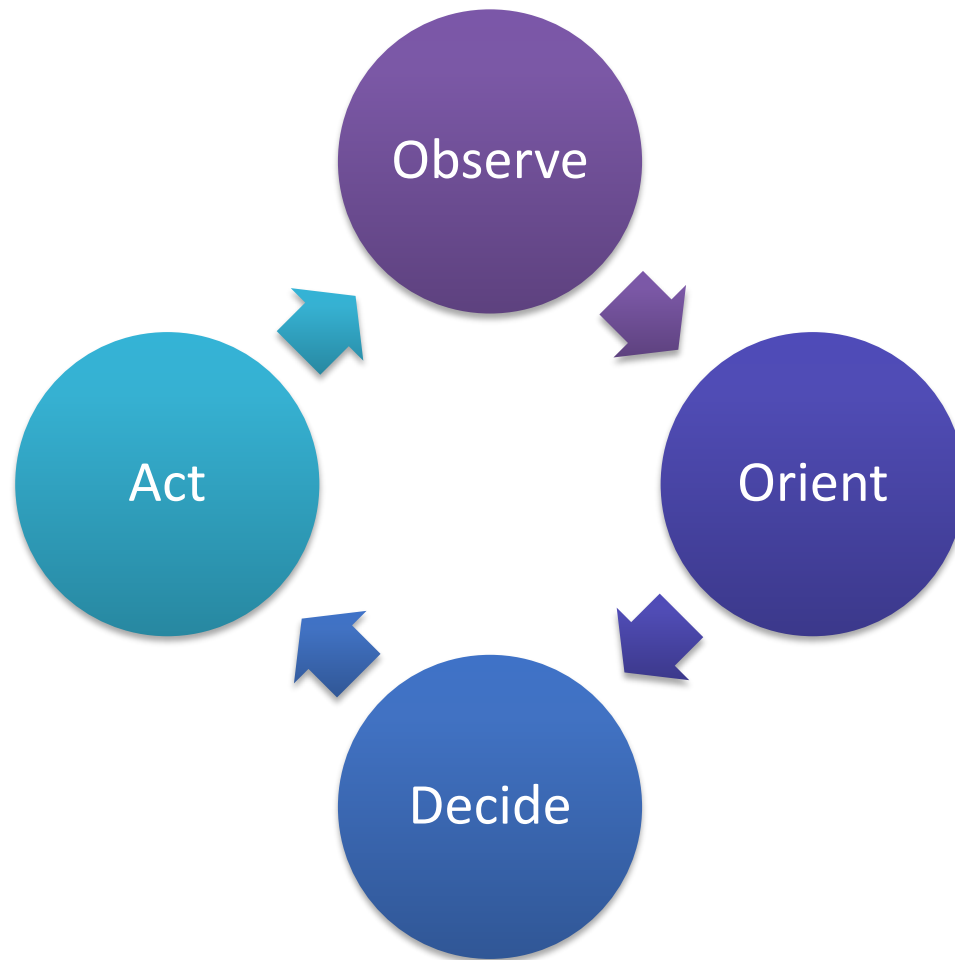


# Cost-Aware Cloud Architectures

Jinesh Varia  
**@jinman**  
Technology Evangelist

**Adrian Cockcroft**  
**@adrianco**  
Director, Architecture





# Experiment Often & Adapt Quickly



- Cost of failure falls dramatically
- Return on (small incremental) Investments is high
- More risk taking, more innovation
- More iteration, faster innovation

« Want to increase innovation?  
**Lower the cost of failure »**

Joi Ito

# Accelerate building a new line of business



# Go Global in Minutes

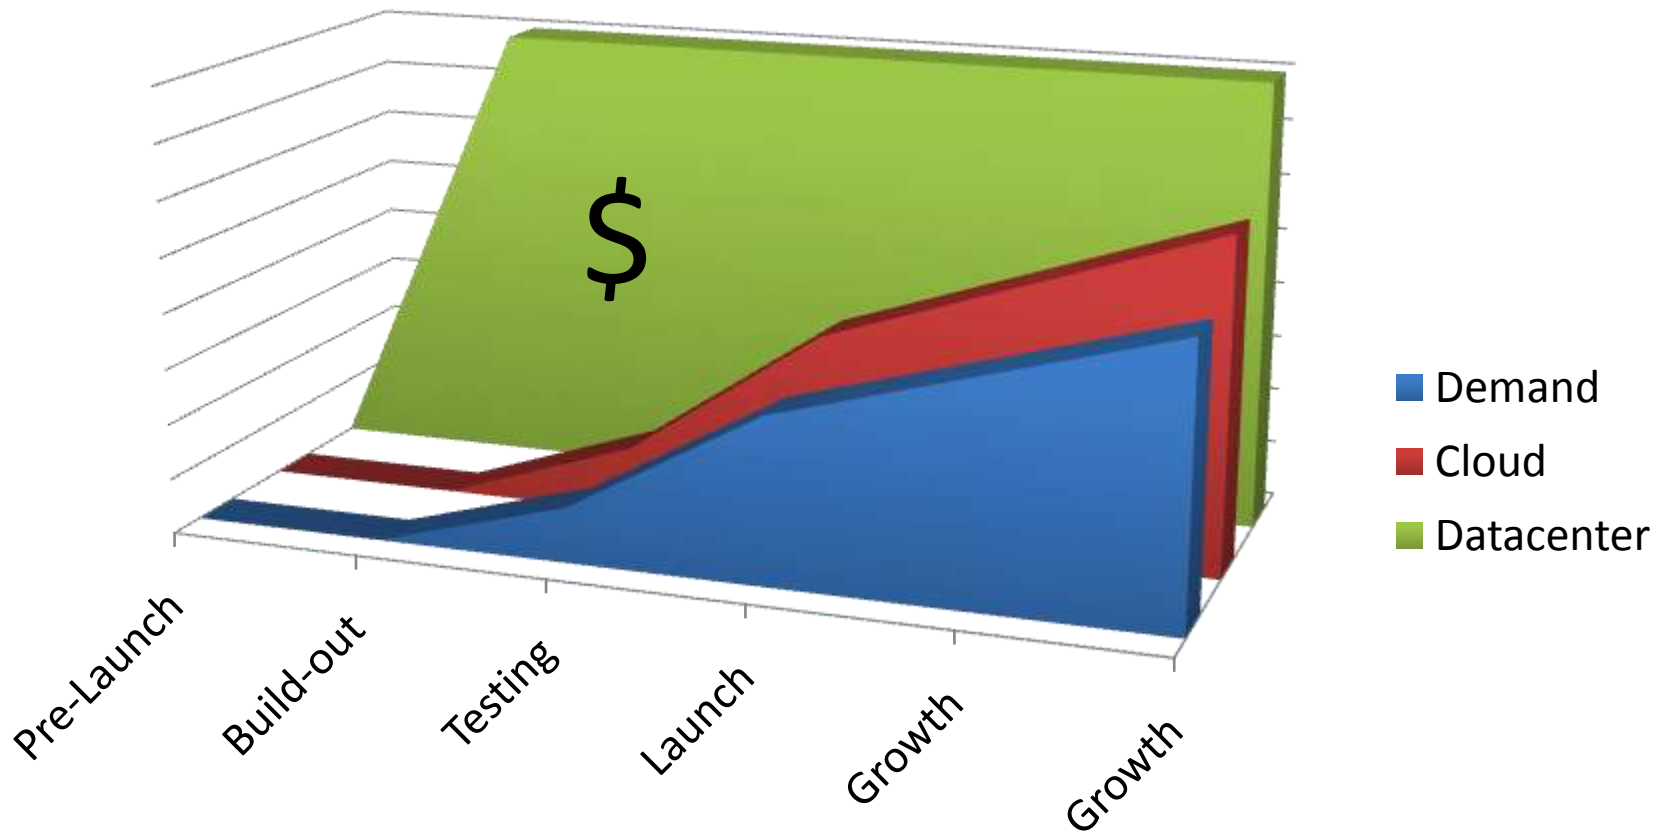




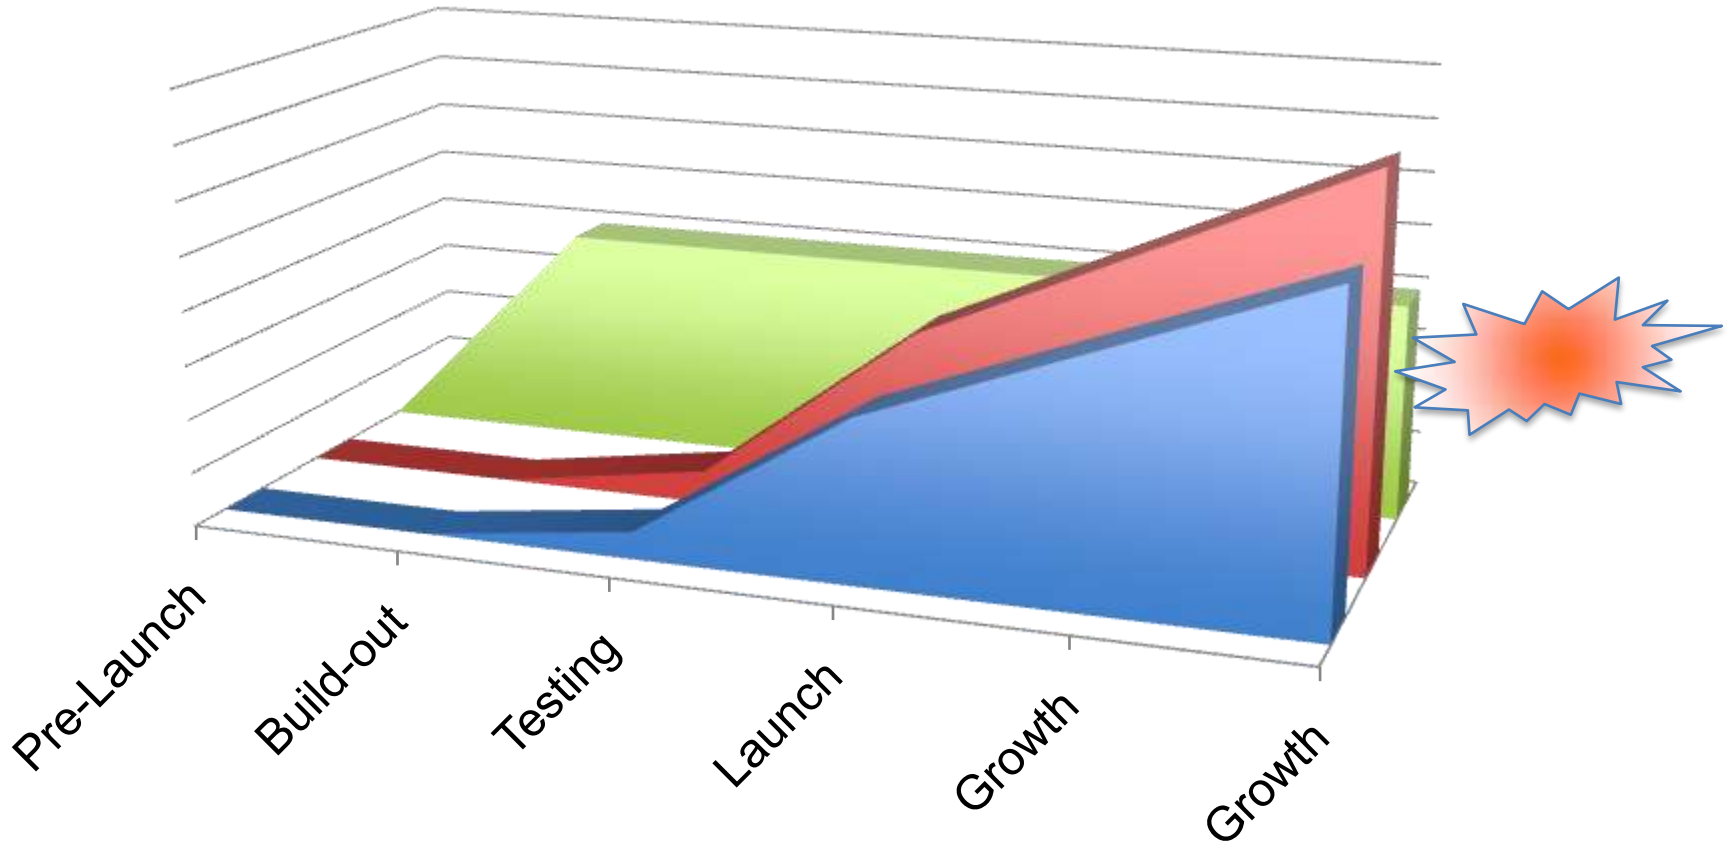
# Netflix Examples

- European Launch using AWS Ireland
  - No employees in Ireland, no provisioning delay, everything worked
  - No need to do detailed capacity planning
  - Over-provisioned on day 1, shrunk to fit after a few days
  - Capacity grows as needed for additional country launches
- Brazilian Proxy Experiment
  - No employees in Brazil, no “meetings with IT”
  - Deployed instances into two zones in AWS Brazil
  - Experimented with network proxy optimization
  - Decided that gain wasn’t enough, shut everything down

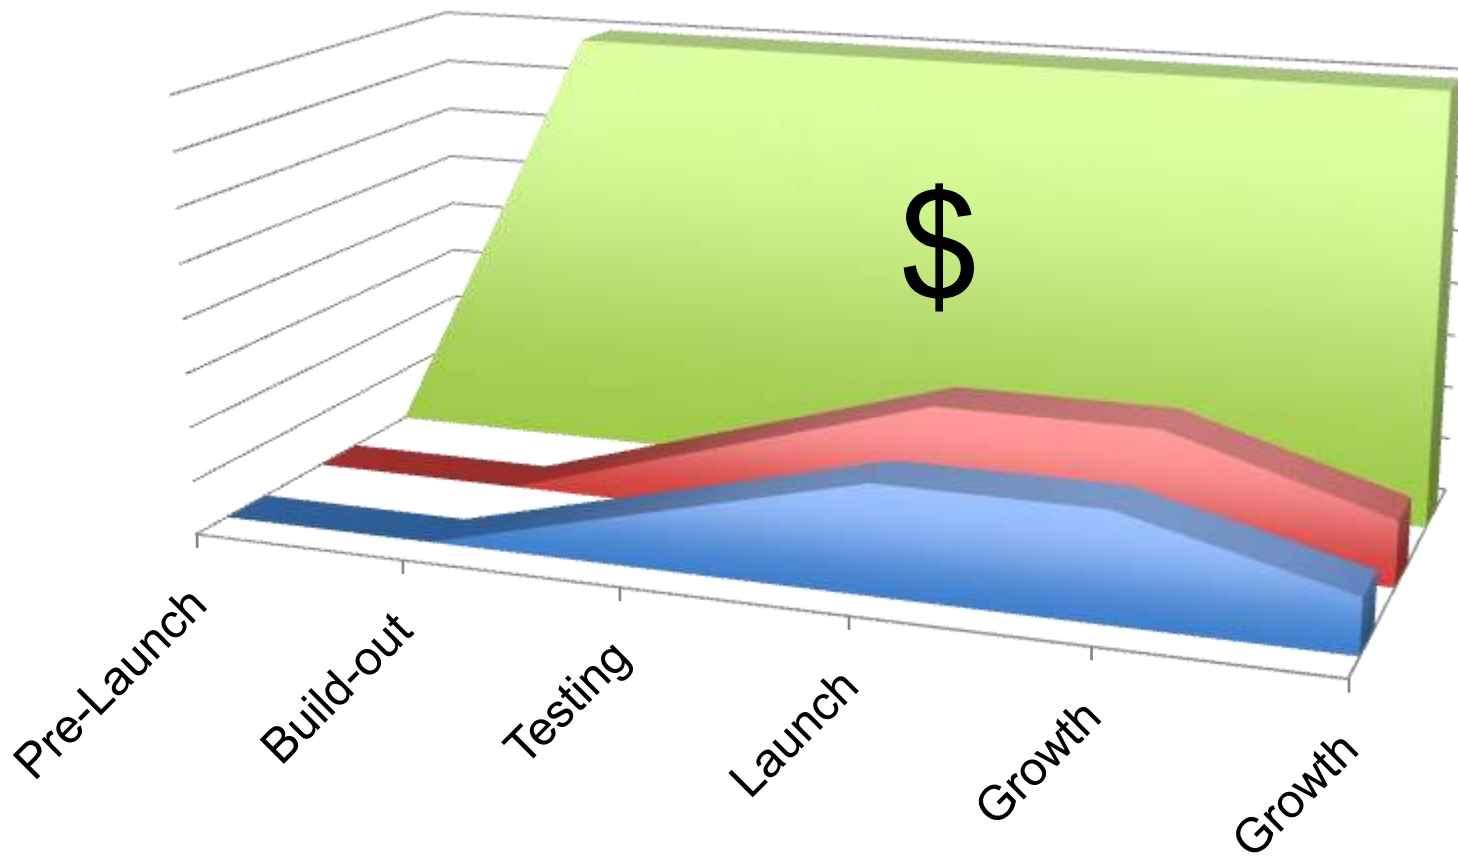
# Product Launch Agility - Rightsized



# Product Launch - Under-estimated



# Product Launch Agility – Over-estimated



**Return on Agility (Agile ROI) = More Revenue**



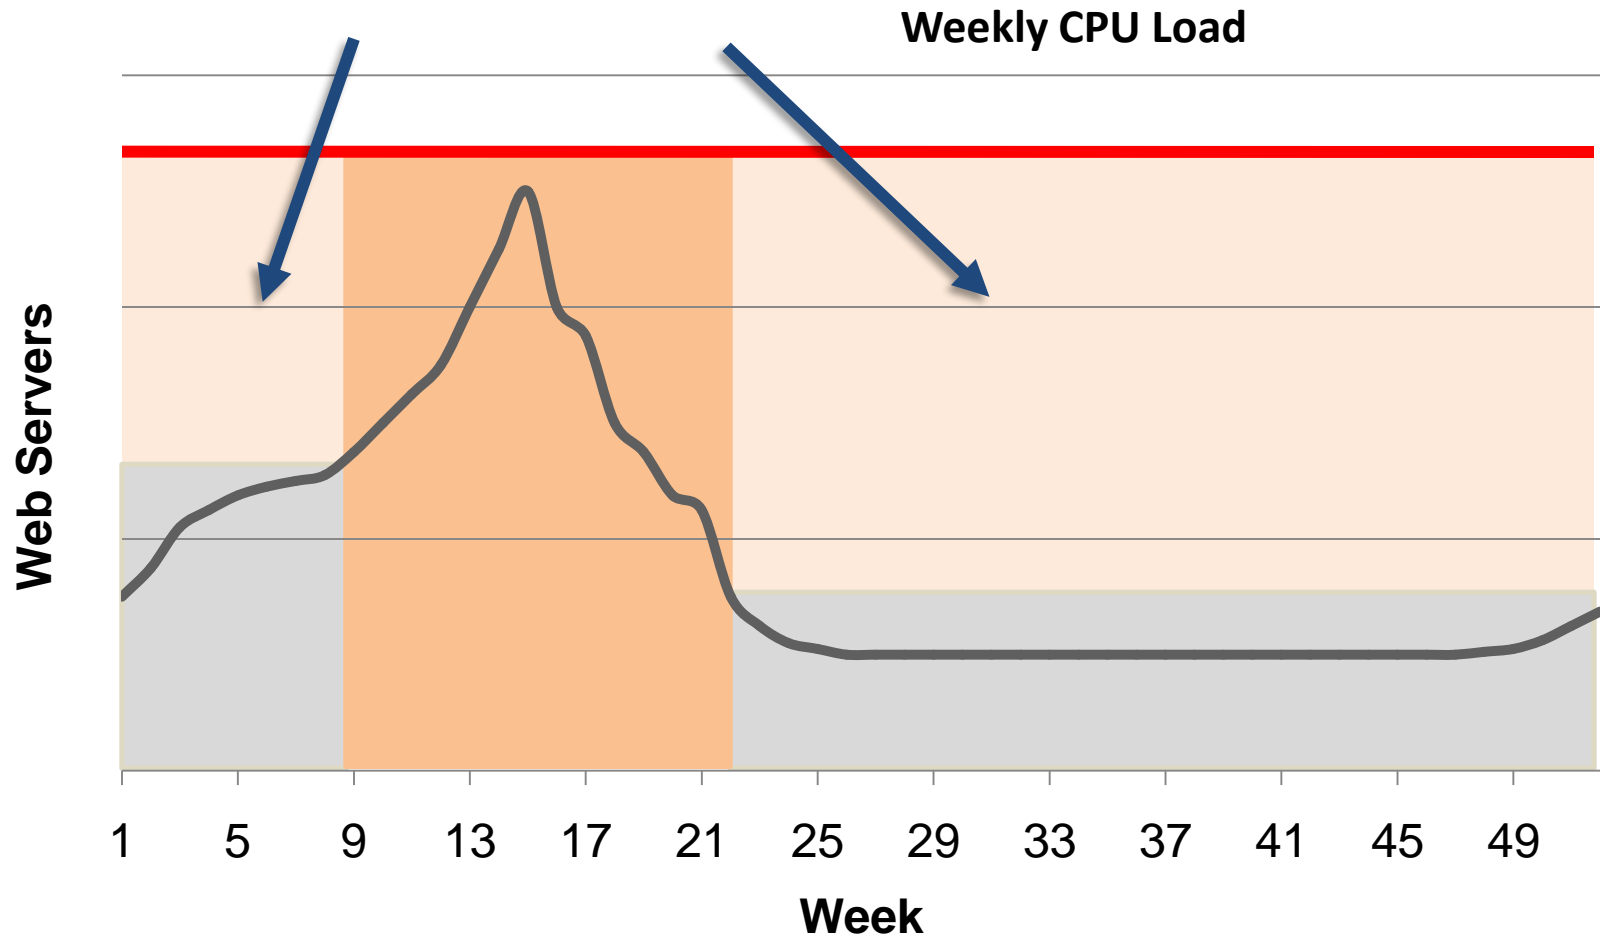
# Key Takeaways on Cost-Aware Architectures....

**#1 Business Agility by Rapid Experimentation = Increased Revenue**



When you turn off your cloud resources, you actually **stop paying for them**

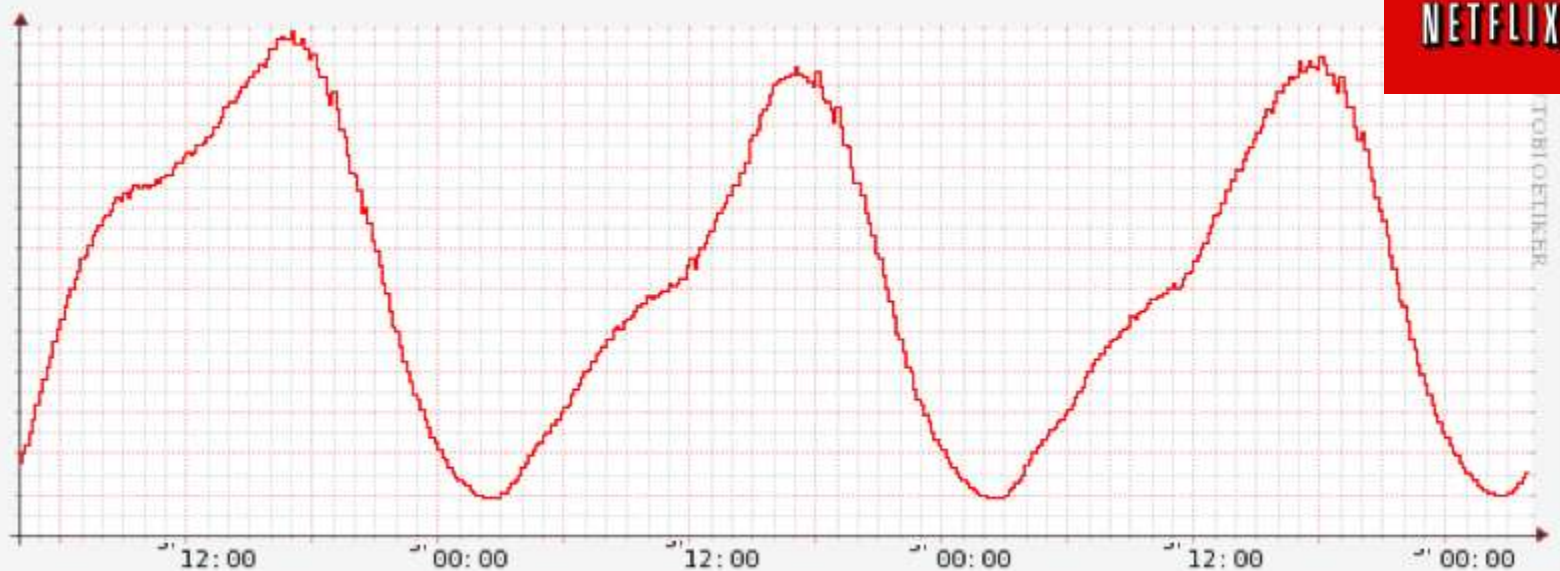
# 50% Savings



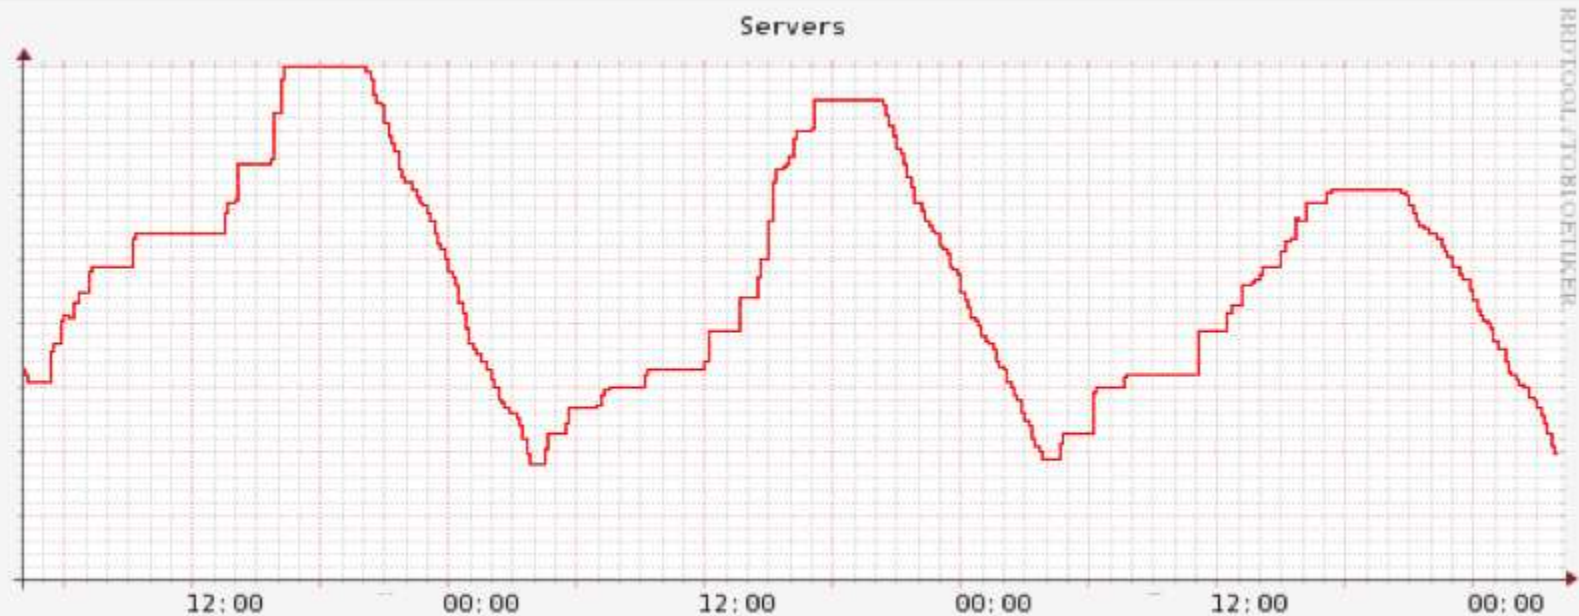
Optimize during a year



Business Throughput



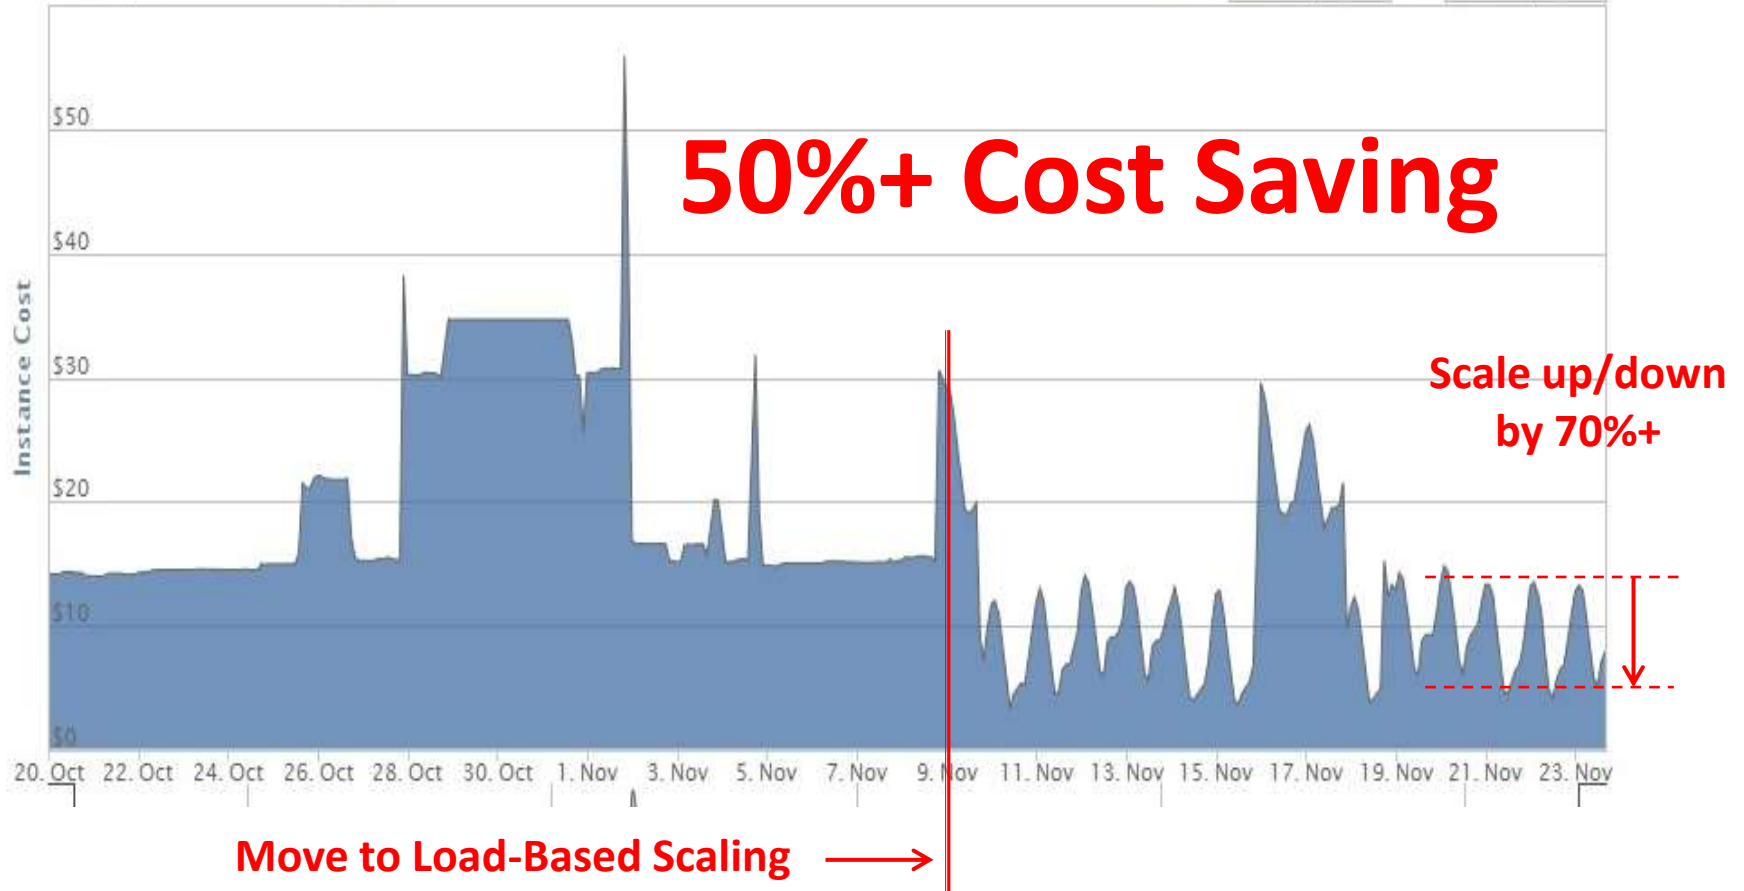
Instances



Cluster - nccp-wii

Zoom 3d 1w 1m 3m All

From: Oct 20, 2011 To: Nov 23, 2011



A public water station is shown outdoors on a concrete pad. It features a grey concrete frame with two horizontal bars. The left bar is attached to a large orange sign that reads "USE ONLY WHAT YOU NEED." in white capital letters, with the "DENVER WATER" logo and name below it. The right bar is attached to a grey concrete post. A wooden plank is mounted on the left bar. The background is a green lawn and a concrete path.

**USE ONLY  
WHAT YOU  
NEED.**

 **DENVER WATER**



Pay as you go

# AWS Support – Trusted Advisor – Your personal cloud assistant

## Trusted Advisor Beta


[Expand All](#)[Download Excel](#)[Refresh All](#)[Contact Support](#)

The AWS Trusted Advisor program monitors AWS infrastructure services, identifies customer configurations, compares them to known best practices, and then notifies customers when opportunities may exist to save money, improve system performance, or close security gaps.

 No issue detected Investigation Recommended Action Recommended


### ▼ Cost Optimizing Checks

#### Unused Elastic IPs

Updated: 2012-06-14 00:00 PDT 

➤ Summary: 0 of 6 Elastic IPs are not in use

#### Underutilized EC2 Instances

Updated: 2012-06-13 22:27 PDT 

➤ Summary: 27 EC2 instances are potentially underutilized



## Other simple optimization tips

- **Don't forget to...**
  - Disassociate unused EIPs
  - Delete unassociated Amazon EBS volumes
  - Delete older Amazon EBS snapshots
  - Leverage Amazon S3 Object Expiration



Janitor Monkey cleans up unused resources

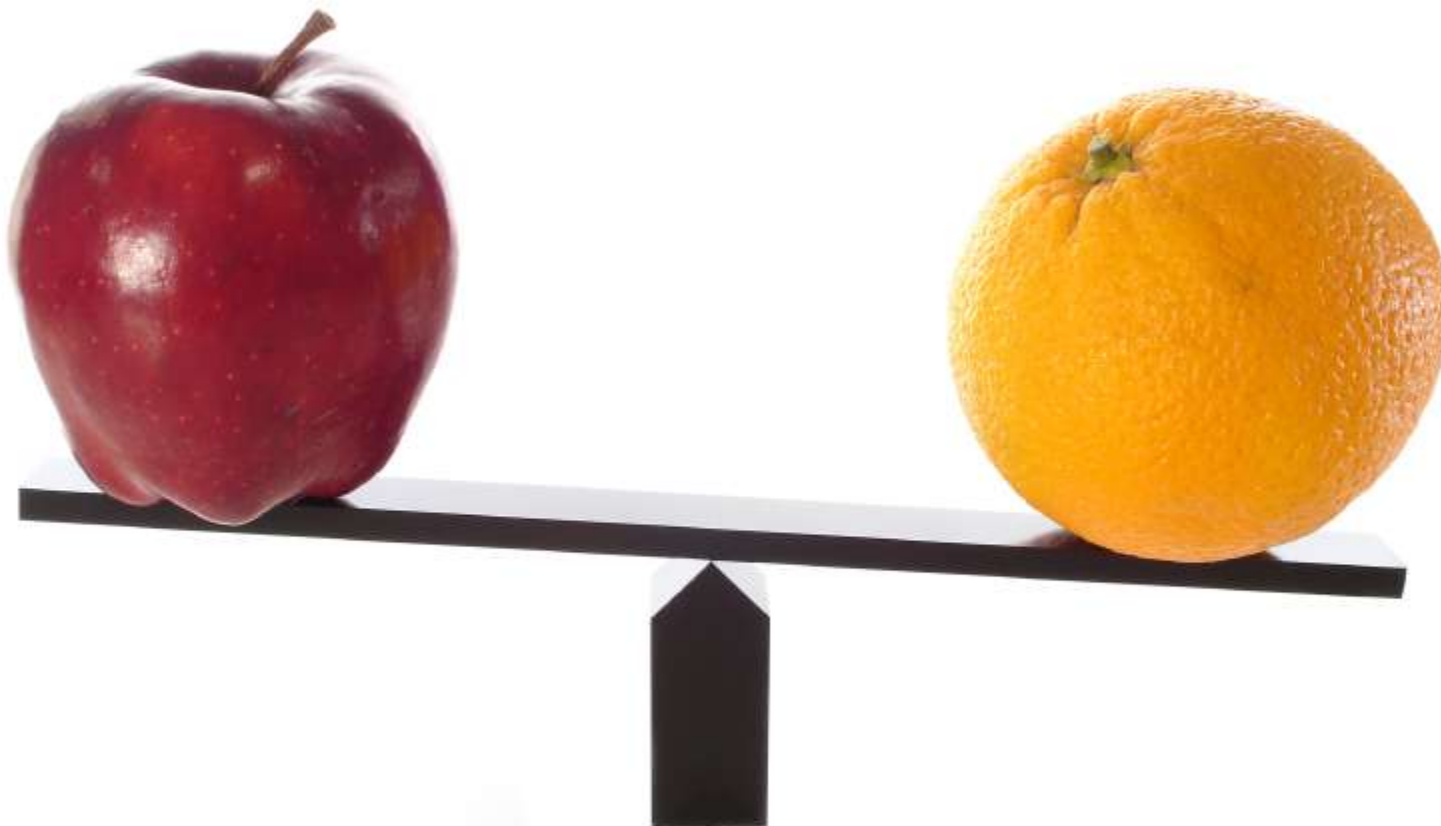


# Building Cost-Aware Cloud Architectures

**#1 Business Agility by Rapid Experimentation = Increased Revenue**

**#2 Business-driven Auto Scaling Architectures = Savings**

## When Comparing TCO...



## When Comparing TCO...

Make sure that  
you are including  
all the cost factors  
into consideration

Place  
Power  
Pipes  
People  
Patterns





# Save more when you reserve

## On-demand Instances

- Pay as you go
- Starts from \$0.02/Hour

## Reserved Instances

- One time low upfront fee + Pay as you go
- \$23 for 1 year term and \$0.01/Hour

1-year and  
3-year terms

Light  
Utilization RI

Medium  
Utilization RI

Heavy  
Utilization RI

# Break-even point

ed  
es  
  
OW  
e + Pay  
  
ear  
  
r

1-year and 3-year terms

Light  
Utilization RI

Medium  
Utilization RI

Heavy  
Utilization RI

**Utilization  
(Uptime)**

10% - 40%  
(>3.5 < 5.5  
months/year)

40% - 75%  
(>5.5 < 7 months/year)

>75%  
(>7 months/year)

**Ideal For**

Disaster Recovery  
(Lowest Upfront)

Standard Reserved  
Capacity

Baseline Servers  
(Lowest Total Cost)

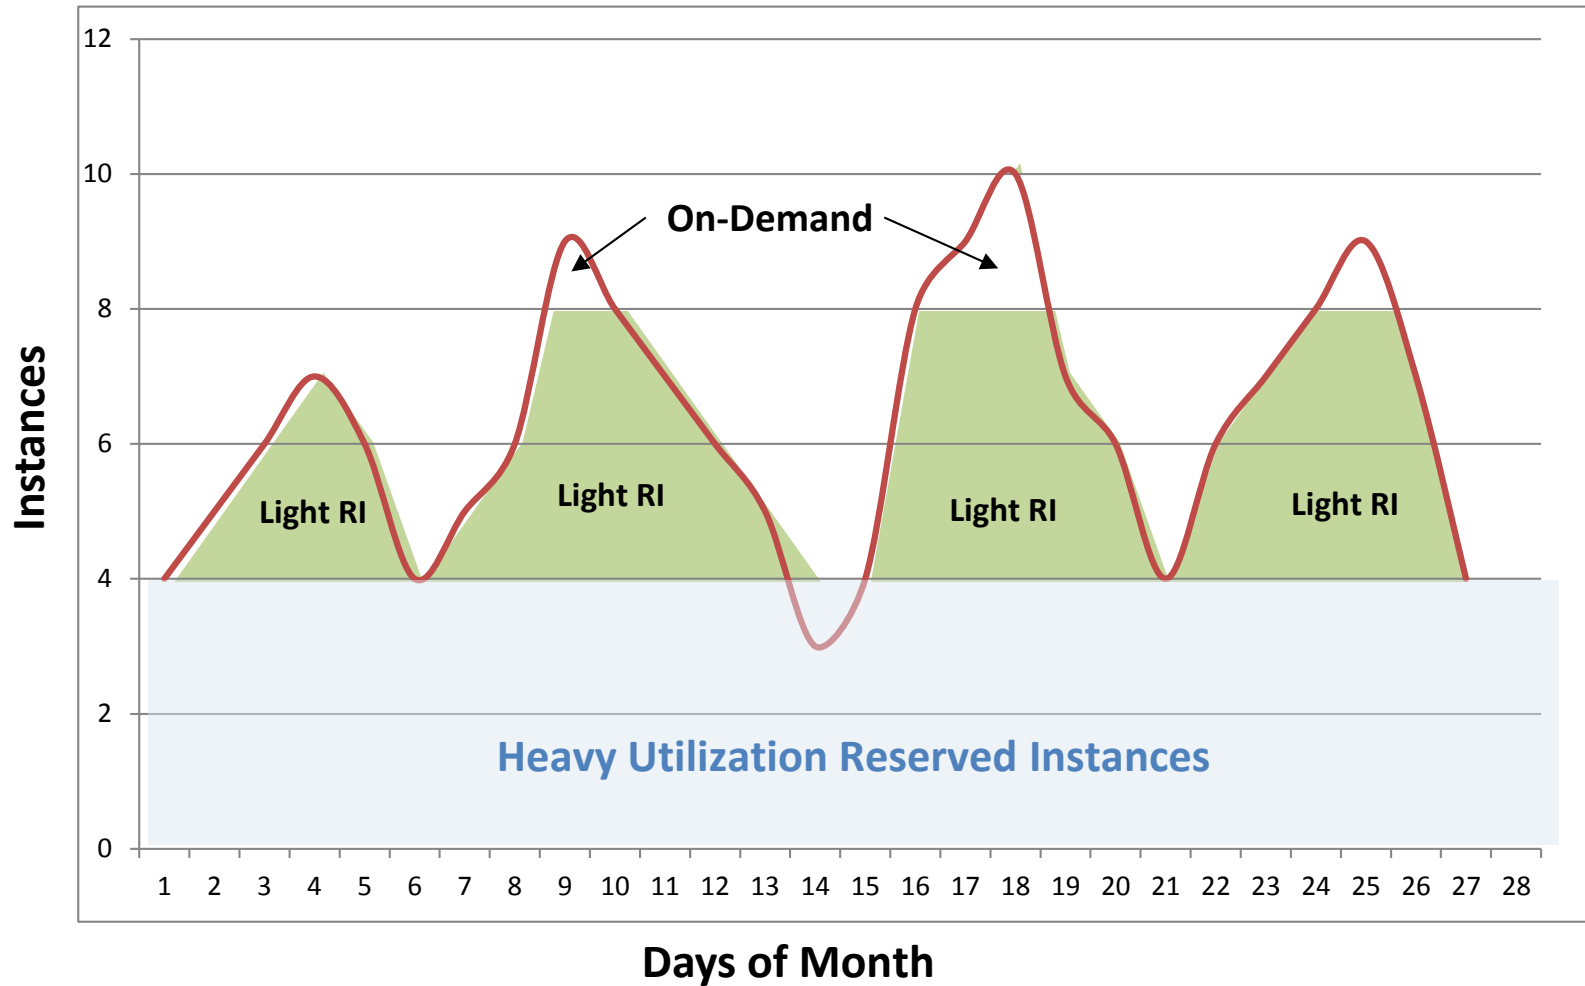
**Savings over  
On-Demand**

**56%**

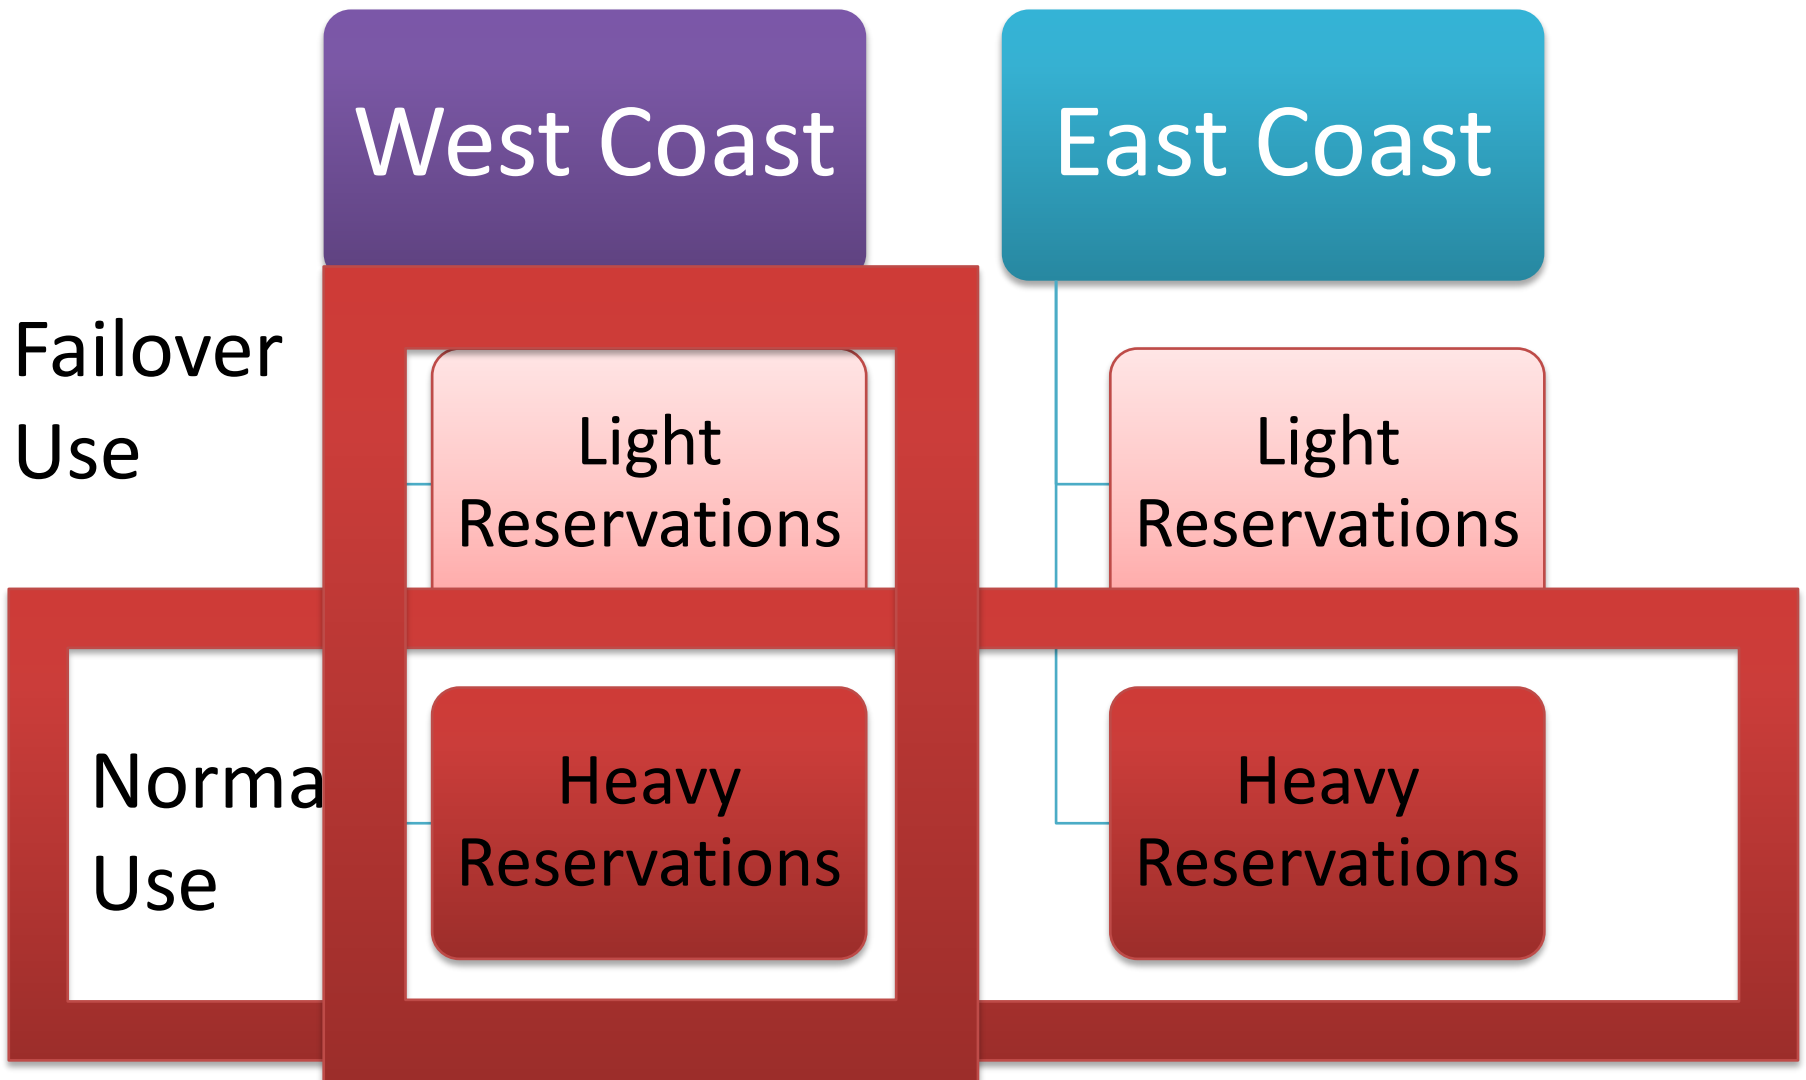
**66%**

**71%**

# Mix and Match Reserved Types and On-Demand



# Netflix Concept for Regional Failover Capacity



# Building Cost-Aware Cloud Architectures

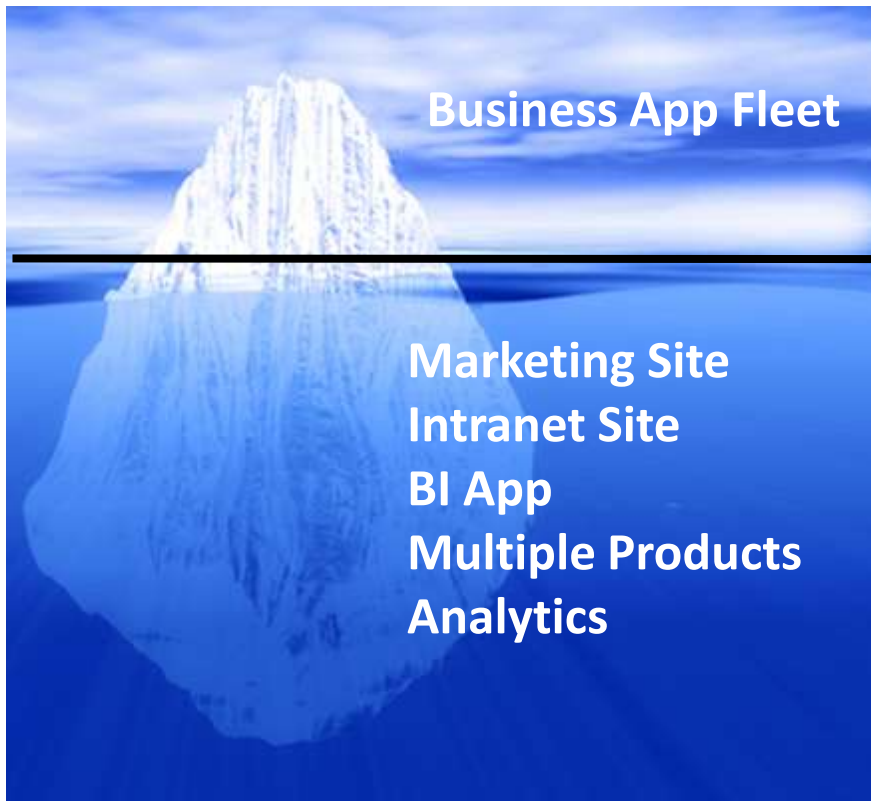
**#1 Business Agility by Rapid Experimentation = Increased Revenue**

**#2 Business-driven Auto Scaling Architectures = Savings**

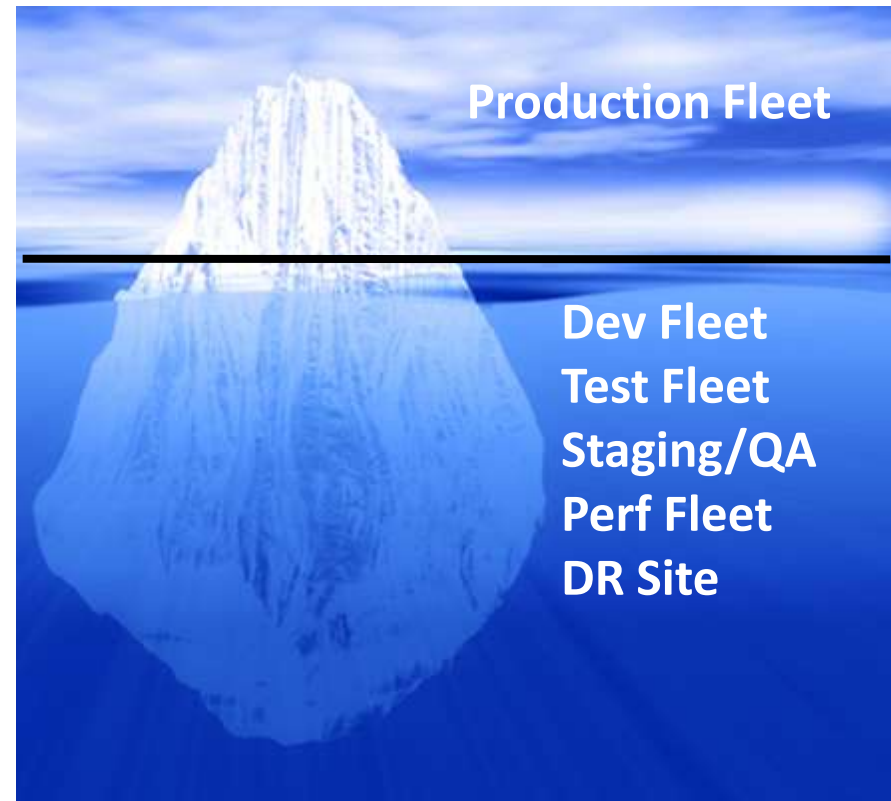
**#3 Mix and Match Reserved Instances with On-Demand = Savings**

# Variety of Applications and Environments

**Every Company has....**



**Every Application has....**

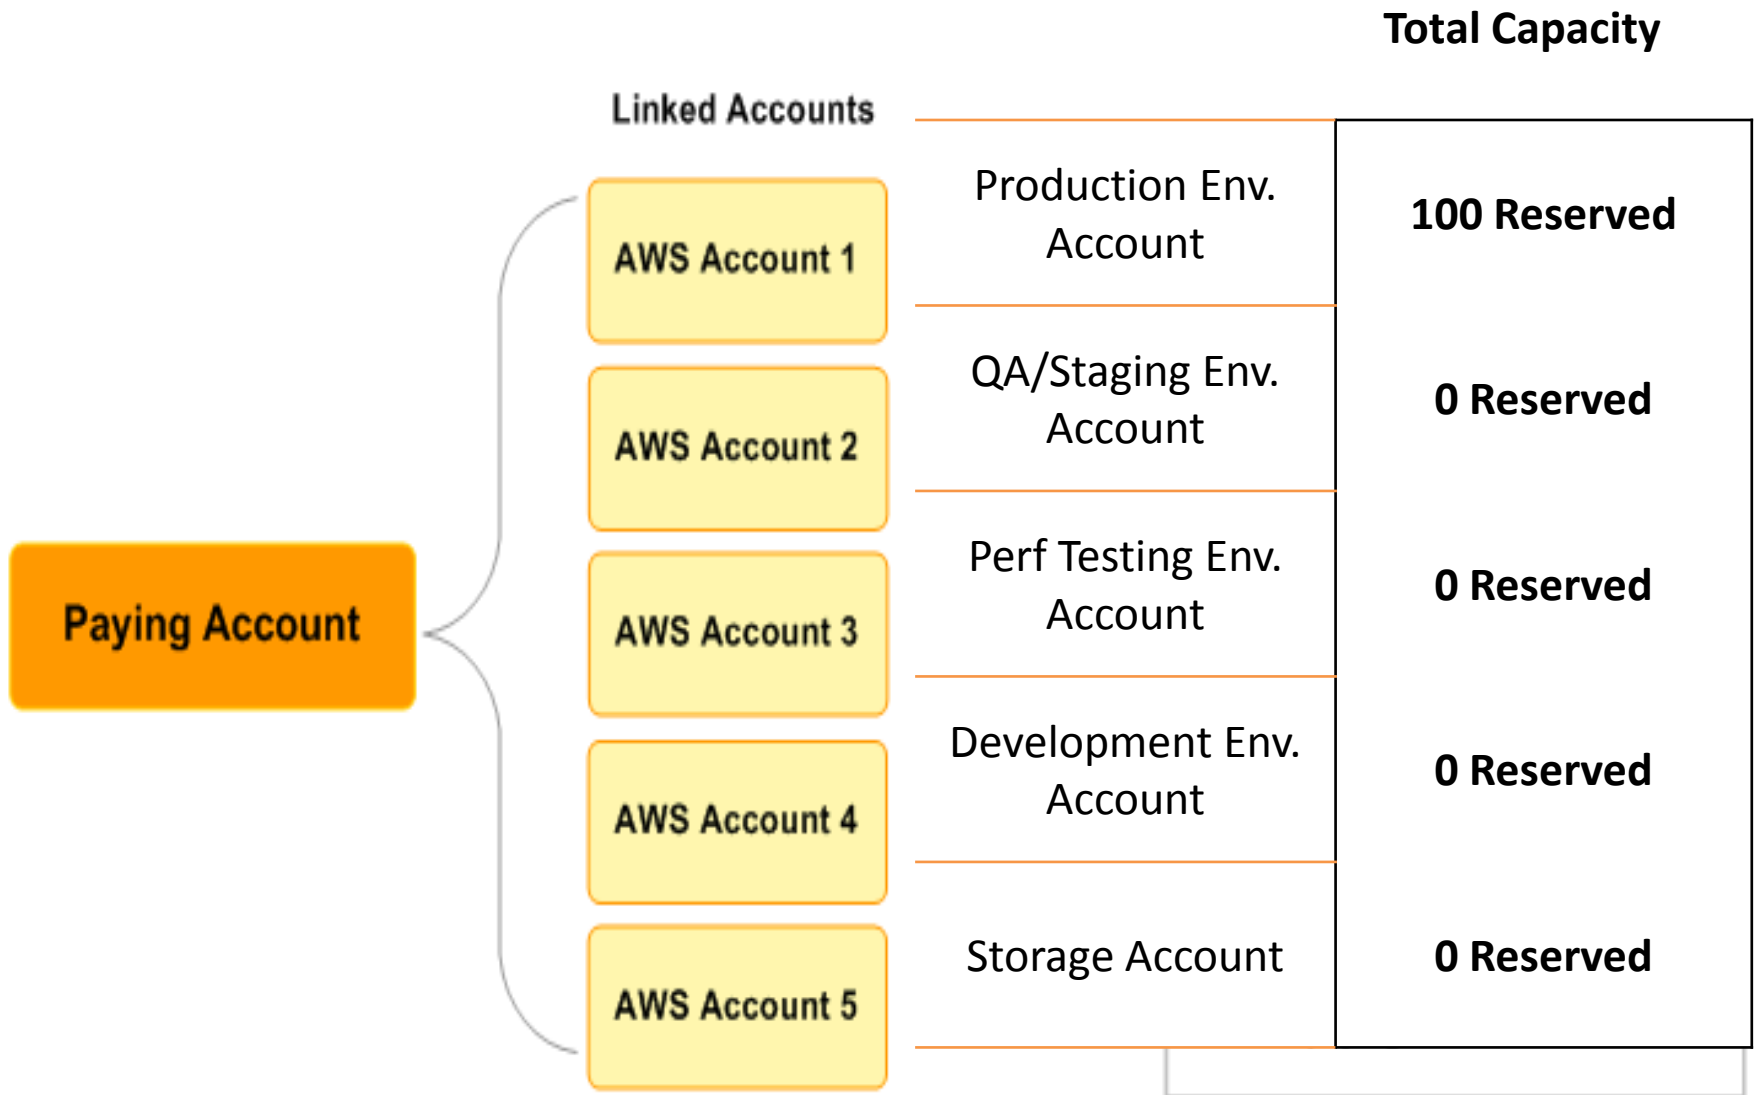


# Consolidated Billing: Single payer for a group of accounts



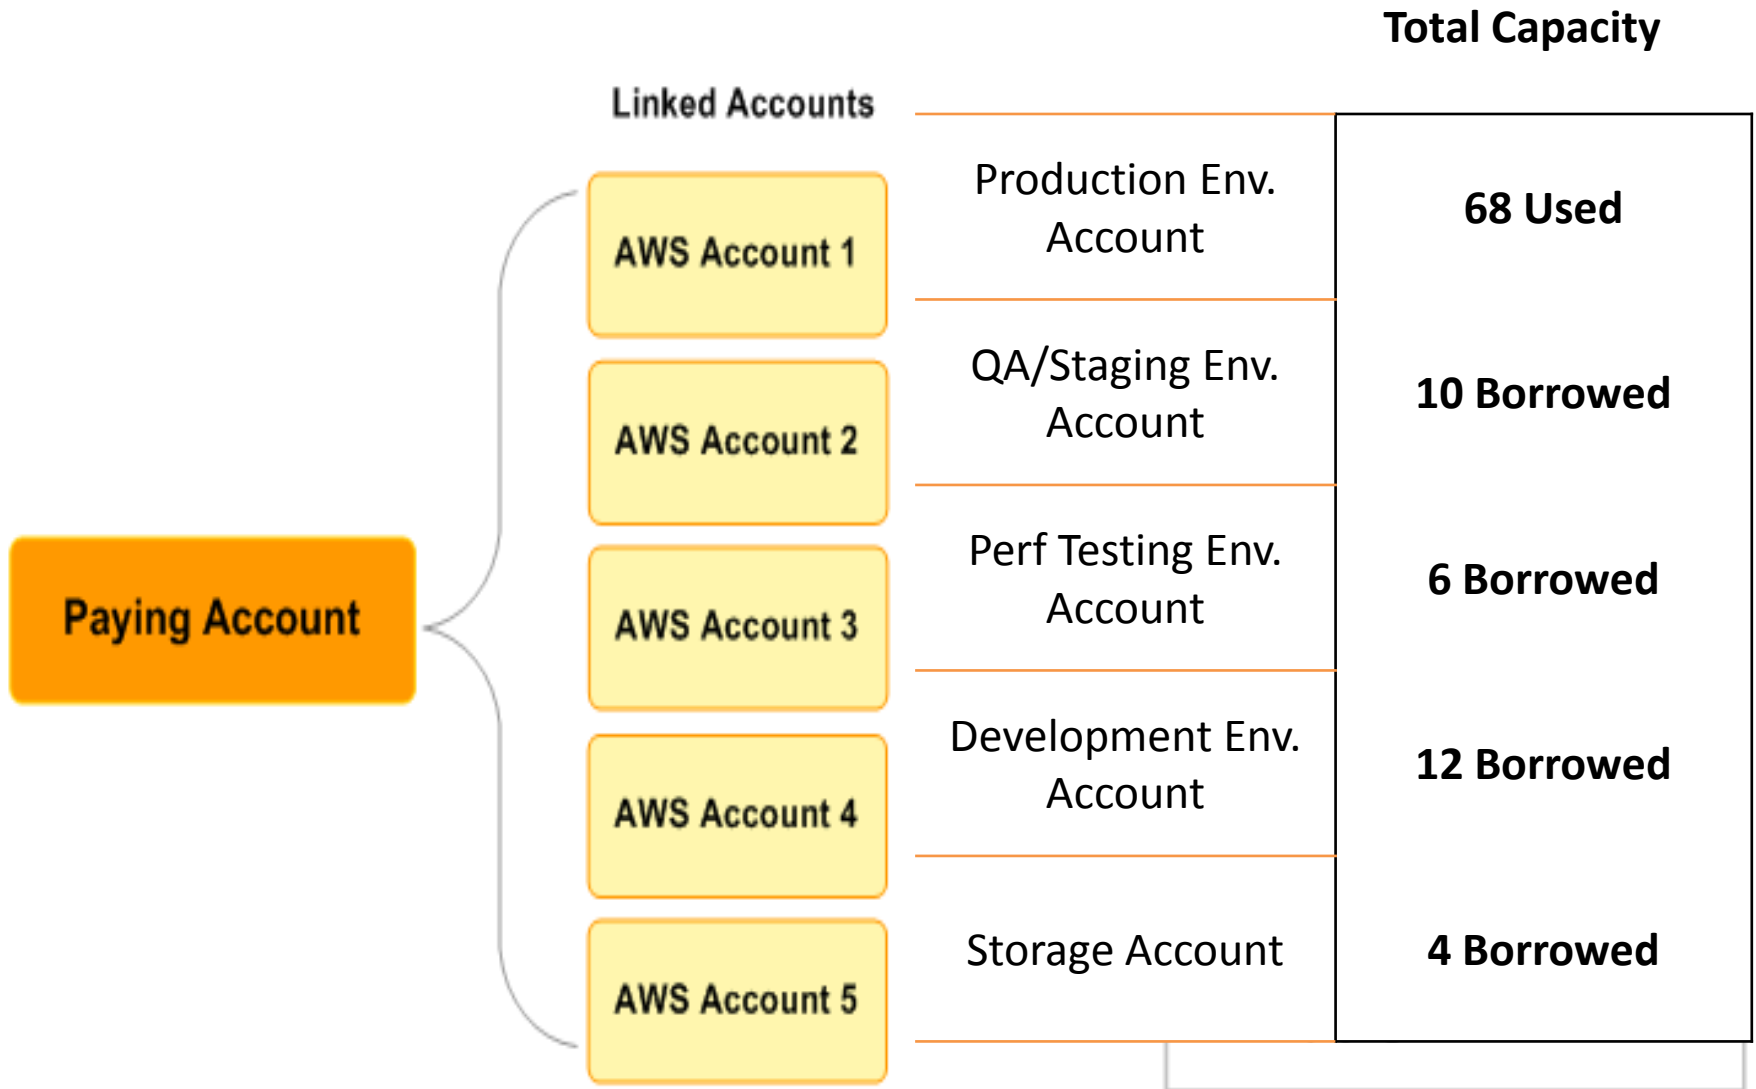
- **One Bill** for multiple accounts
- **Easy Tracking** of account charges (e.g., download CSV of cost data)
- **Volume Discounts** can be reached faster with combined usage
- **Reserved Instances** are shared across accounts (including RDS Reserved DBs)

# Over-Reserve the Production Environment





# Consolidated Billing Borrows Unused Reservations



## Consolidated Billing Advantages

- Production account is guaranteed to get burst capacity
  - Reservation is higher than normal usage level
  - Requests for more capacity always work up to reserved limit
  - Higher availability for handling unexpected peak demands
- No additional cost
  - Other lower priority accounts soak up unused reservations
  - Totals roll up in the monthly billing cycle

# Building Cost-Aware Cloud Architectures

**#1 Business Agility by Rapid Experimentation = Increased Revenue**

**#2 Business-driven Auto Scaling Architectures = Savings**

**#3 Mix and Match Reserved Instances with On-Demand = Savings**

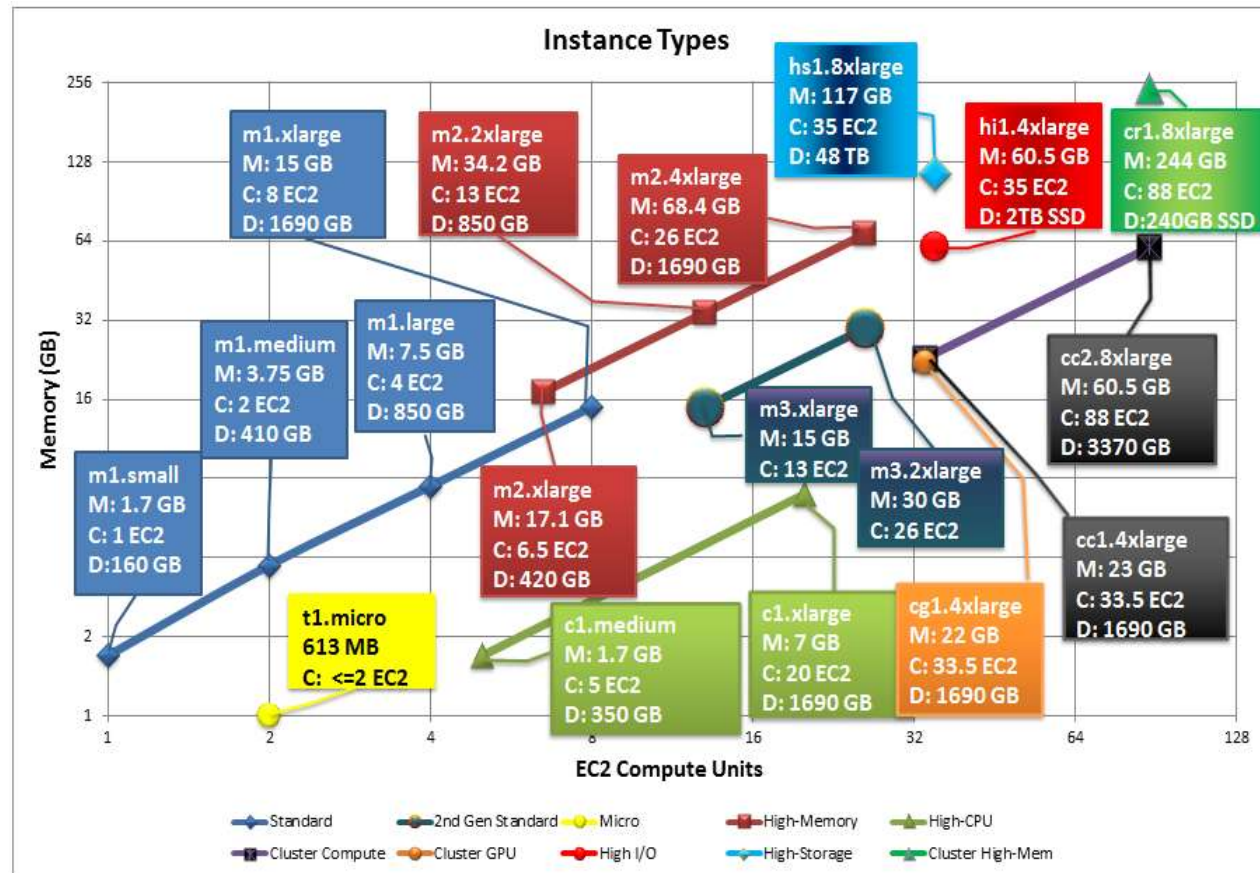
**#4 Consolidated Billing and Shared Reservations = Savings**

Continuous optimization in your  
architecture results in  
*recurring savings*  
as early as your next month's bill



# Right-size your cloud: Use only what you need

- An instance type for every purpose
- Assess your memory & CPU requirements
  - ~~Fit your application to the resource~~
  - Fit the resource to your application
- Only use a larger instance when needed



# Reserved Instance Marketplace

Buy a smaller term instance  
Buy instance with different OS or type  
Buy a Reserved instance in different region

Sell your unused Reserved Instance  
Sell unwanted or over-bought capacity  
Further reduce costs by optimizing

## Purchase Reserved Instances

[Cancel](#) ✕

Platform:

Term:

Instance Type:

Tenancy:

Availability Zone:

Offering Type:

Search

Seller	Term	Effective Rate	Upfront Price	Hourly Rate	Availability Zone	Offering Type	Quantity Available	Desired Quantity	
AWS	36 months	\$0.239	\$2320.00	\$0.151	ap-southeast-1a	Heavy Utilization	Unlimited	<input type="text" value="1"/>	<a href="#">Add to Cart</a>
AWS	36 months	\$0.239	\$2320.00	\$0.151	ap-southeast-1b	Heavy Utilization	Unlimited	<input type="text" value="1"/>	<a href="#">Add to Cart</a>
3rd Party	4 months	\$0.339	\$400.00	\$0.20	ap-southeast-1a	Heavy Utilization	2	<input type="text" value="1"/>	<a href="#">Add to Cart</a>
AWS	12 months	\$0.359	\$1478.00	\$0.19	ap-southeast-1a	Heavy Utilization	Unlimited	<input type="text" value="1"/>	<a href="#">Add to Cart</a>
AWS	12 months	\$0.359	\$1478.00	\$0.19	ap-southeast-1b	Heavy Utilization	Unlimited	<input type="text" value="1"/>	<a href="#">Add to Cart</a>

# Instance Type Optimization

## Older m1 and m2 families

- Slower CPUs
- Higher response times
- Smaller caches (6MB)
- Oldest m1.xl 15GB/8ECU/48c
- Old m2.xl 17GB/6.5ECU/41c
- ~16 ECU/\$/hr

## Latest m3 family

- Faster CPUs
- Lower response times
- Bigger caches (20MB)
- Even faster for Java vs. ECU
- New m3.xl 15GB/13 ECU/50c
- 26 ECU/\$/hr – 62% better!
- Java measured even higher
- Deploy fewer instances

# Building Cost-Aware Cloud Architectures

**#1 Business Agility by Rapid Experimentation = Increased Revenue**

**#2 Business-driven Auto Scaling Architectures = Savings**

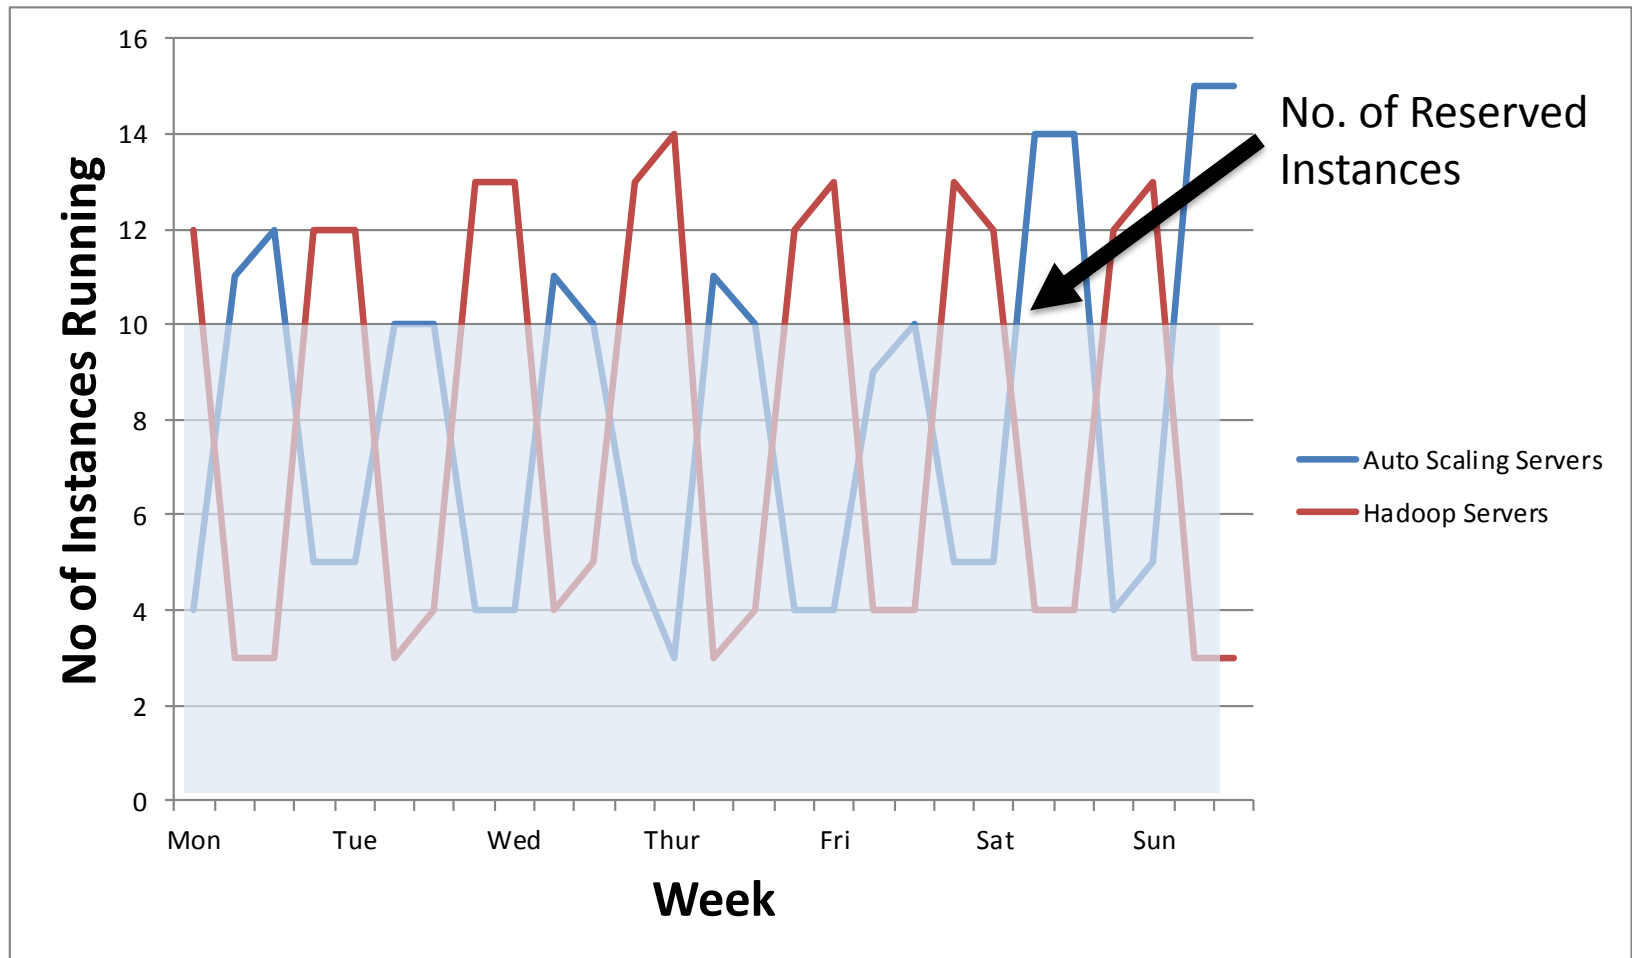
**#3 Mix and Match Reserved Instances with On-Demand = Savings**

**#4 Consolidated Billing and Shared Reservations = Savings**

**#5 Always-on Instance Type Optimization = Recurring Savings**

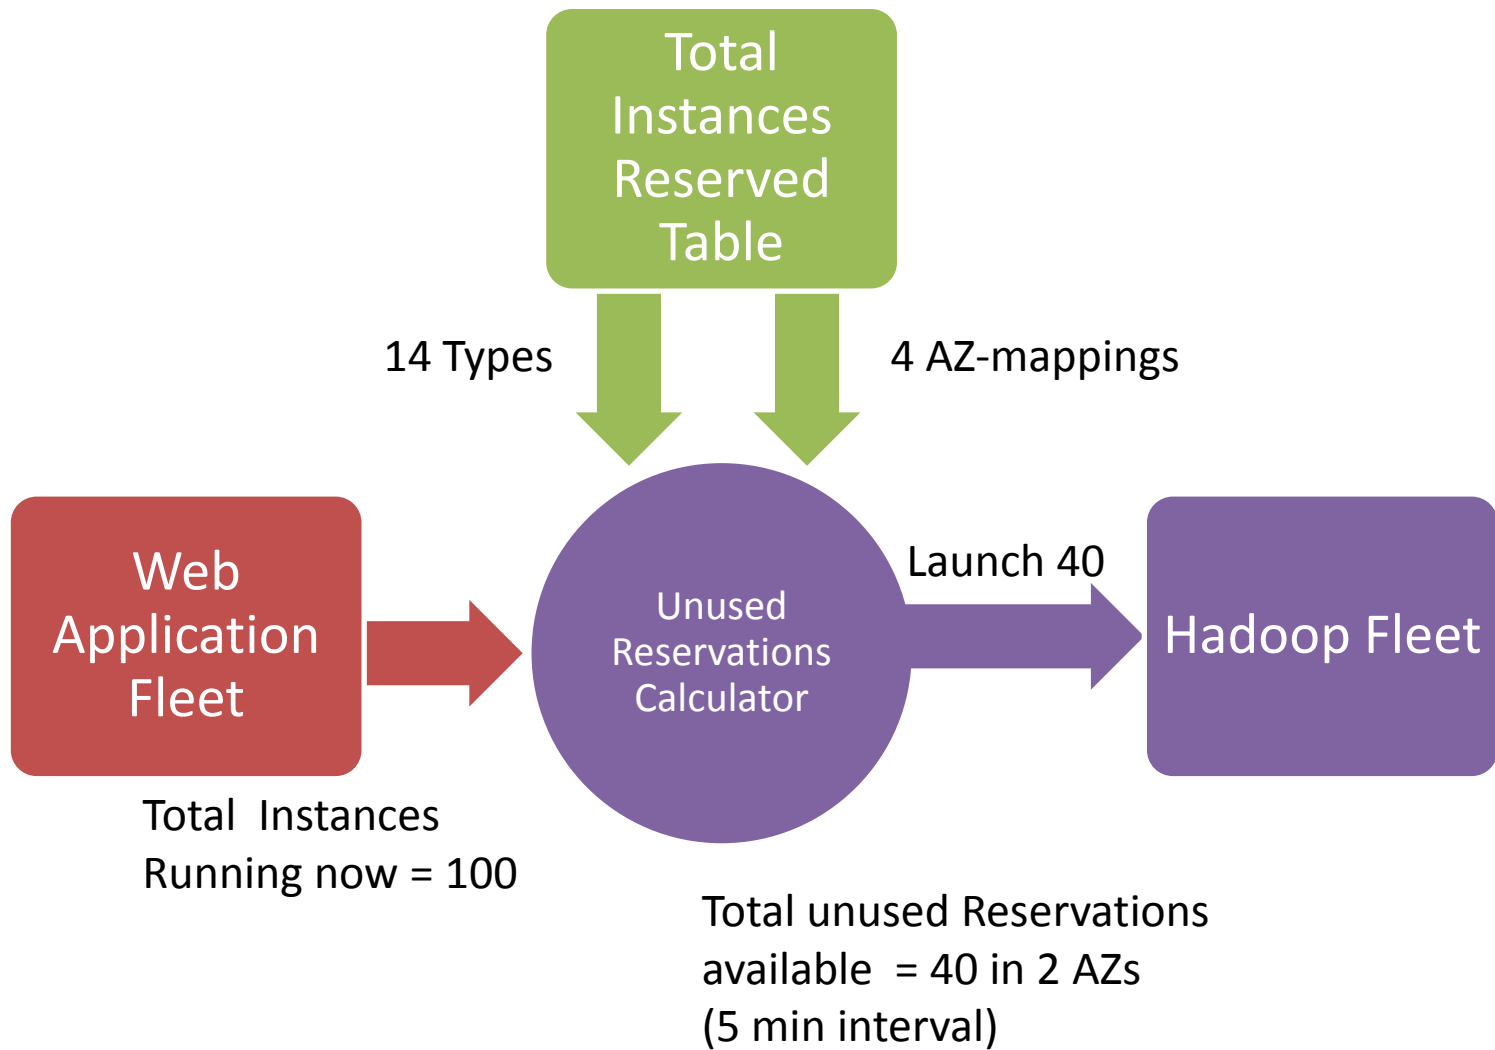


## Follow the Customer (Run web servers) during the day



Follow the Money (Run Hadoop clusters) at night





## Soaking up unused reservations

Unused reserved instances is published as a metric

### Netflix Data Science ETL Workload

- Daily business metrics roll-up
- Starts after midnight
- EMR clusters started using hundreds of instances

### Netflix Movie Encoding Workload

- Long queue of high and low priority encoding jobs
- Can soak up 1000's of additional unused instances

# Building Cost-Aware Cloud Architectures

**#1 Business Agility by Rapid Experimentation = Increased Revenue**

**#2 Business-driven Auto Scaling Architectures = Savings**

**#3 Mix and Match Reserved Instances with On-Demand = Savings**

**#4 Consolidated Billing and Shared Reservations = Savings**

**#5 Always-on Instance Type Optimization = Recurring Savings**

**#6 Follow the Customer (Run web servers) during the day  
Follow the Money (Run Hadoop clusters) at night**



# Thank you!

Jinesh Varia and Adrian Cockcroft

jvaria@amazon.com @jinman

acockcroft@netflix.com @adrianco

# Slideshare.net/Netflix Details

- Meetup S1E3 July – Featuring Contributors Eucalyptus, IBM, Paypal, Riot Games
  - <http://techblog.netflix.com/2013/07/netflixoss-meetup-series-1-episode-3.html>
- Lightning Talks March S1E2
  - <http://www.slideshare.net/RuslanMeshenberg/netflixoss-meetup-lightning-talks-and-roadmap>
- Lightning Talks Feb S1E1
  - <http://www.slideshare.net/RuslanMeshenberg/netflixoss-open-house-lightning-talks>
- Asgard In Depth Feb S1E1
  - <http://www.slideshare.net/joesondow/asgard-overview-from-netflix-oss-open-house>
- Security Architecture
  - [http://www.slideshare.net/jason\\_chan/resilience-and-security-scale-lessons-learned/](http://www.slideshare.net/jason_chan/resilience-and-security-scale-lessons-learned/)

# Takeaways

*Cloud Native Manages Scale and Complexity at Speed*

*NetflixOSS makes it easier for everyone to become Cloud Native*

*Rethink deployments and turn things off to save money!*

<http://netflix.github.com>

<http://techblog.netflix.com>

<http://slideshare.net/Netflix>

<http://www.linkedin.com/in/adriancockcroft>



@adrianco #netflixcloud @NetflixOSS