

Decomposing SMACK Stack

Spark & Mesos Internals

Anton Kirillov
intro by Sebastian Stoll

Apache Spark Meetup
Oooyala, March 2016

Who is this guy?

[@antonkirillov](#)

- Staff Engineer in Data Team @ Ooyala
 - Scala programmer
 - Focused on distributed systems
 - Building data platforms with SMACK/Hadoop
 - Ph.D. in Computer Science
-
- blog: datastrophic.io
 - github: github.com/datastrophic

Roadmap

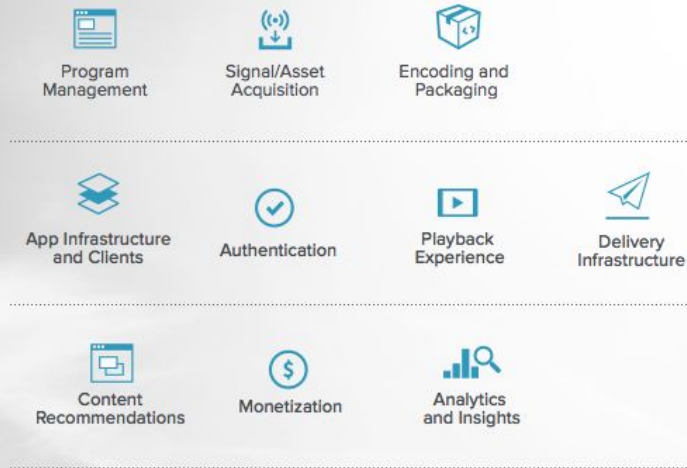
- Intro to Ooyala stack
- SMACK stack
 - Overview
 - Architecture design options
- Spark
 - Core concepts & execution workflow
 - Architecture
- Mesos
 - Cluster resource management
 - Architecture and scheduling
 - Frameworks
 - Spark on Mesos



REVOLUTIONIZING DIGITAL TV

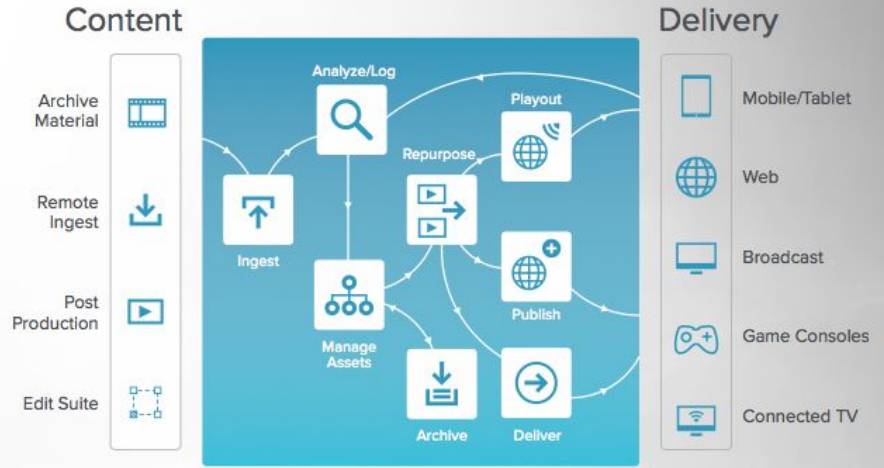


What Ooyala does



Online Video Platform

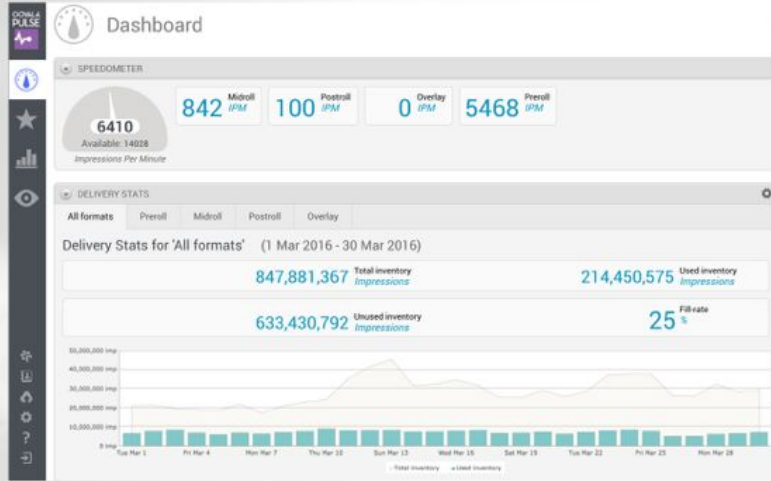
MIO



Ooyala Media Logistics

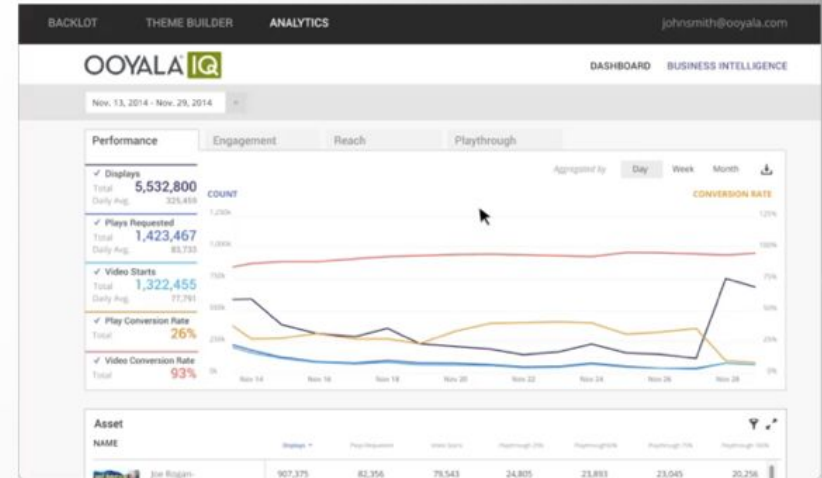
What Ooyala does

OOYALA® PULSE 



Ad Tech

OOYALA® IQ 



Analytics

Ooyala does that



MESOS



akka



kafka

- Ooyala IQ
- Data operations at scale
- Stream analytics

- Microservice infrastructure
- Spark cluster resource manager

- Load testing
- General tooling

- Ooyala IQ
- Ooyala Pulse Insight
- Delivery metrics
- Forecasting

- Ooyala Now
- Data Flow

SMACK Stack Overview

components and architecture designs

SMACK Stack



- **Spark** - a generalized framework for distributed data processing supporting in-memory data caching and reuse across computations



- **Mesos** - cluster resource management system that provides efficient resource isolation and sharing across distributed applications



- **Akka** - a toolkit and runtime for building highly concurrent, distributed, and resilient message-driven applications on the JVM




- **Cassandra** - distributed, highly available database designed to handle large amounts of data across multiple datacenters



- **Kafka** - a high-throughput, low-latency distributed messaging system designed for handling real-time data feeds

Storage Layer: Cassandra



- Pros:
 - optimized for heavy write loads
 - configurable CA (CAP)
 - linearly scalable
 - XDCR support
 - easy cluster resizing and inter-DC data migration
- Cons:
 - data model (distributed nested sorted map)
 - designed for fast serving but not batch processing
 - not well-suited for ad-hoc queries against historical raw data

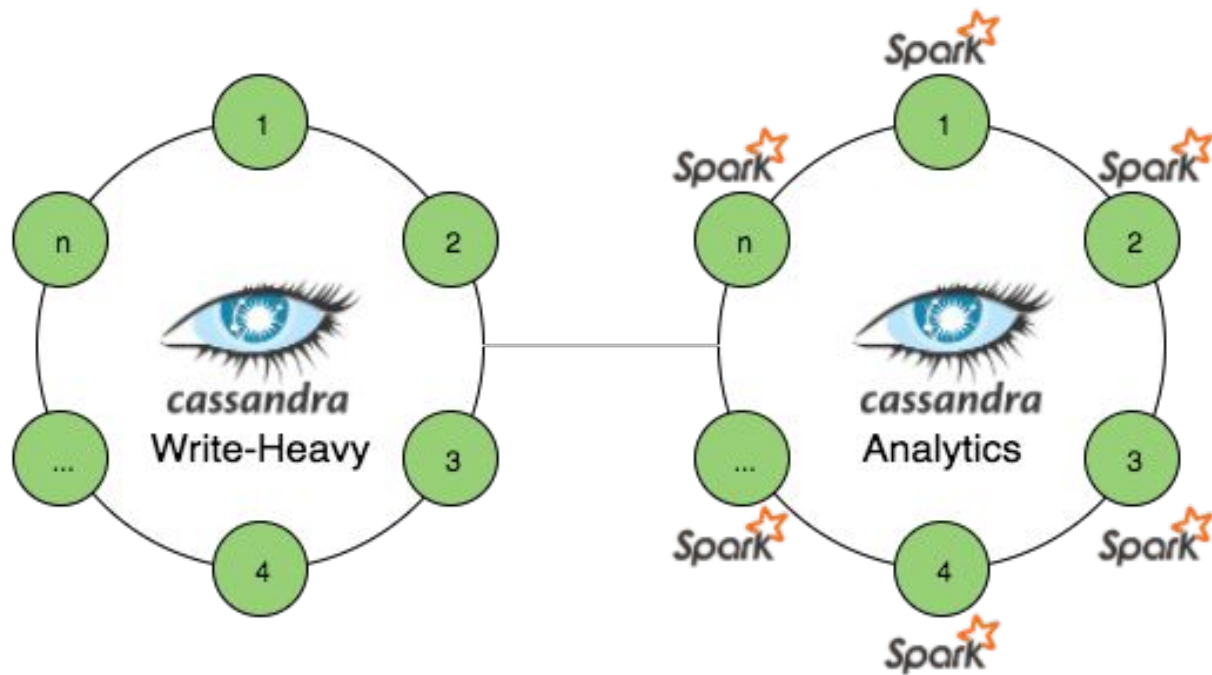
Fixing NoSQL limitations with Spark

```
//joining raw events with rolled-up and grouping by type
```

```
sqlContext.sql {  
  SELECT  
    events.campaignId,  
    events.eventType,  
    events.value + campaigns.total as total_events  
  FROM events  
  JOIN campaigns  
  ON events.campaignId = campaigns.id AND events.eventType = campaigns.eventType  
}.stripMargin  
}.registerTempTable("joined")
```

```
sqlContext.sql {  
  SELECT campaignId, eventType, sum(total_events) as total  
  FROM joined  
  GROUP BY campaignId, eventType  
}.stripMargin  
}.saveAsCassandraTable("keyspace", "totals")
```

Architecture of Spark/Cassandra Clusters

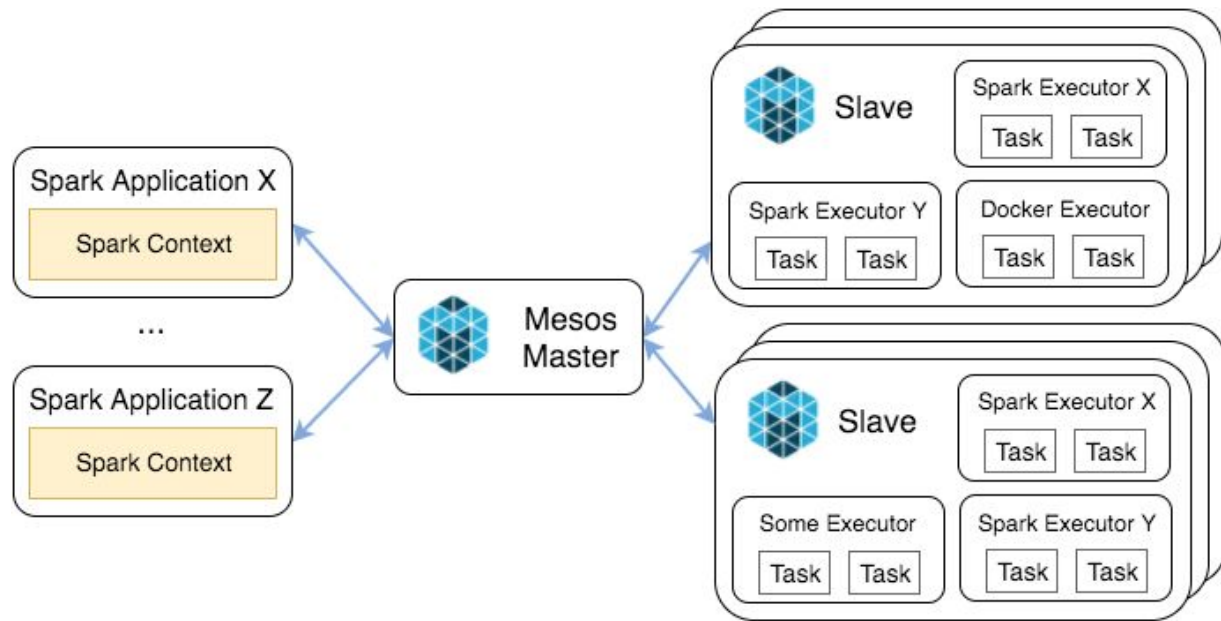


Separate Write & Analytics:

- clusters can be scaled independently
- data is replicated by Cassandra asynchronously
- Analytics has different Read/Write load patterns
- Analytics contains additional data and processing results
- Spark resource impact limited to only one DC

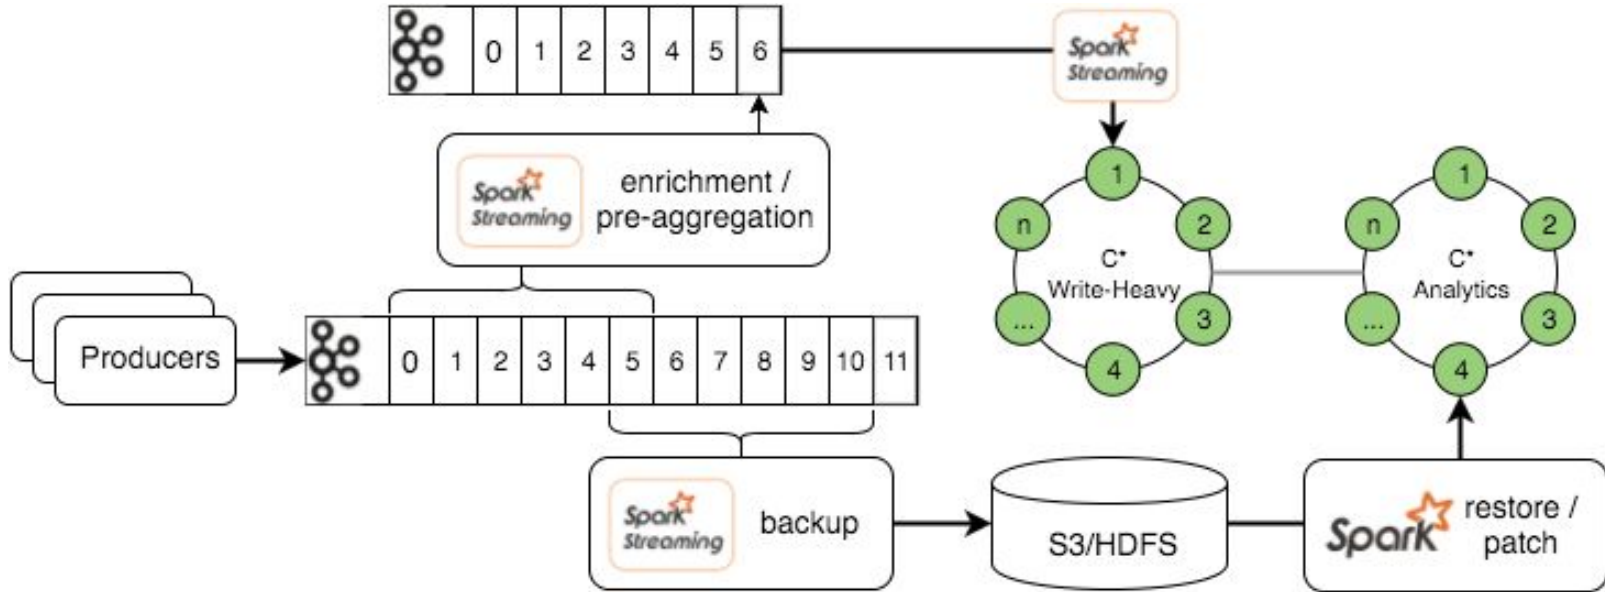
To fully facilitate Spark-C* connector data locality awareness,
Spark workers should be collocated with Cassandra nodes (**gotcha**: CL=ONE)

Mesos as Spark cluster manager



- fine-grained resource sharing between Spark and other applications
- scalable partitioning between multiple instances of Spark
- unified platform for running various applications (frameworks)
- fault-tolerant and scalable

Stream Processing with Kafka and Spark



- be prepared for failures and broken data
- backup and patching strategies should be designed upfront
- patch/restore if time interval could be done by replay if store is idempotent

Spark Streaming with Kafka


```
val streamingContext = new StreamingContext(sc.getConf, Seconds(10))
```

```
val eventStream = KafkaUtils.createStream(  
  ssc = streamingContext,  
  zkQuorum = "zoo01,zoo02,zoo03",  
  groupId = "spark_consumer",  
  topics = Map("raw_events" -> 3)  
)
```

```
eventStream.map(_._2.toEvent)  
             .saveToCassandra(keyspace, table)
```

```
streamingContext.start()  
streamingContext.awaitTermination()
```

Data Ingestion with Akka



- actor model implementation for JVM
- message-based and asynchronous
- easily scalable from one process to cluster of machines
- actor hierarchies with parental supervision
- easily packages in Docker to be run on Mesos

Akka Http microservice

```
val config = new ProducerConfig(KafkaConfig())
lazy val producer = new KafkaProducer[A, A](config)

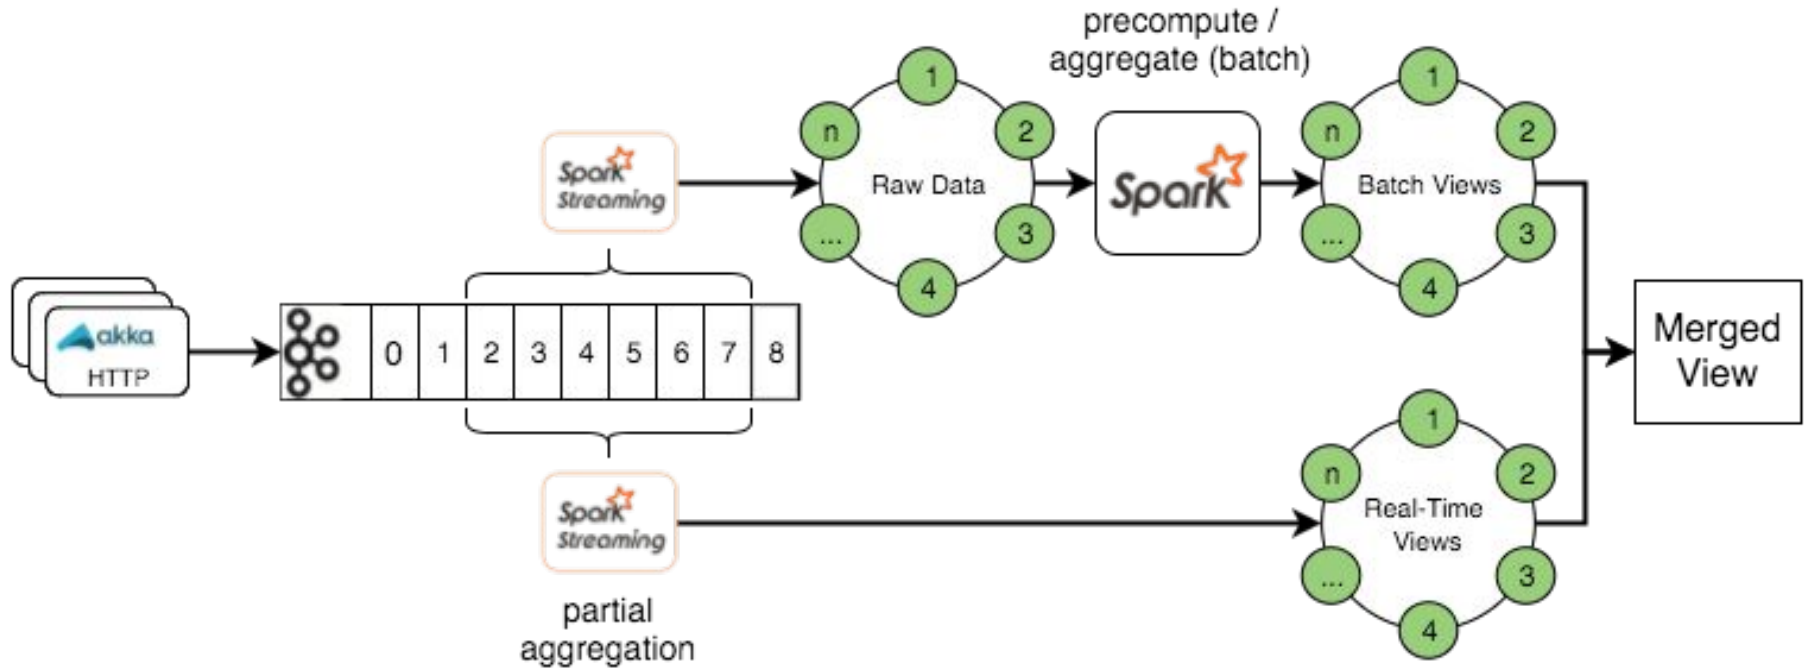
val routes: Route = {
  post{
    decodeRequest{
      entity(as[String]){ str =>
        JsonParser.parse(str).validate[Event] match {
          case s: JsSuccess[String] =>
            producer.send(new KeyedMessage(topic, str))
            system.actorOf(Props[CassandraWriter]) ! s.get
          case e: JsError => BadRequest -> JsError.toFlatJson(e).toString()
        }
      }
    }
  }
}

object AkkaHttpMicroservice extends App with Service {
  Http().bindAndHandle(routes, config.getString("http.interface"), config.getInt("http.port"))
}
```

Writing to Cassandra with Akka

```
class CassandraWriterActor extends Actor with ActorLogging {  
  
  //for demo purposes, session initialized here  
  val session = Cluster.builder()  
    .addContactPoint("cassandra.host")  
    .build()  
    .connect()  
  
  override def receive: Receive = {  
    case event: Event =>  
      val statement = new SimpleStatement(event.createQuery)  
        .setConsistencyLevel(ConsistencyLevel.QUORUM)  
  
      Try(session.execute(statement)) match {  
        case Failure(ex) => //error handling code  
        case Success => sender ! WriteSuccessful  
      }  
    }  
  }  
}
```

Lambda Architecture with SMACK



- when design meets reality it's hard to implement canonical architecture
- depending on the use case it's easy to implement Kappa architecture as well

SMACK stack:

- concise toolbox for wide variety of data processing scenarios
- battle-tested and widely used software with large communities
- easy scalability and replication of data while preserving low latencies
- unified cluster management for heterogeneous loads
- single platform for any kind of applications
- implementation platform for different architecture designs
- really short time-to-market (e.g. for MVP verification)

Apache Spark in Depth

core concepts, architecture & internals

Meet Spark

- Generalized framework for distributed data processing (batch, graph, ML)
- Scala collections functional API for manipulating data at scale
- In-memory data caching and reuse across computations
- Applies set of coarse-grained transformations over partitioned data
- Failure recovery relies on lineage to recompute failed tasks
- Supports majority of input formats and integrates with Mesos / YARN

Spark makes data engineers happy

Backup/restore of Cassandra tables in Parquet

```
def backup(config: Config) {  
  sc.cassandraTable(config.keyspace, config.table).map(_._toEvent).toDF()  
  .write.parquet(config.path)  
}  
  
def restore(config: Config) {  
  sqlContext.read.parquet(config.path)  
  .map(_._toEvent).saveToCassandra(config.keyspace, config.table)  
}
```

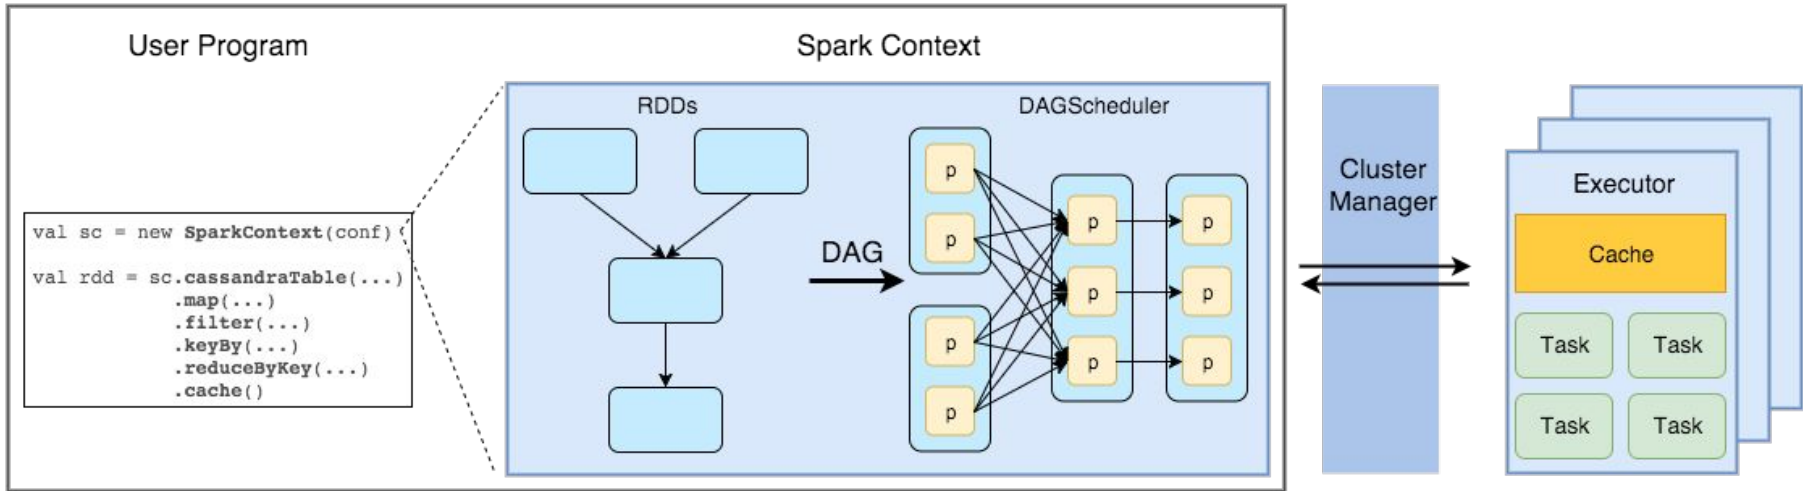
Query different data sources to identify discrepancies

```
sqlContext.sql {  
  """  
  SELECT count()  
  FROM cassandra_event_rollups  
  JOIN mongo_event_rollups  
  ON cassandra_event_rollups.uuid = mongo_event_rollups.uuid  
  WHERE cassandra_event_rollups.value != mongo_event_rollups.value  
  """  
  .stripMargin  
}
```

Core Concepts

Spark Application

Workers



RDD: Resilient Distributed Dataset

- A fault-tolerant, immutable, parallel data structure
- Provides API for
 - manipulating the collection of elements (transformations and materialization)
 - persisting intermediate results in memory for later reuse
 - controlling partitioning to optimize data placement
- Can be created through deterministic operation
 - from storage (distributed file system, database, plain file)
 - from another RDD
- Stores information about parent RDDs
 - for execution optimization and operations pipelining
 - to recompute the data in case of failure

RDD: a developer's view

- Distributed immutable data + lazily evaluated operations
 - partitioned data + iterator
 - transformations & actions
- An interface defining 5 main properties

a list of partitions (e.g. splits in Hadoop)

def getPartitions: Array[Partition]

a list of dependencies on other RDDs

def getDependencies: Seq[Dependency[_]]

a function for computing each split

def compute(split: Partition, context: TaskContext): Iterator[T]

(optional) a list of preferred locations to compute each split on

def getPreferredLocations(split: Partition): Seq[String] = Nil

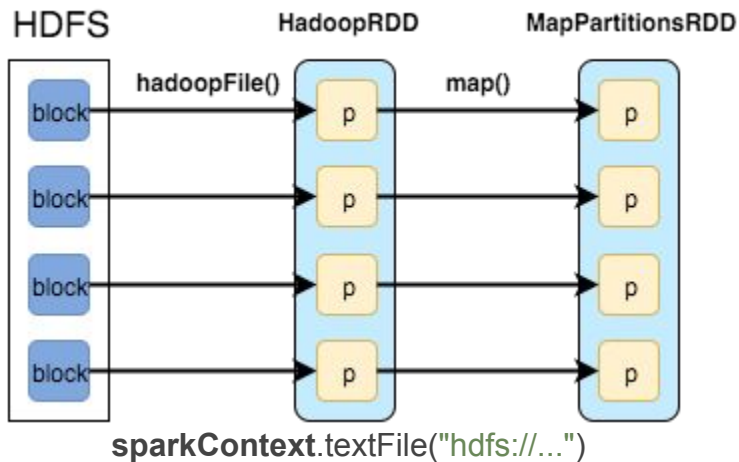
(optional) a partitioner for key-value RDDs

val partitioner: Option[Partitioner] = None

lineage

execution optimization

RDDs Example



- **HadoopRDD**

- `getPartitions` = HDFS blocks
- `getDependencies` = None
- `compute` = load block in memory
- `getPreferredLocations` = HDFS block locations
- `partitioner` = None

- **MapPartitionsRDD**

- `getPartitions` = same as parent
- `getDependencies` = parent RDD
- `compute` = compute parent and apply `map()`
- `getPreferredLocations` = same as parent
- `partitioner` = None

RDD Operations

- Transformations

- apply user function to every element in a partition (or to the whole partition)
- apply aggregation function to the whole dataset (groupBy, sortBy)
- introduce dependencies between RDDs to form DAG
- provide functionality for repartitioning (repartition, partitionBy)

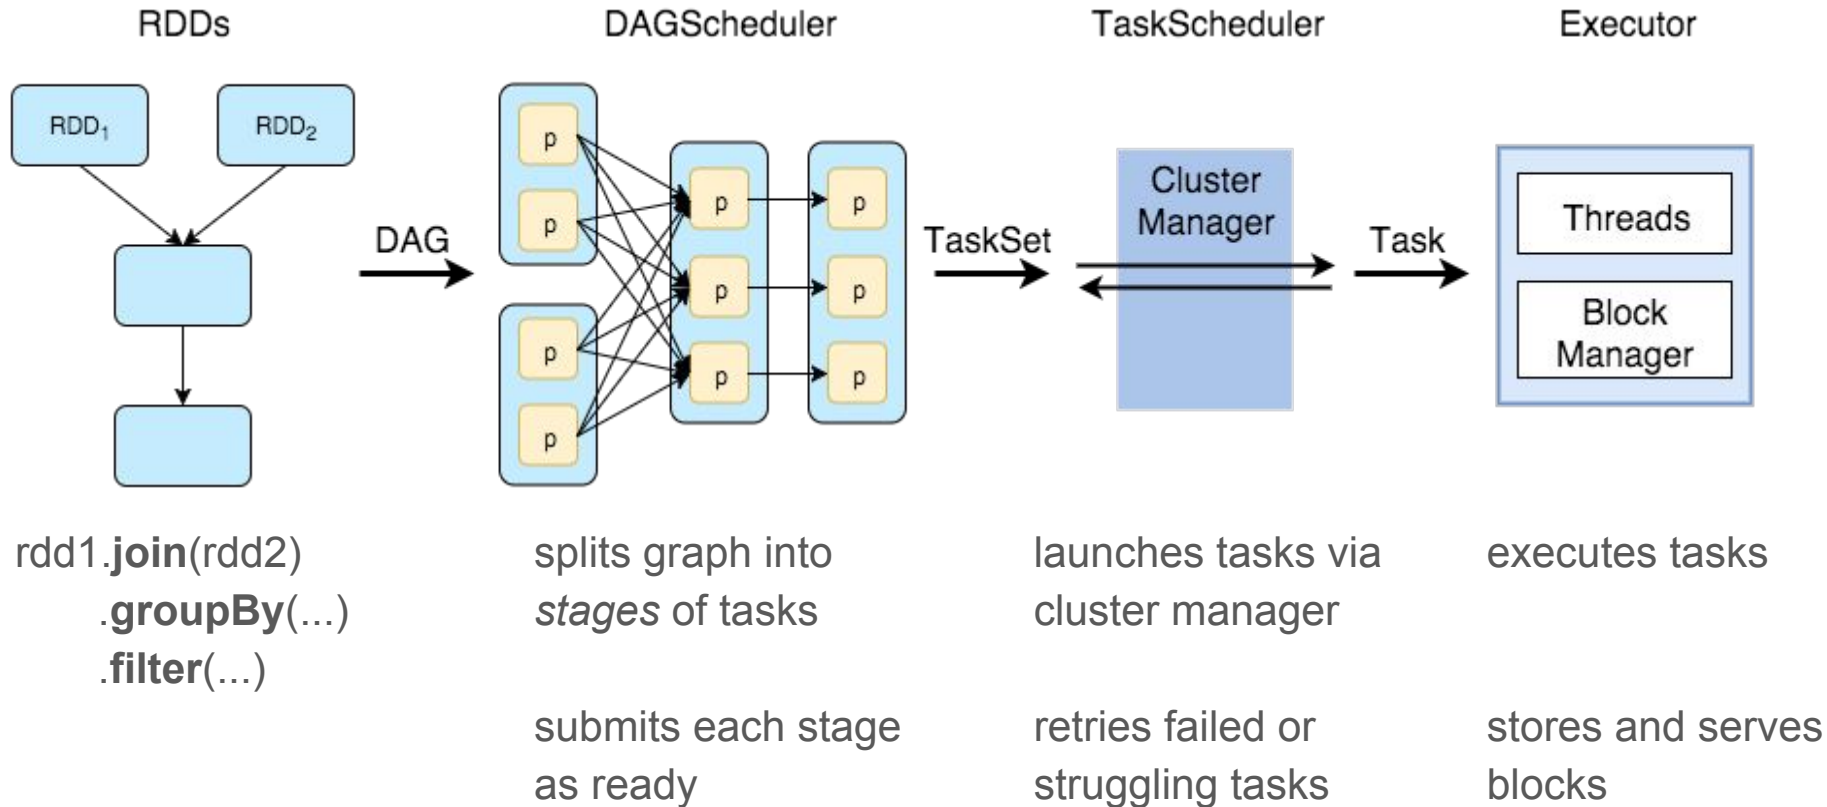
- Actions

- trigger job execution
- used to materialize computation results

- Extra: persistence

- explicitly store RDDs in memory, on disk or off-heap (cache, persist)
- checkpointing for truncating RDD lineage

Execution workflow



Code sample: joining aggregated and raw data

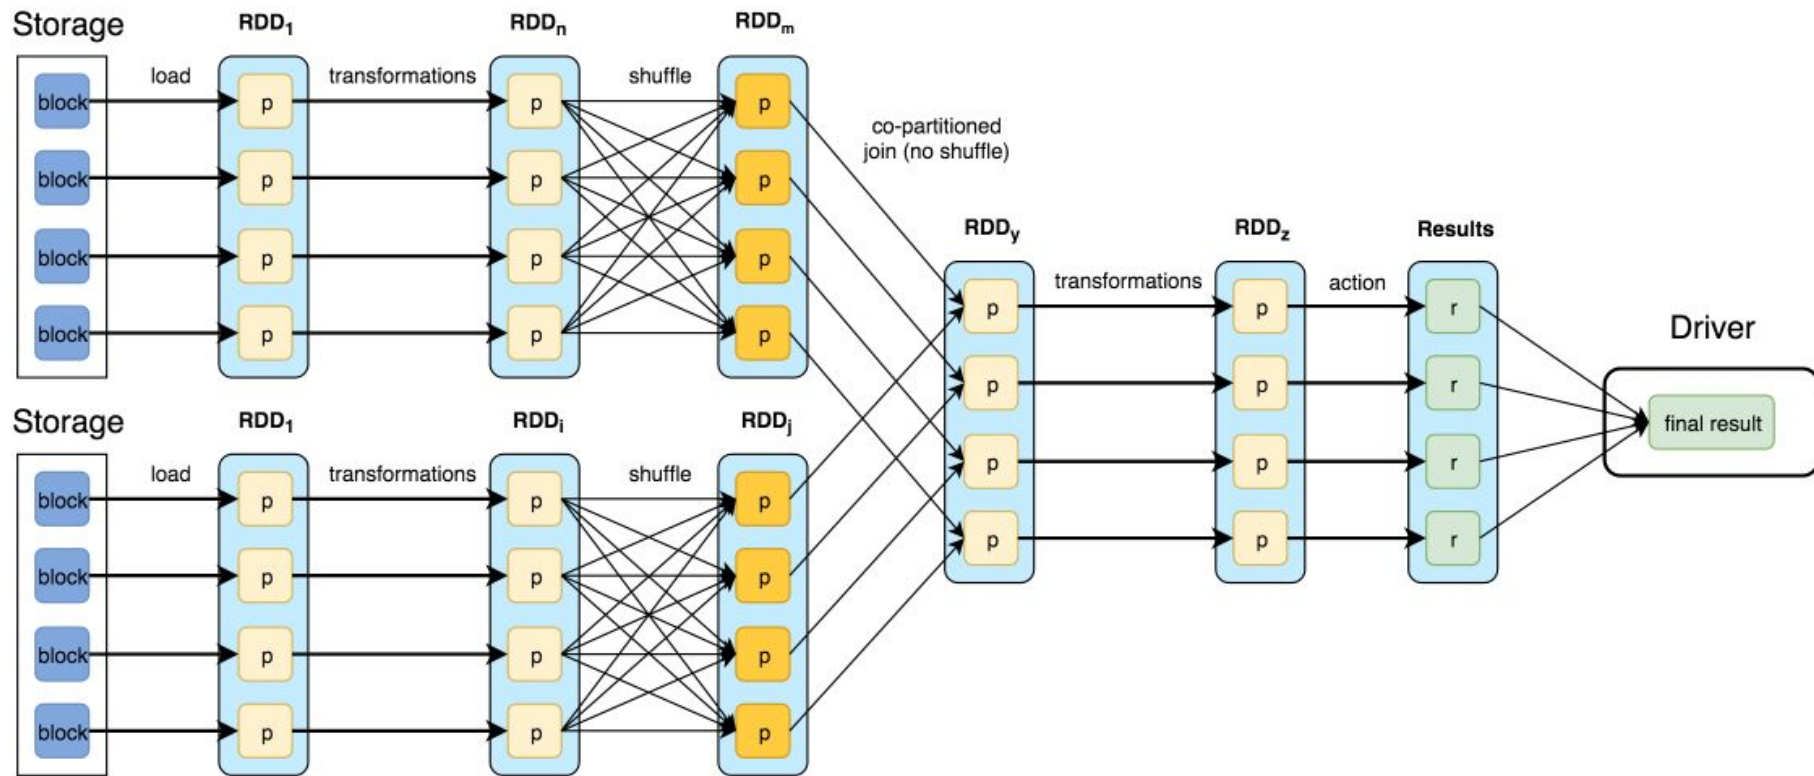
```
//aggregate events after specific date for given campaign
val events = sc.cassandraTable("demo", "event")
    .map(_._toEvent)
    .filter(event => event.campaignId == campaignId && event.time.isAfter(watermark))
    .keyBy(_._eventType)
    .reduceByKey(_ + _)
    .cache()

//aggregate campaigns by type
val campaigns = sc.cassandraTable("demo", "campaign")
    .map(_._toCampaign)
    .filter(campaign => campaign.id == campaignId && campaign.time.isBefore(watermark))
    .keyBy(_._eventType)
    .reduceByKey(_ + _)
    .cache()






//joined rollups and raw events
val joinedTotals = campaigns.join(events)
    .map { case (key, (campaign, event)) => CampaignTotals(campaign, event) }
    .collect()

//count totals separately
val eventTotals = events.map { case (t, e) => s"$t -> ${e.value}" }.collect()
val campaignTotals = campaigns.map { case (t, e) => s"$t -> ${e.value}" }.collect()
```

DAG




Dependency types




- **Narrow (pipelineable)**
 - each partition of the parent RDD is used by at most one partition of the child RDD
 - allow for pipelined execution on one cluster node
 - failure recovery is more efficient as only lost parent partitions need to be recomputed
- **Wide (shuffle)**
 - multiple child partitions may depend on one parent partition
 - require data from all parent partitions to be available and to be shuffled across the nodes
 - if some partition is lost from all the ancestors a complete recomputation is needed

Stages and Tasks



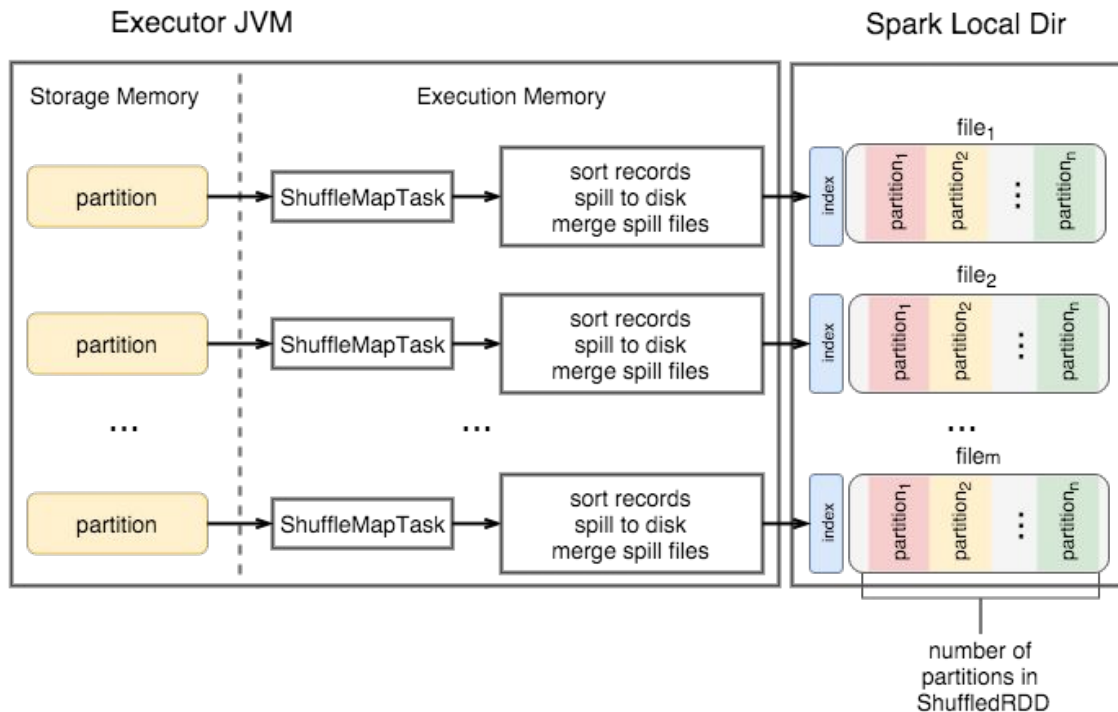
- Stages breakdown strategy
 - check backwards from final RDD
 - add each “narrow” dependency to the current stage
 - create new stage when there’s a shuffle dependency
- Tasks
 - *ShuffleMapTask* partitions its input for shuffle
 - *ResultTask* sends its output to the driver

Shuffle



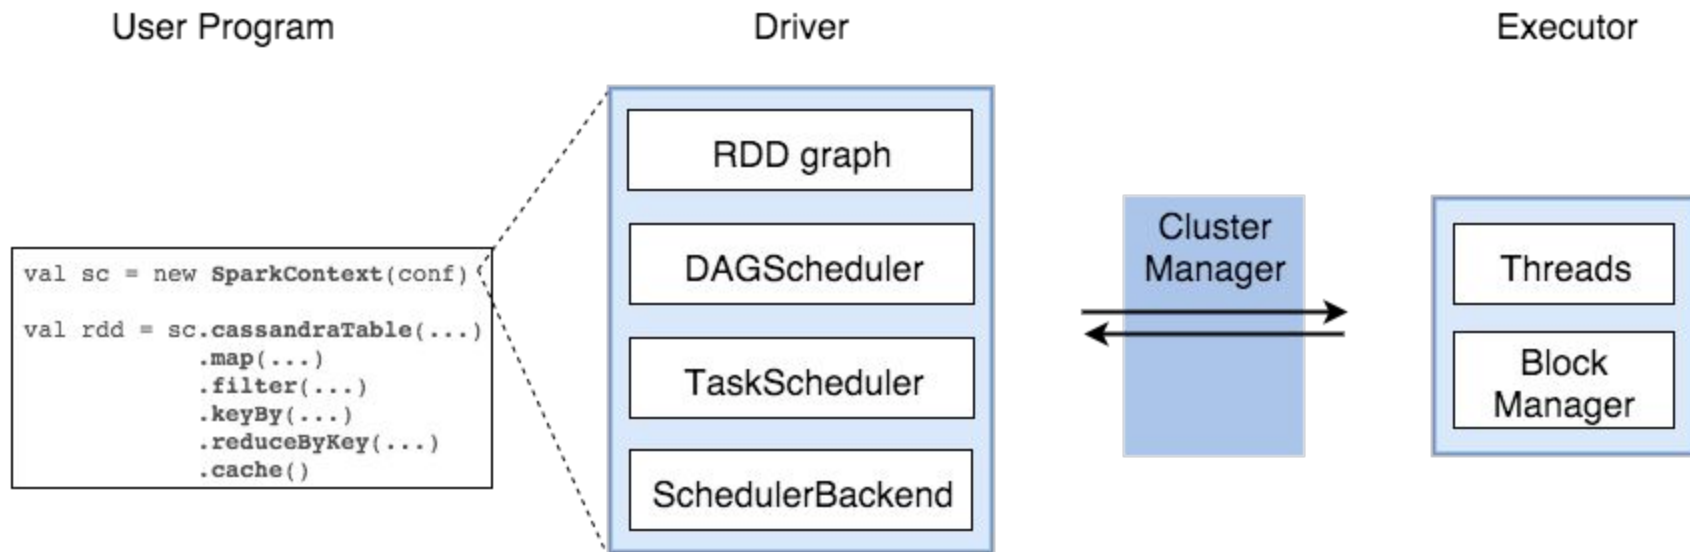
- Shuffle Write
 - redistributes data among partitions and writes files to disk
 - each *shuffle* task creates one file with regions assigned to reducer
 - *sort shuffle* uses in-memory sorting with spillover to disk to get final result
- Shuffle Read
 - fetches the files and applies *reduce()* logic
 - if data ordering is needed then it is sorted on “reducer” side for any type of shuffle

Sort Shuffle



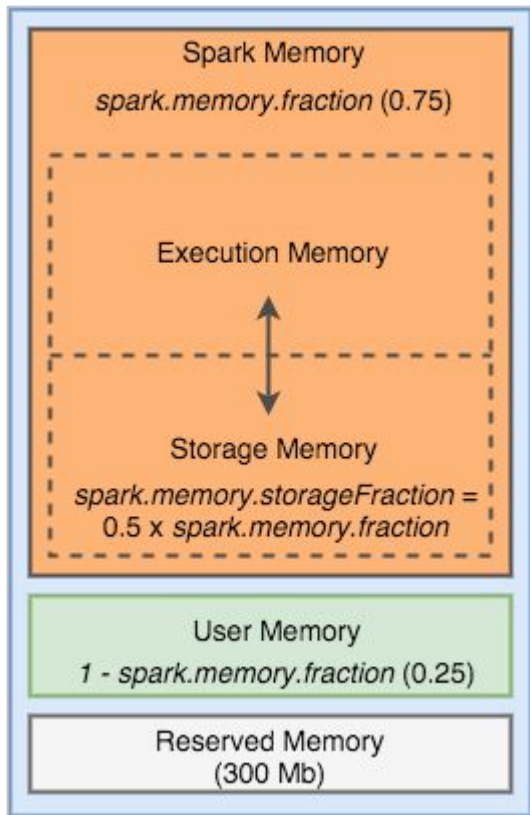
- Incoming records accumulated and sorted in memory according to their target partition ids
- Sorted records are written to file or multiple files if spilled and then merged
- *index* file stores offsets of the data blocks in the data file
- Sorting without deserialization is possible under certain conditions (SPARK-7081)

Spark Components



Memory Management in Spark 1.6

JVM Heap



- **Execution Memory**
 - storage for data needed during tasks execution
 - shuffle-related data
- **Storage Memory**
 - storage of cached RDDs and broadcast variables
 - possible to borrow from execution memory (spill otherwise)
 - safeguard value is 0.5 of Spark Memory when cached blocks are immune to eviction
- **User Memory**
 - user data structures and internal metadata in Spark
 - safeguarding against OOM
- **Reserved memory**
 - memory needed for running executor itself and not strictly related to Spark

Execution Modes

- `spark-shell --master [local | spark | yarn-client | mesos]`
 - launches REPL connected to specified cluster manager
 - always runs in client mode
- `spark-submit --master [local | spark:// | mesos:// | yarn] spark-job.jar`
 - launches assembly jar on the cluster
- Masters
 - **local[k]** - run Spark locally with K worker threads
 - **spark** - launches driver app on Spark Standalone installation
 - **mesos** - driver will spawn executors on Mesos cluster (**deploy-mode: client | cluster**)
 - **yarn** - same idea as with Mesos (**deploy-mode: client | cluster**)
- Deploy Modes
 - **client** - driver executed as a separate process on the machine where it has been launched and spawns executors
 - **cluster** - driver launched as a container using underlying cluster manager

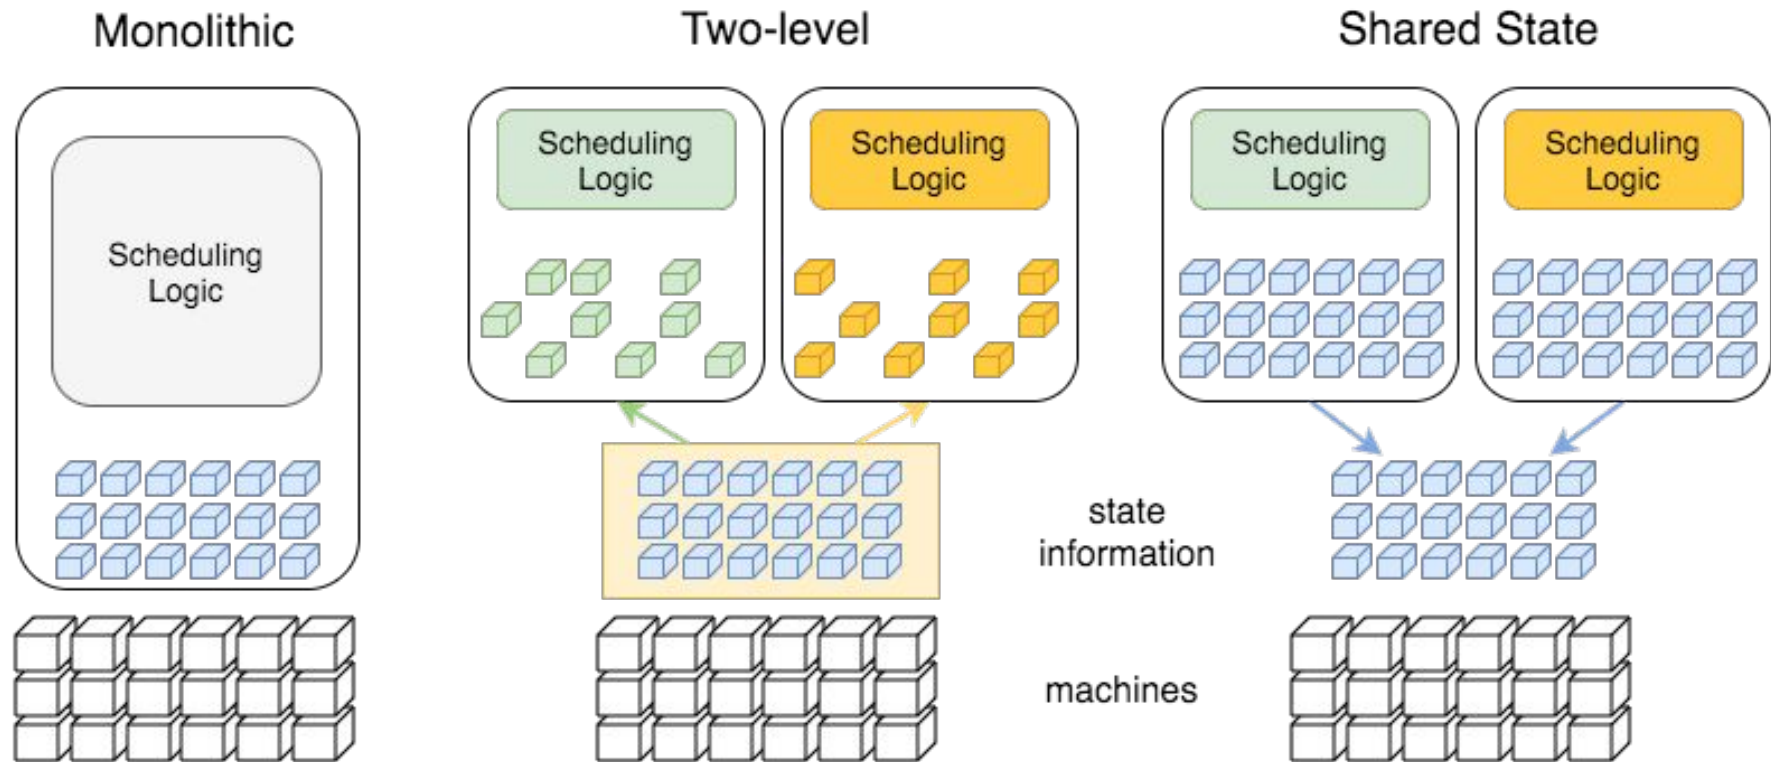
Apache Mesos

architecture, scheduling, frameworks & Spark

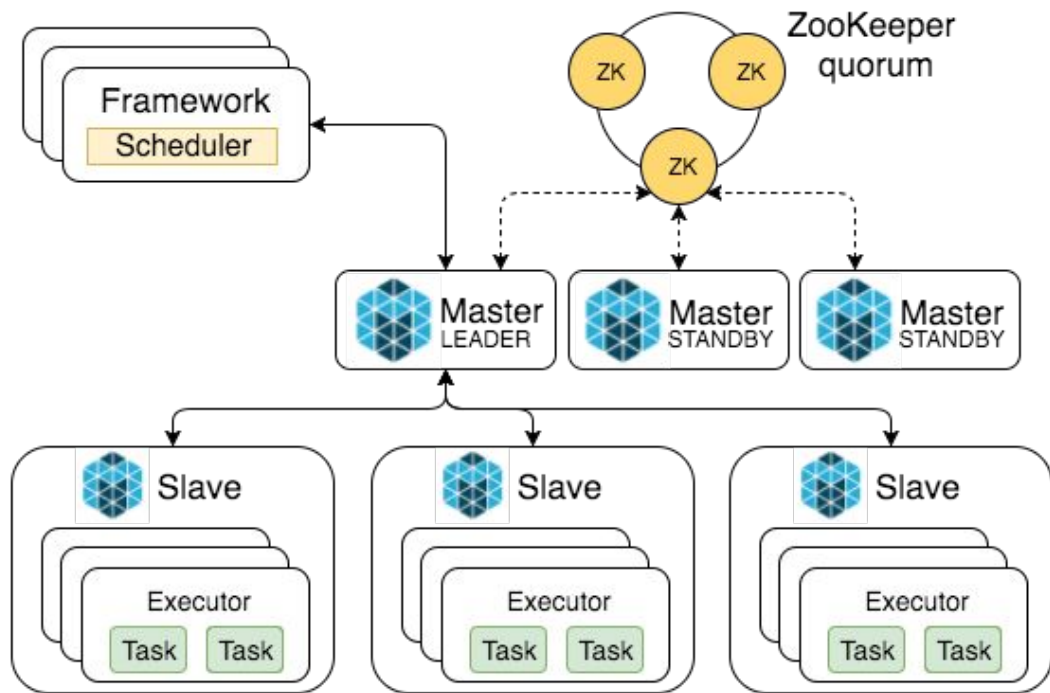
Cluster Resource Managers: Requirements

- Efficiency
 - efficient sharing of resources across applications
 - utilization of cluster resources in the most optimal manner
- Flexibility
 - support of wide array of current and future frameworks
 - dealing with hardware heterogeneity
 - support of resource requests of different types
- Scalability
 - scaling to clusters of dozens of thousands of nodes
 - scheduling system's response times must remain acceptable while increasing number of machines and applications
- Robustness
 - fault-tolerant guarantees for the system and applications
 - high availability of central scheduler component

Cluster Manager Architectures

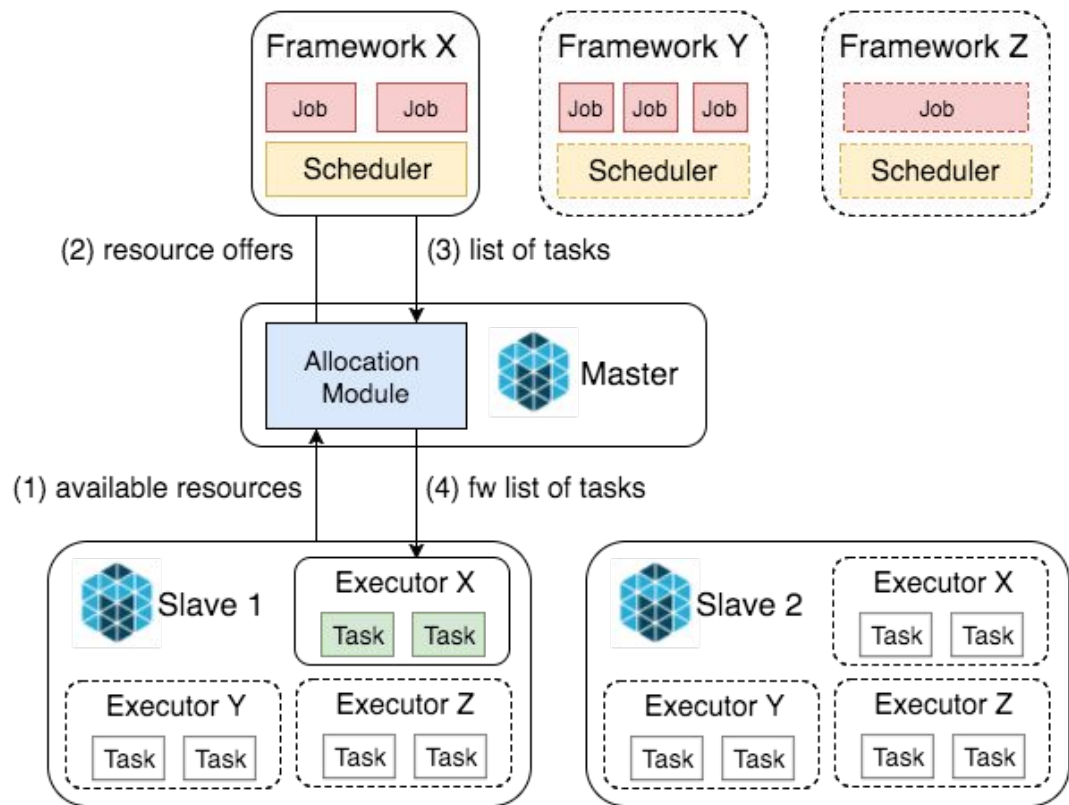


Mesos Architecture



- **Master**
 - a mediator between slave resources and frameworks
 - enables fine-grained sharing of resources by making resource offers
- **Slave**
 - manages resources on physical node and runs executors
- **Framework**
 - application that solves a specific use case
 - *Scheduler* negotiates with master and handles resource offers
 - *Executors* consume resources and run tasks on slaves

Two-Level Scheduling



- Slave nodes report to Master amount of available resources
- Allocation module starts offering resources to frameworks
- Framework receives offers
 - if resources do not satisfy its needs - rejects the offer
 - if resources satisfy its demands - creates list of *tasks* and sends to master
- Master verifies tasks and forwards to executor (and launches the executor if it's not running)

Resource offer

```
id: { value: "0cb2328a-61c2-4316-91ef-cbbb6ebbf504-O1" }
framework_id: { value: "0cb2328a-61c2-4316-91ef-cbbb6ebbf504-0001" }
slave_id: { value: "0cb2328a-61c2-4316-91ef-cbbb6ebbf504-S0" }
hostname: "mesos-slave"
resources { name: "cpus", type: SCALAR, scalar { value: 6.0 }, role: "*" }
resources { name: "mem", type: SCALAR, scalar { value: 6762.0 }, role: "*" }
resources { name: "disk", type: SCALAR, scalar { value: 13483.0 }, role: "*" }
resources { name: "ports", type: RANGES, ranges { range { begin: 31000, end: 32000 } }, role: "*" }
url {
  scheme: "http"
  address {
    hostname: "mesos-slave"
    ip: "172.18.0.5"
    port: 5151
  }
  path: "/slave(1)"
}
```

Framework Scheduler

```
class SomeMesosScheduler extends Scheduler {  
  override def resourceOffers(driver: SchedulerDriver, offers: List[Offer]): Unit = {  
    for(offer <- offers){  
      stateLock.synchronized {  
        if(isOfferValid(offer)){  
          val executorInfo = buildExecutorInfo(driver, "Executor A")  
          //amount of tasks is calculated to fully use resources from the offer  
          val tasks = buildTasks(offer, executorInfo)  
          driver.launchTasks(List(offer.getId), tasks)  
        } else {  
          driver.declineOffer(offer.getId)  
        }  
      }  
    }  
  }  
  //rest of the methods implementations go here  
}
```

Dominant Resource Fairness (DRF)

- *Dominant resource*
 - a resource of specific type (cpu, ram, etc.) which is most demanded by a framework among other resources it needs
 - the resource is identified as a share of the total cluster resources of the same type
- *Dominant share*
 - a share of dominant resource allocated to a framework in the cluster
- Example:
 - Cluster total: 9 CPU & 18 GB RAM
 - Framework A tasks need < 3 CPU, 1 GB > (or < 33% CPU, 5% RAM >)
 - Framework B tasks need < 1 CPU, 4 GB > (or < 11% CPU, 22% RAM >)
- DRF algorithm computes frameworks' dominant shares and tries to maximize the smallest dominant share in the system

DRF Demo

- 3 frameworks with < 8% CPU, 7.5% RAM > demand each

Name	Active Tasks	CPUs	Mem	Disk	Max Share
Framework C	4	2	2.0 GB	0 B	33.333%
Framework B	4	2	2.0 GB	0 B	33.333%
Framework A	4	2	2.0 GB	0 B	33.333%

Name	Active Tasks	CPUs	Mem	Disk	Max Share
Framework C	4	2	2.0 GB	0 B	33.333%
Framework B	5	2.5	2.5 GB	0 B	41.667%
Framework A	3	1.5	1.5 GB	0 B	25%

- Framework A < 33% CPU, 15% RAM >, Framework B < 16% CPU, 30% RAM >)

Name	Active Tasks	CPUs	Mem	Disk	Max Share
Framework A	2	4	2.0 GB	0 B	66.667%
Framework B	2	2	3.9 GB	0 B	59.154%

- Framework A < 33% CPU, 15% RAM >, Framework B < 16% CPU, 36% RAM >)

Name	Active Tasks	CPUs	Mem	Disk	Max Share
Framework A	2	4	2.0 GB	0 B	66.667%
Framework B	1	1	2.4 GB	0 B	36.971%

Name	Active Tasks	CPUs	Mem	Disk	Max Share
Framework A	1	2	1000 MB	0 B	33.333%
Framework B	2	2	4.9 GB	0 B	73.943%

DRF properties

- Sharing incentive
 - Each user should be better off sharing the cluster, than exclusively using her own partition of the cluster. Consider a cluster with identical nodes and n users. Then a user should not be able to allocate more tasks in a cluster partition consisting of $1/n$ of all resources.
- Strategy-proofness
 - Users should not be able to benefit by lying about their resource demands. This provides incentive compatibility, as a user cannot improve her allocation by lying.
- Envy-freeness
 - A user should not prefer the allocation of another user. This property embodies the notion of fairness.
- Pareto efficiency
 - It should not be possible to increase the allocation of a user without decreasing the allocation of at least another user. This property is important as it leads to maximizing system utilization subject to satisfying the other properties.

Resource Reservation

- Goals:
 - allocate all single slave resources to one type of framework
 - divide cluster between several framework types or organisations
 - framework groups prioritization and guaranteed allocation
- Static reservation
 - slave node is configured on start (cannot be reserved for another role or unreserved)
`--resources="cpus:4;mem:2048;cpus(spark):8;mem(spark):4096"`
- Dynamic reservation
 - resources are reserved/unreserved within a respond to resource offer
`Offer::Operation::Reserve`
 - [MESOS-2018](#)
- Extras:
 - persistent volumes
 - multiple disk resources

Resource Isolation

- Goals:
 - running tasks isolation and capping of runtime resources
 - programmatic control over task resources
 - use images to allow different environments
- Docker containerizer
 - executed tasks are docker containers (e.g. microservices packed in Docker)
- Mesos containerizer (default)
 - Mesos-native (no dependencies on other technologies)
 - provides fine-grained controls (cgroups/namespaces)
 - provides disk usage limits controls
- Composing
 - allows using multiple containerizers together
 - the first containerizer supporting task configuration will be used to launch it

Ubiquitous frameworks: Marathon

The screenshot displays the Marathon web interface. At the top, there's a navigation bar with the Marathon logo, 'Apps' (selected), 'Deployments', 'About', and 'Docs'. Below this, a breadcrumb shows 'Apps > /spark-log-uploader'. The main section shows the app '/spark-log-uploader' in a 'Running' state. There are buttons for 'Suspend', 'Scale', 'Restart App', and 'Destroy App'. Below these are tabs for 'Tasks' and 'Configuration' (selected). The 'Configuration' tab shows the 'Current Version' with a 'Refresh' button. A table lists various configuration parameters:

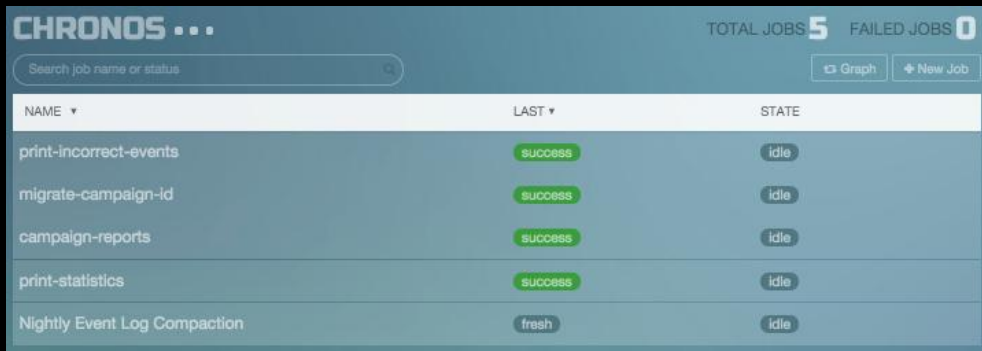
Command	spark-submit --class io.datastrophic.S3CassandraLoaderJob --executor-memory 4G --driver-memory 1G /datastrophic/jars/spark-uploader.jar --bucket events --prefix 2015/"'.gz --host 172.31.0.0 --keyspace events --table raw_events
Constraints	Unspecified
Container	Unspecified
CPUs	1
Environment	Unspecified
Executor	Unspecified
Instances	1
Memory	512 MB
Disk Space	0 MB
Ports	10002
Backoff Factor	1.15
Backoff	1 seconds
Max Launch Delay	3600 seconds
URIs	Unspecified
Version	2015-09-08T20:13:09.370Z

- distributed init.d
- long running tasks execution
- HA mode with ZooKeeper
- Docker executor
- REST API

Marathon: launching Chronos in Docker

```
curl -XPOST 'http://marathon:8080/v2/apps' -H 'Content-Type: application/json' -d '{
  "id": "chronos",
  "container": {
    "type": "DOCKER",
    "docker": {
      "network": "HOST",
      "image": "datastrophic/chronos:mesos-0.27.1-chronos-2.5",
      "parameters": [
        { "key": "env", "value": "CHRONOS_HTTP_PORT=4400" },
        { "key": "env", "value": "CHRONOS_MASTER=zk://zookeeper:2181/mesos" },
        { "key": "env", "value": "CHRONOS_ZK_HOSTS=zookeeper:2181" }
      ]
    }
  },
  "ports": [ 4400 ],
  "cpus": 1,
  "mem": 512,
  "instances": 1
}'
```

Ubiquitous frameworks: Chronos



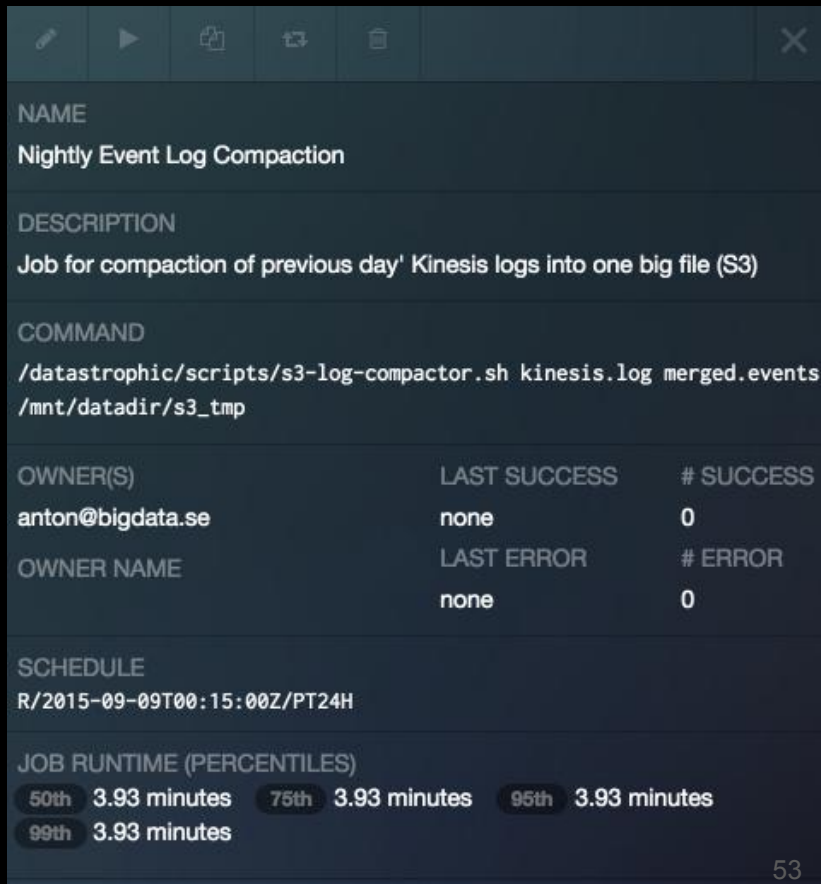
CHRONOS ... TOTAL JOBS 5 FAILED JOBS 0

Search job name or status

Graph New Job

NAME	LAST	STATE
print-incorrect-events	success	idle
migrate-campaign-id	success	idle
campaign-reports	success	idle
print-statistics	success	idle
Nightly Event Log Compaction	fresh	idle

- distributed cron
- HA mode with ZooKeeper
- supports graphs of jobs
- sensitive to network failures



NAME
Nightly Event Log Compaction

DESCRIPTION
Job for compaction of previous day's Kinesis logs into one big file (S3)

COMMAND
`/datastrophic/scripts/s3-log-compactor.sh kinesis.log merged.events /mnt/datadir/s3_tmp`

OWNER(S)	LAST SUCCESS	# SUCCESS
anton@bigdata.se	none	0
OWNER NAME	LAST ERROR	# ERROR
	none	0

SCHEDULE
R/2015-09-09T00:15:00Z/PT24H

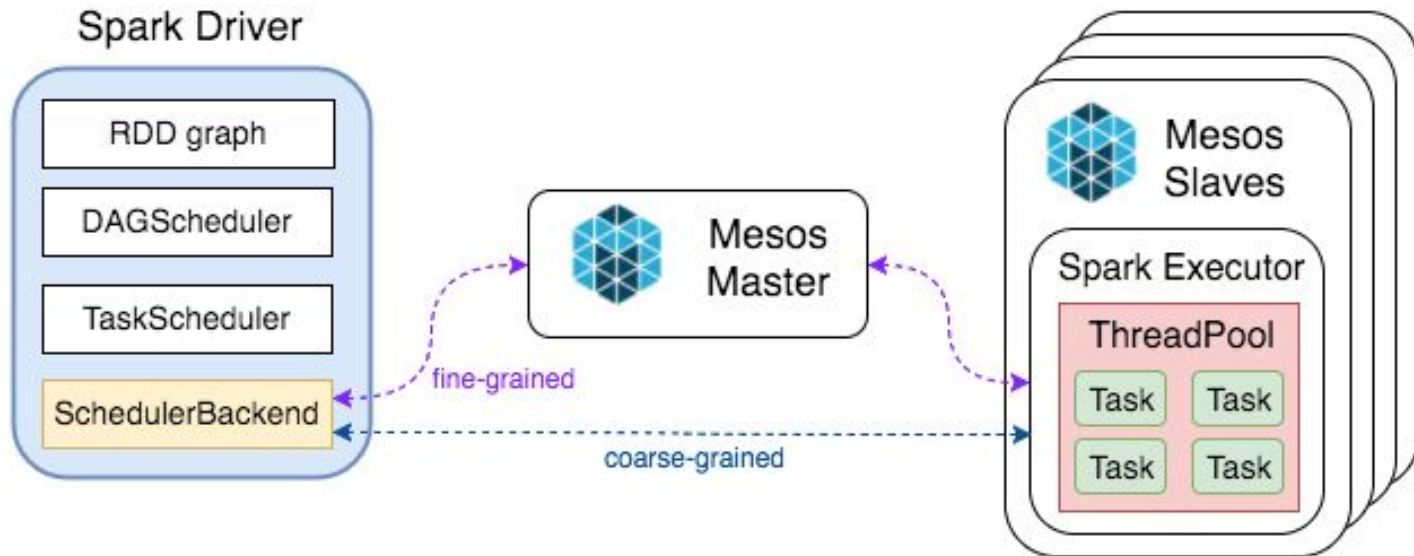
JOB RUNTIME (PERCENTILES)

50th	3.93 minutes	75th	3.93 minutes	95th	3.93 minutes
99th	3.93 minutes				

More Mesos frameworks

- Spark
- Hadoop
- Cassandra
- Kafka
- Myriad: YARN on Mesos
- Storm
- Samza

Spark on Mesos



- **Coarse-grained mode(default)**
 - Spark Executor is launched one per Slave and acquires all available cores in cluster
 - Tasks are scheduled by Spark relying on its RPC mechanism (Akka)
- **Fine-grained mode**
 - Spark Executor is launched one per Slave with minimal resources needed (1 core)
 - Spark tasks are executed as Mesos tasks and use Mesos semantics

Spark on Mesos

- Coarse-grained mode

```
/opt/spark/bin/spark-submit \  
  --class io.datastrophic.demo.SparkJob \  
  --master mesos://zk://zookeeper:2181/mesos \  
  --conf "spark.cores.max=10" \  
  /opt/jobs/spark-jobs-assembly.jar
```

- Fine-grained mode

```
/opt/spark/bin/spark-submit \  
  --class io.datastrophic.demo.SparkJob \  
  --master mesos://zk://zookeeper:2181/mesos \  
  --conf "spark.mesos.coarse=false" \  
  /opt/jobs/spark-jobs-assembly.jar
```


Spark on Mesos vs. YARN

- Mesos (coarse-grained)

```
/opt/spark/bin/spark-submit \  
  --class io.datastrophic.demo.SparkJob \  
  --master mesos://zk://zookeeper:2181/mesos \  
  --conf "spark.cores.max=100" \  
  /opt/jobs/spark-jobs-assembly.jar
```

- YARN

```
/opt/spark/bin/spark-submit \  
  --class io.datastrophic.demo.SparkJob \  
  --master yarn \  
  --num-executors 25 \  
  --executor-cores 4 \  
  /opt/jobs/spark-jobs-assembly.jar
```

Running Spark via Marathon

```
curl -XPOST 'http://marathon:8080/v2/apps' -H 'Content-Type: application/json' -d '{  
  "cmd": "/opt/spark/bin/spark-submit  
  --class io.datastrophic.demo.SparkJob  
  --master mesos://zk://zookeeper:2181/mesos  
  --deploy-mode client  
  /opt/jobs/spark-jobs-assembly.jar",  
  "id": "spark-pi",  
  "cpus": 1,  
  "mem": 1024,  
  "instances": 1  
}'
```

Running Spark via Chronos

```
curl -L -H 'Content-Type: application/json' -X POST http://mesos:4400/scheduler/iso8601 -d '{  
  "name": "Scheduled Spark Submit Job",  
  "/opt/spark/bin/spark-submit  
  --class io.datastrophic.demo.SparkJob  
  --master mesos://zk://zookeeper:2181/mesos  
  /opt/jobs/spark-jobs-assembly.jar",  
  "shell": true,  
  "async": false,  
  "cpus": 0.1,  
  "disk": 256,  
  "mem": 1024,  
  "owner": "anton@datastrophic.io",  
  "description": "Spark Job executed every 3 minutes",  
  "schedule": "R/2016-03-14T12:35:00.000Z/PT3M"  
}'
```

Spark deployment strategies

- Binaries distribution
 - every node in the cluster must have Spark libraries installed in the same locations
 - pros: easy to start with
 - cons: hard to upgrade, hard to have several Spark versions simultaneously
- Edge nodes
 - use nodes with specific environment setup which are reachable from Mesos cluster and keep Spark executor jars in accessible location like S3, HTTP or HDFS
 - pros: easy to use multiple Spark versions, minimal dependencies on Mesos
 - cons: hard to maintain in case of multi-tenancy
- Dockerized environment
 - Instead of edge nodes use Docker containers with environment configured for specific needs (hosts still need to be reachable from Mesos cluster) and use Docker Spark executor
 - pros: highly isolated environments for specific needs, could be upgraded independently, zero impact on cluster nodes
 - cons: could be hard to properly setup and configure

Mesos Framework Walkthrough

- **Throttler**
 - a demo framework for load testing Cassandra
 - load intensity is controlled by parameters: total queries, queries per task and parallelism (how many Mesos tasks to run in parallel)
- Goals
 - take a look at working (simple) Mesos application
 - see how Scheduler, Executor and framework launcher could be implemented
- Sources:
 - source code and dockerized Mesos cluster configuration are available at [github/datastrophic/mesos-workshop](https://github.com/datastrophic/mesos-workshop)
 - all the examples (and even more) available as well

Questions

[@antonkirillov](#)

[datastrophic.io](#)