

Ingesting streaming data into Graph Database

Guido Schmutz

San Francisco, 22.3.2018



Analytics and
Data Summit



@gschmutz



guidoschmutz.wordpress.com

BASEL ▪ BERN ▪ BRUGG ▪ DÜSSELDORF ▪ FRANKFURT A.M. ▪ FREIBURG I.BR. ▪ GENF
HAMBURG ▪ KOPENHAGEN ▪ LAUSANNE ▪ MÜNCHEN ▪ STUTTGART ▪ WIEN ▪ ZÜRICH

trivadis
makes **IT** easier. A logo for Trivadis featuring the company name in a bold, red, sans-serif font. Below it, the tagline "makes IT easier." is written in a smaller, red, lowercase font. To the right of the tagline are three small red dots.

■ Guido Schmutz

Working at Trivadis for more than 21 years

Oracle ACE Director for Fusion Middleware and SOA



Consultant, Trainer Software Architect for Java, Oracle, SOA and Big Data / Fast Data



Head of Trivadis Architecture Board

Technology Manager @ Trivadis

More than 30 years of software development experience

Contact: guido.schmutz@trivadis.com

Blog: <http://guidoschmutz.wordpress.com>

Slideshare: <http://www.slideshare.net/gschmutz>

Twitter: [@gschmutz](https://twitter.com/gschmutz)

gschmutz 24.08.17 on April 18, 2017
Tags: flink (1), kafka (59), kafka-connect (4), kafka-streams (17), spark-streaming (31), storm (39), streams (41)

Last week in Stream Processing & Analytics – 18.4.2017

This is the 62nd edition of my blog series blog series around Stream Processing and Analytics!

Every week I'm also updating the following two lists with the presentations/videos of the current week:

- [Presentations from Slideshare](#)
- [Videos from YouTube](#)

As usual, find below the new blog articles, presentations, videos and software releases from last week:

News and Blog Posts

General

- [Multi Master Replication For Geo-Distributed Data: It's more than you think by Ellen Friedman](#)
- [Understanding Indicators of Attack \(IOAs\): The Power of Event Stream Processing in CrowdStrike Falcon by Dan Brown](#)
- [Stream processing and messaging systems for the IoT age by Ben Lorica](#)

Apache Kafka / Kafka Streams / Confluent Platform

- [Creating a Data Pipeline with Kafka Connect API – from Architecture to Operations by Alexandra Wang](#)
- [Streaming Spring Boot Application Logs to ELK Stack – Part 1 by kaadayamuthu](#)
- [Streaming Spring Boot Application Logs to Apache Kafka – ELK/Kafka Stack – Part 2 by kaadayamuthu](#)

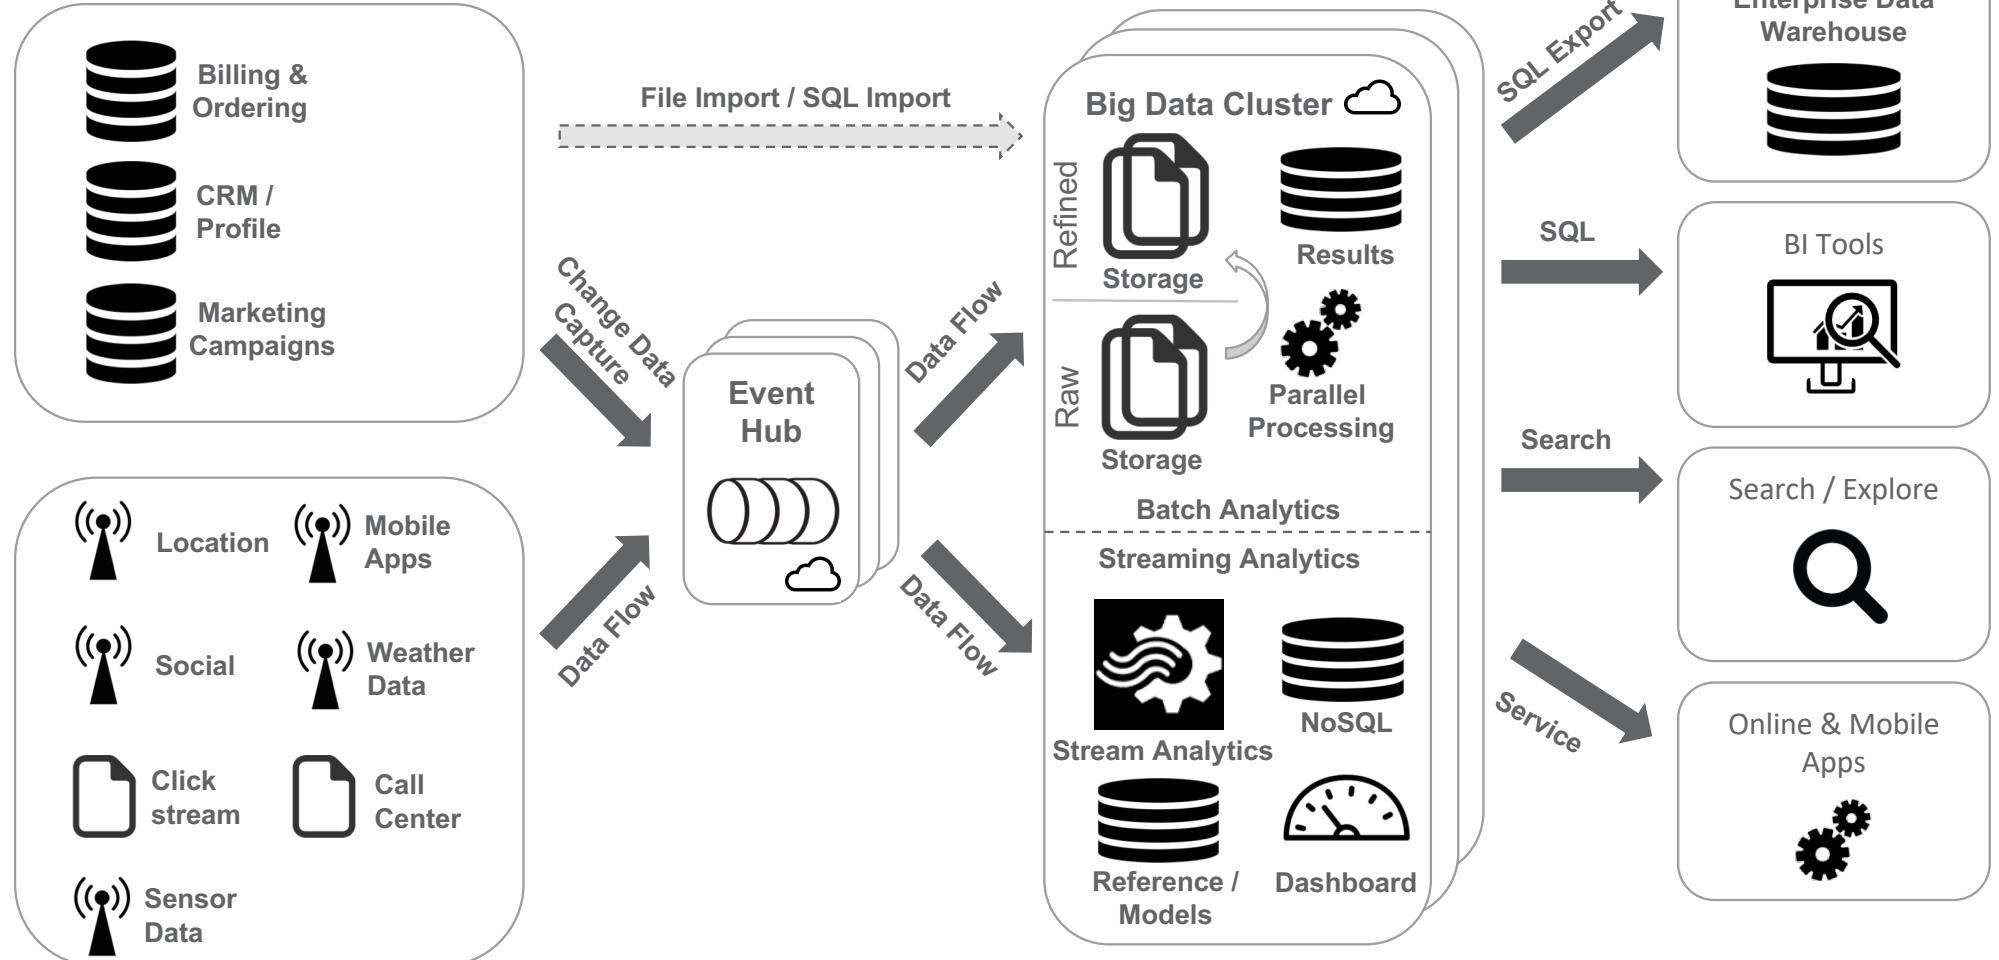


■ Agenda

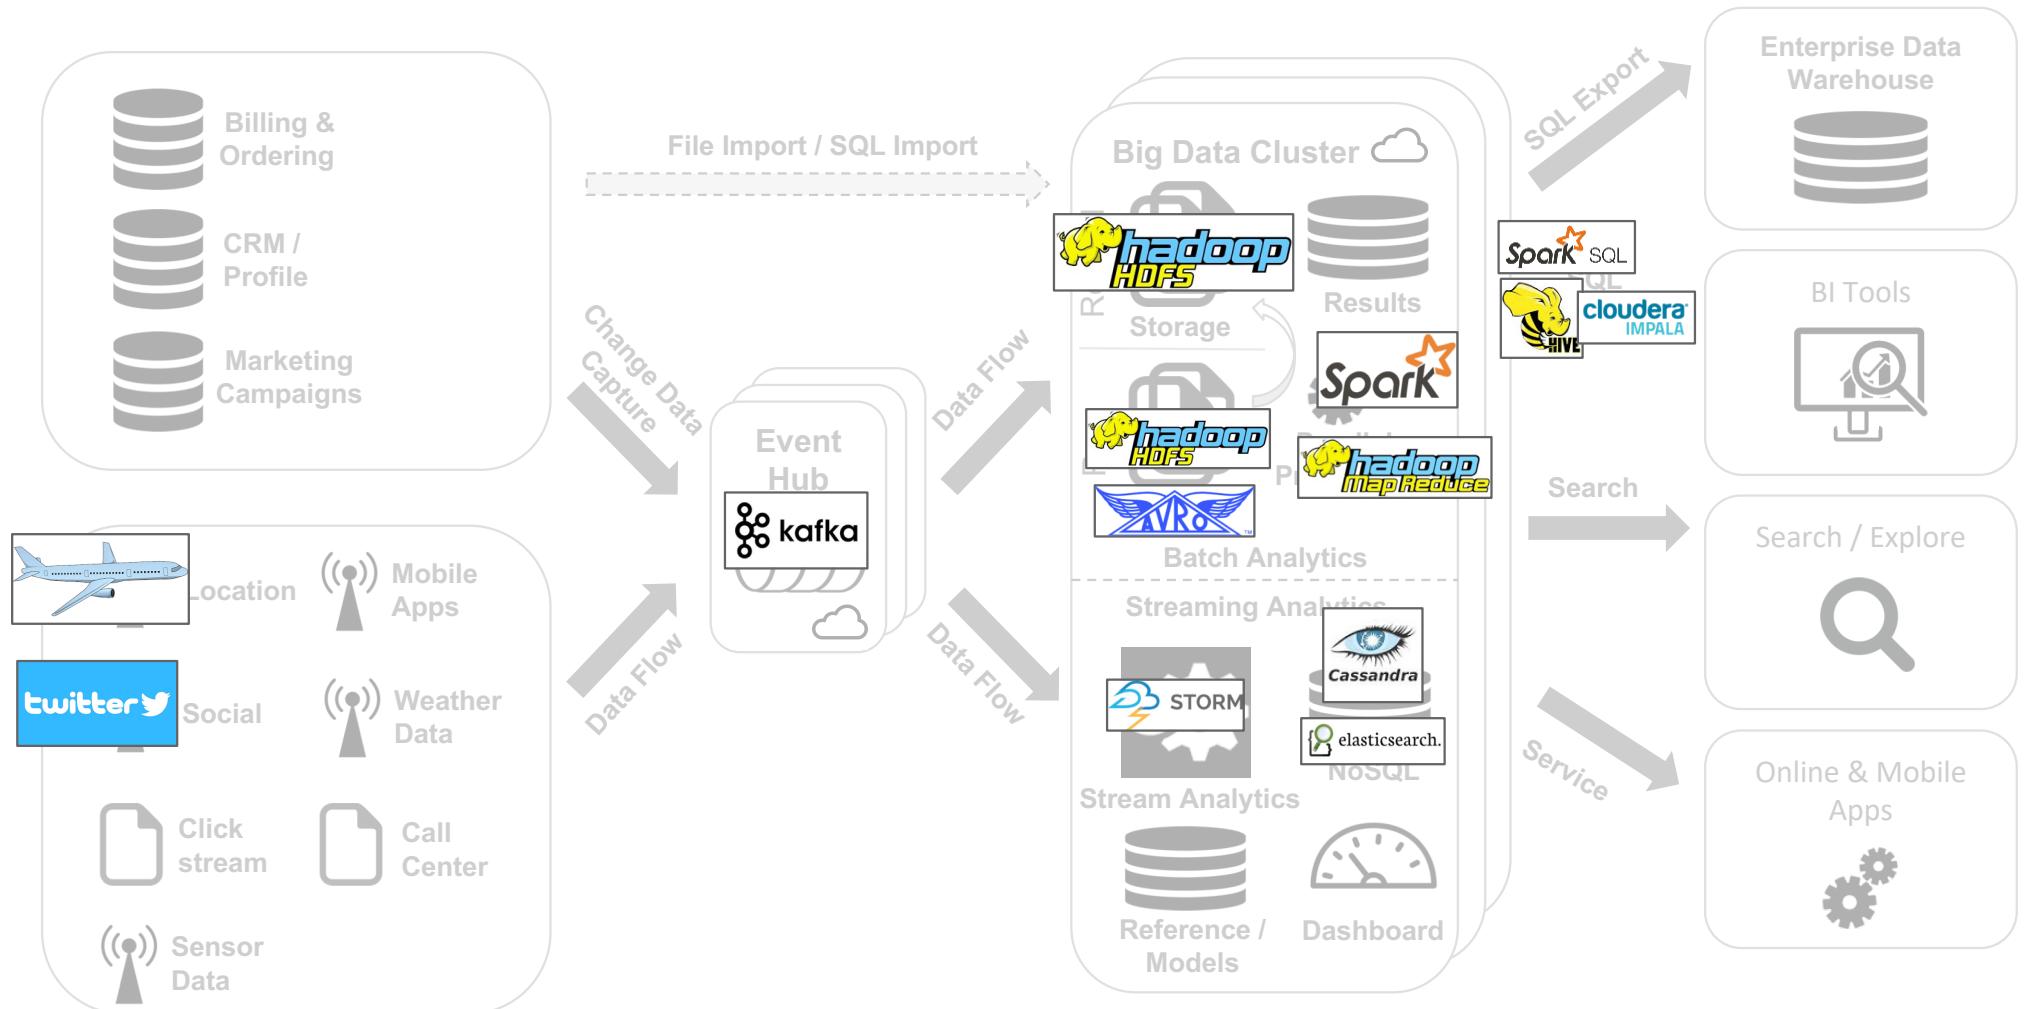
1. Introduction
2. DataStax (DSE) Graph
3. Use-Case 1 - Batch Data Ingestion of Social Graph
4. Use-Case 2 - Streaming Ingestion of Social Graph
5. Use-Case 3 - Streaming Ingestion into RDF Graph
6. Summary

Introduction

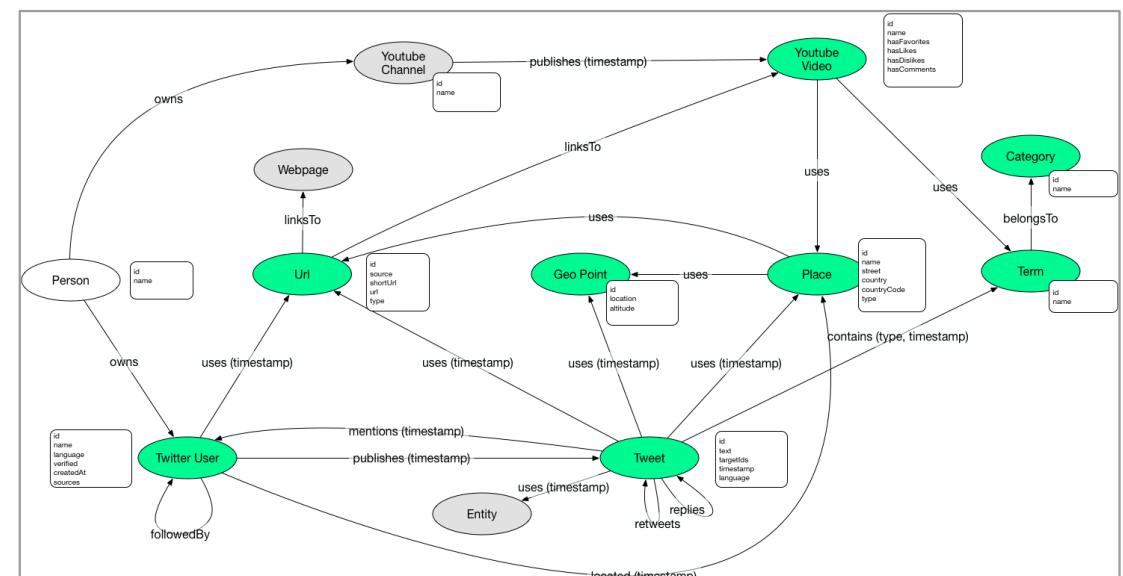
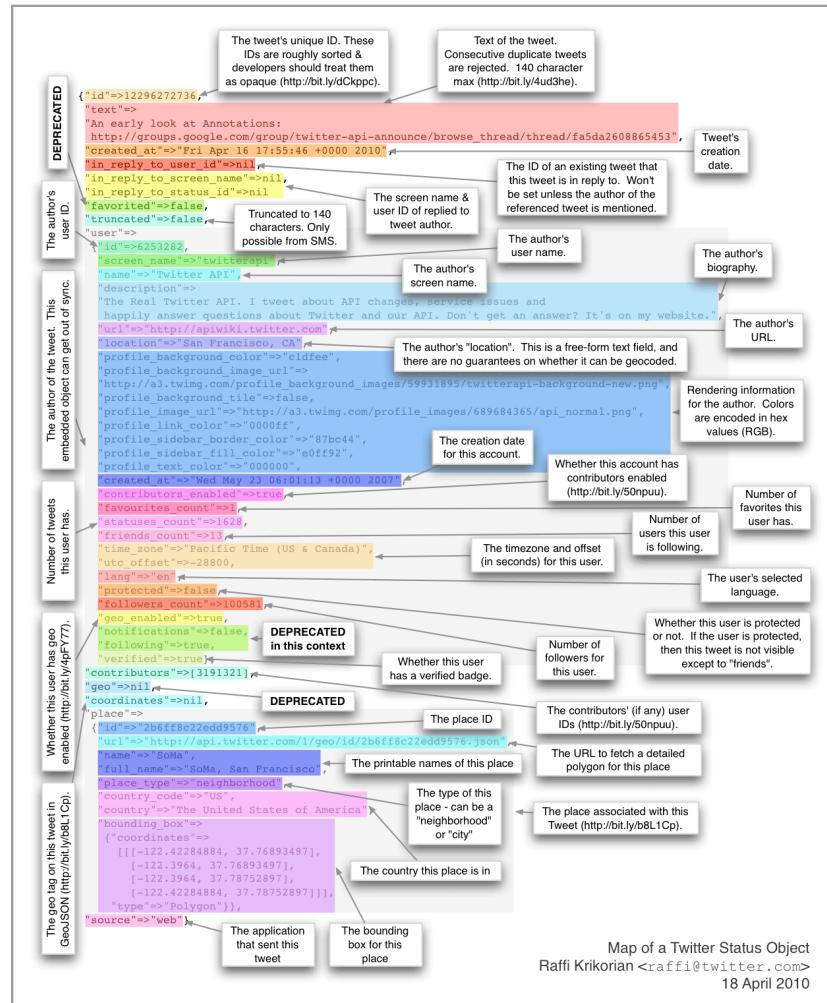
Customer Context – Big Data Lab



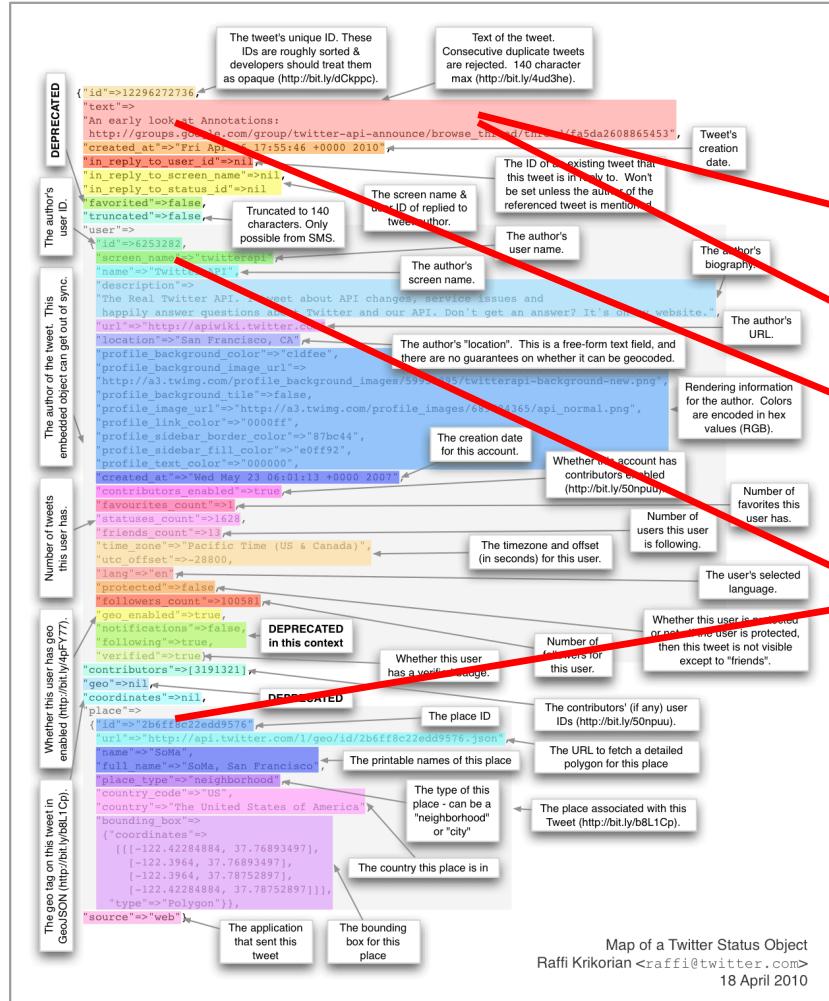
Customer Context – Big Data Lab



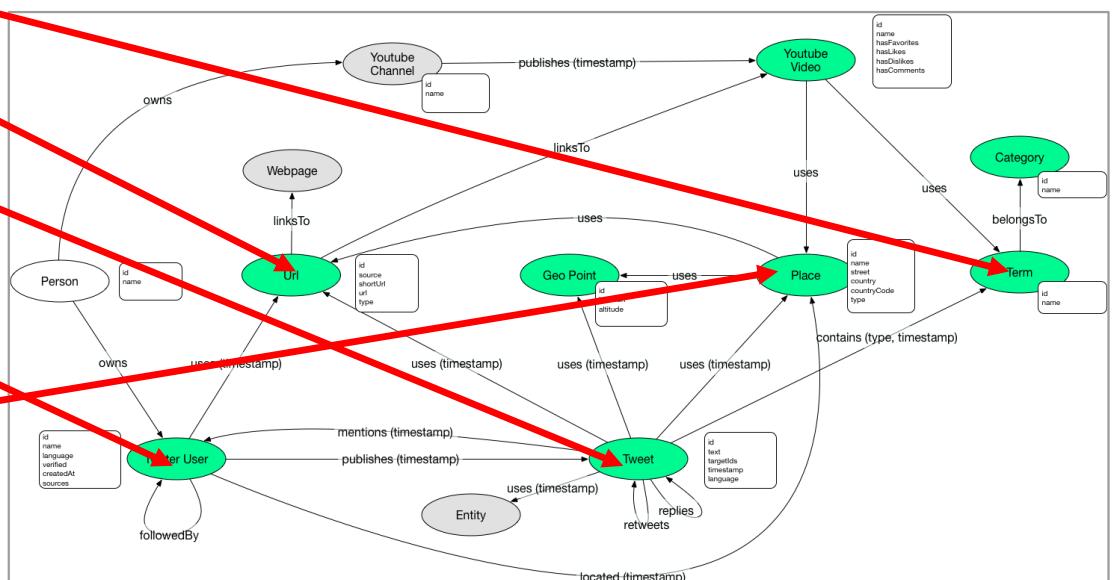
Map of a Twitter Status object



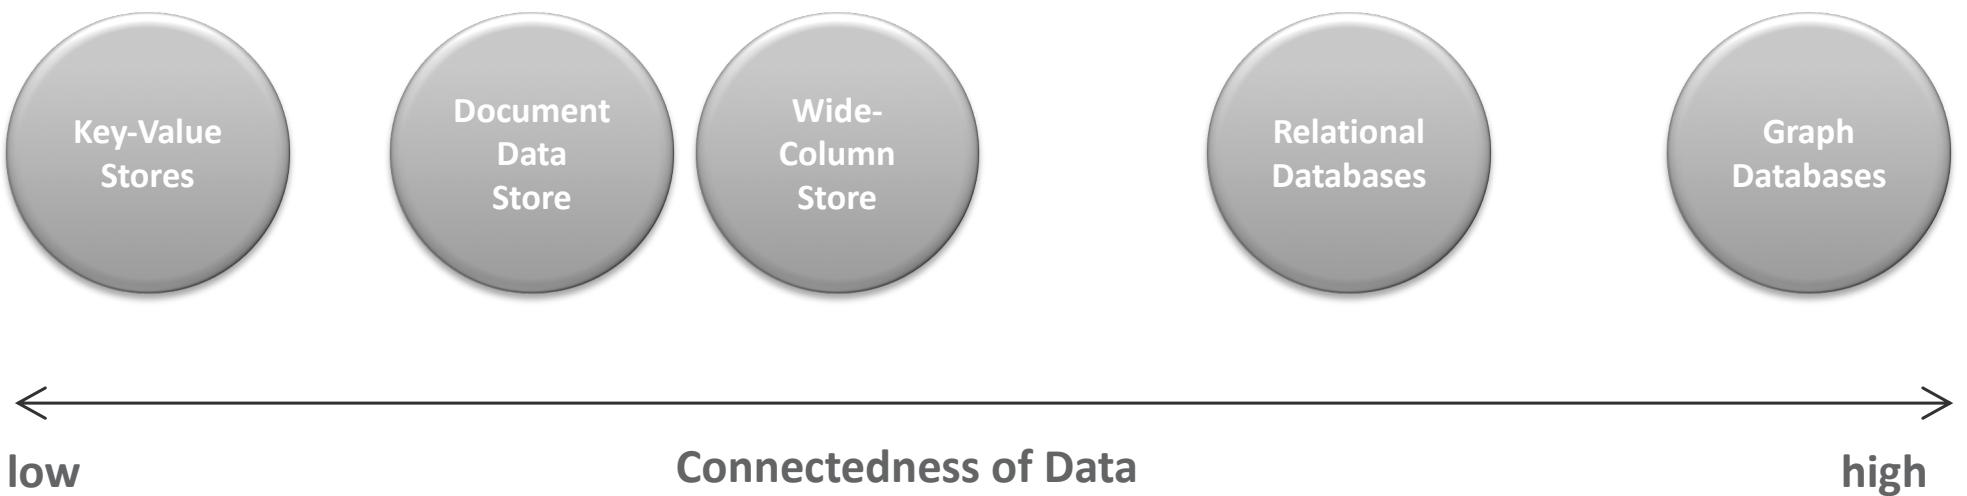
Map of a Twitter Status object



Map of a Twitter Status Object
Raffi Krikorian <raffi@twitter.com>
18 April 2010



■ Know your domain and choose the right data store!



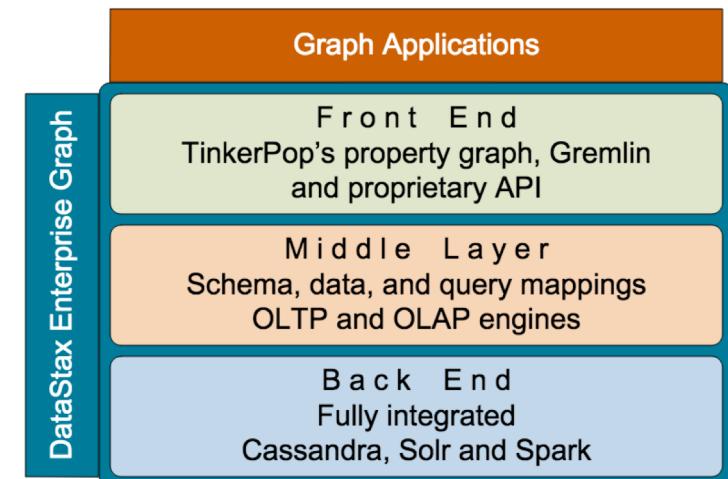
DataStax (DSE) Graph

■ DataStax Enterprise Edition – DSE Graph

A scale-out property graph database

Store & find relationships in data fast and easy
in large graphs.

- Adopts **Apache TinkerPop** standards for data and traversal
- Uses **Apache Cassandra** for scalable storage and retrieval
- Leverages **Apache Solr** for full-text search and indexing
- Integrates **Apache Spark** for fast analytic traversal
- Supports data security



Images: DataStax

■ Lab cluster setup

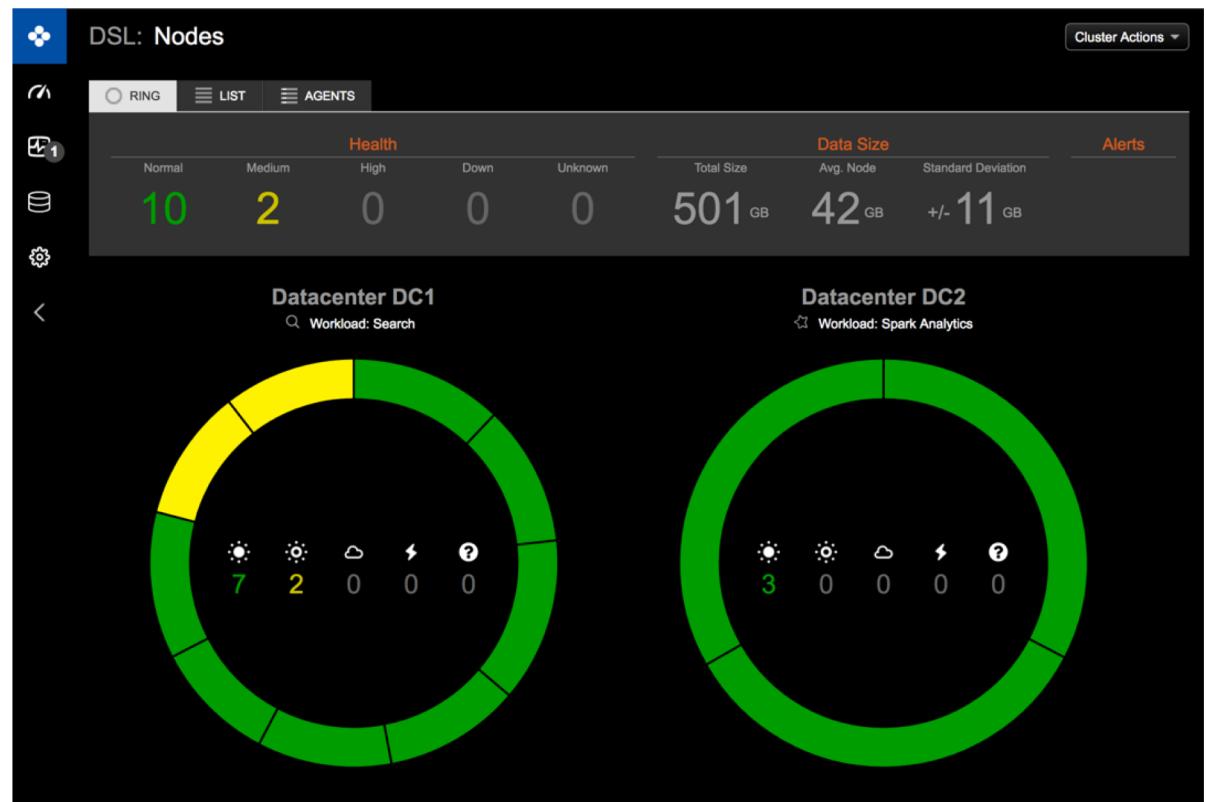
2 local data center for workload segregation

DC1

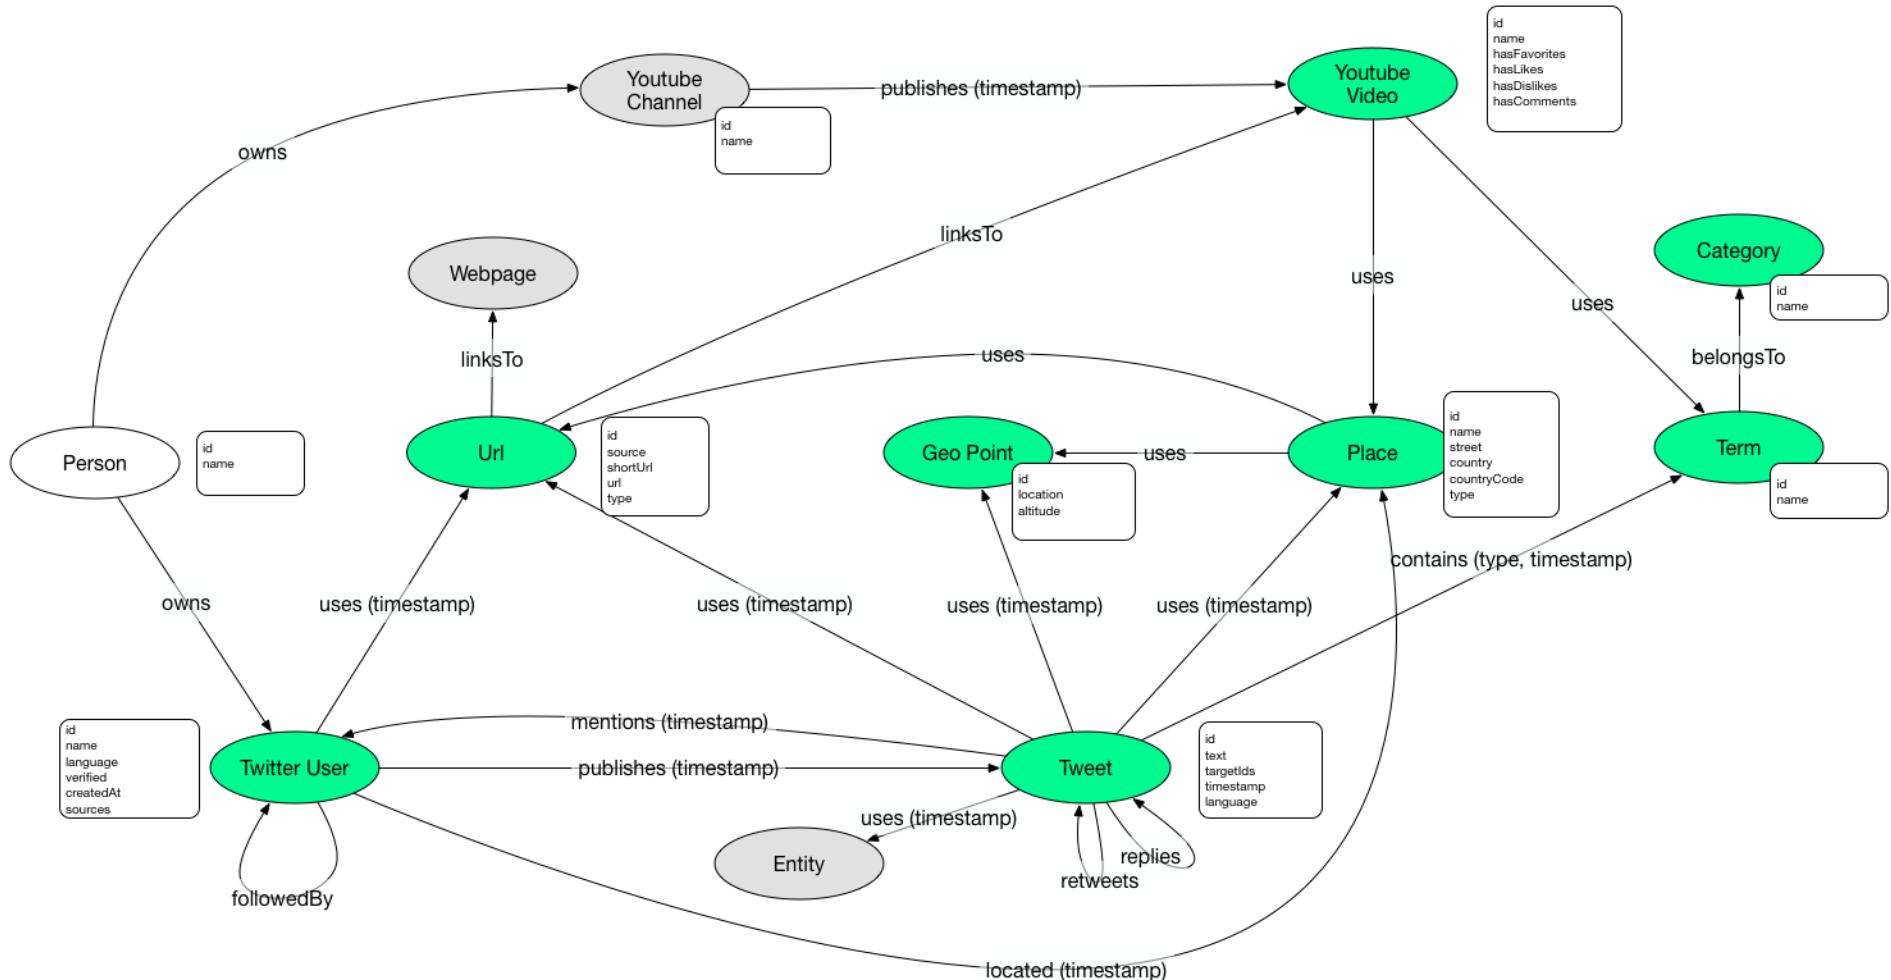
- 9 nodes
- OLTP workload incl. Search

DC2

- 3 nodes
- OLAP workload with Spark



Social Network Graph Model



Creating the Graph Schema (I)

Define types of properties to be used with Vertices and Edges

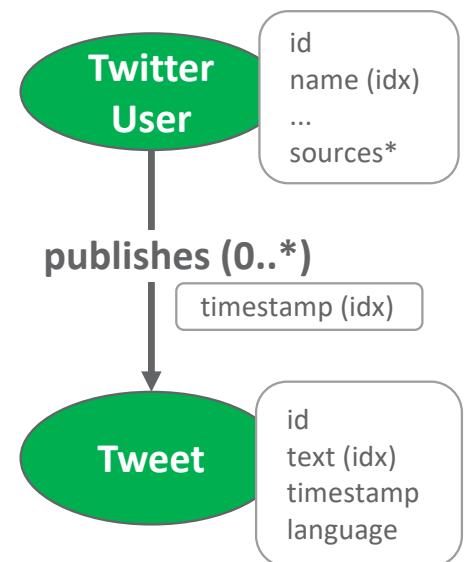
```
schema.propertyKey("id").Long().single().create();
schema.propertyKey("name").Text().create();
schema.propertyKey("sources").Text().multiple().create();
```

Define types of vertices to be used

```
schema.vertexLabel("twitterUser").
    properties("id","name","language","verified","createdAt","sources").create();
schema.vertexLabel("tweet".
    properties("id","text","targetIds","timestamp","language").create()
```

Define types of edges to be used

```
schema.edgeLabel("publishes").properties("timestamp")
    multiple().properties("rating").
    connection("twitterUser","tweet").create();
```



■ Creating the Graph Schema (II)

Define global materialized index on vertex (good for high selectivity)

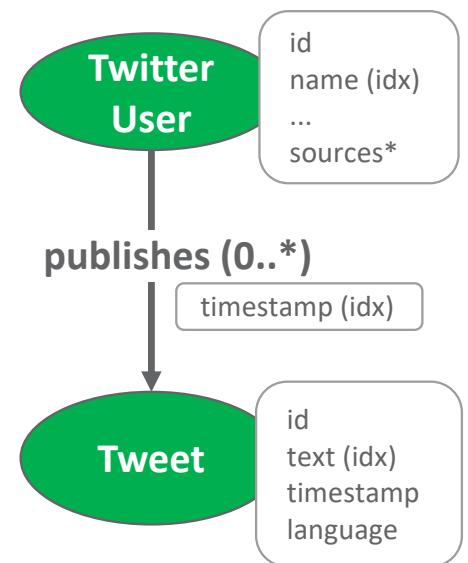
```
schema.vertexLabel("twitterUser") .  
    index('twitterUserNameIdx') .  
    materialized() .by('name') .ifNotExists() .add()
```

Defining global full-text search index (SolR) on text of Tweet

```
schema.vertexLabel("tweet") .  
    index("search") .  
    search() .by("text") .asText() .ifNotExists() .add()
```

Define vertex-centric index on edge

```
schema.vertexLabel("twitterUser") .  
    index("usesAtTimestamp") .  
    outE("publishes") .by("timestamp") .ifNotExists() .add()
```



■ Creating Vertex Instances

Adding a new **TwitterUser** vertex

```
Vertex u = graph.addVertex("twitterUser");
u.property("id", "1323232");
u.property("name", "realDonaldTrump");
u.property("language", "en");
u.property("sources", "tweet");
```

userId: 1323232
name: realDonaldTrump
language: en
sources: tweet

Twitter
User

Adding a new **Tweet** vertex

```
Vertex m = graph.addVertex(label, "tweet",
                            "id", "1213232323",
                            "text", "I called President Putin of
Russia to congratulate him on his election ...",
                            "timestamp", "2018-03-21",
                            "language", "en");
```

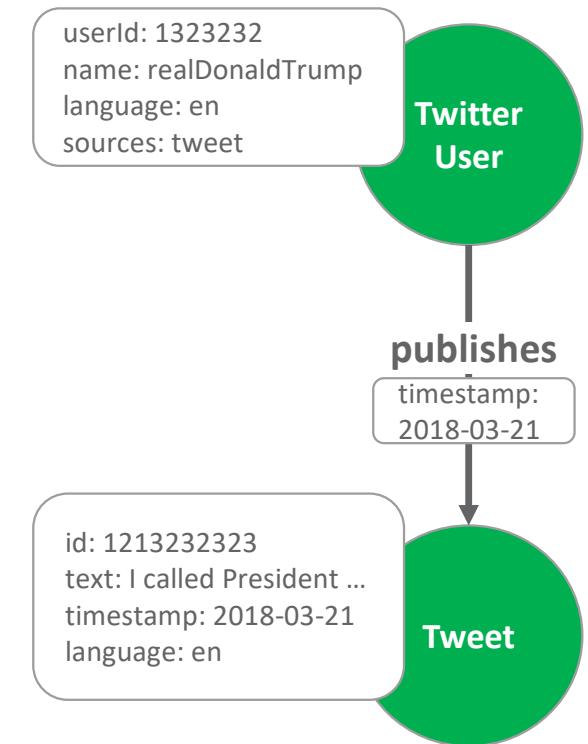
id: 1213232323
text: I called President ...
timestamp: 2018-03-21
language: en

Tweet

■ Creating Edges between Vertices

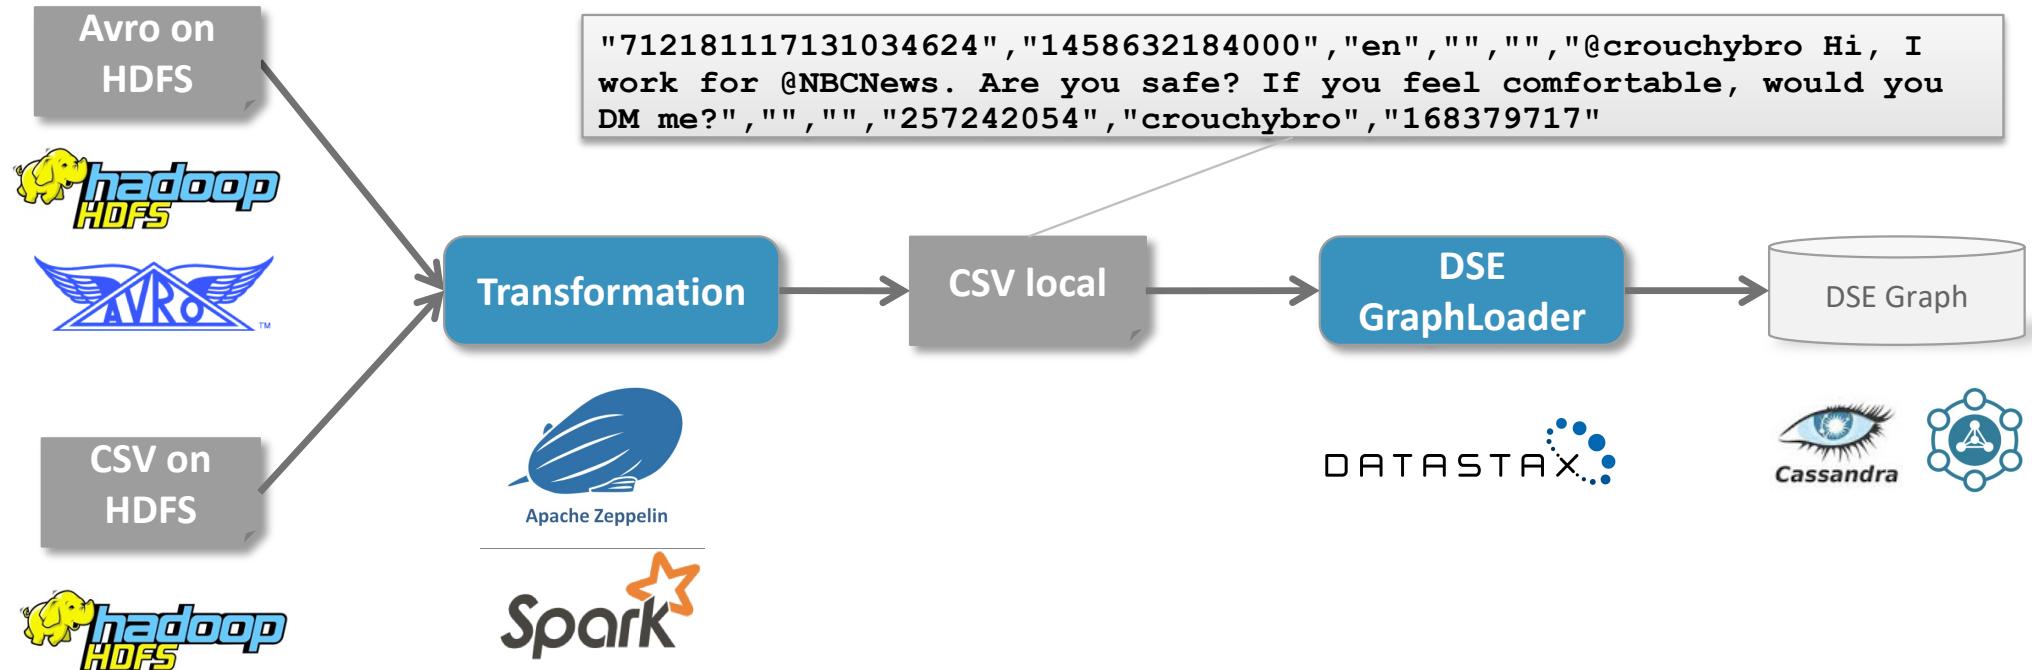
Adding an edge between the **TwitterUser** and the **Tweet** vertex

```
Vertex u = graph.addVertex("twitterUser", ...);  
Vertex t = graph.addVertex("tweet", ...);  
  
Edge e = u.addEdge("publishes", t);  
e.property("timestamp", "2018-03-21");
```



Use Case 1 – Batch Data Ingestion of Social Graph

Batch Data Ingestion of Social Graph



■ Extracting and Preparing Data through Zeppelin

The screenshot shows a Zeppelin notebook interface with a single job titled "Tweet". The job consists of two cells:

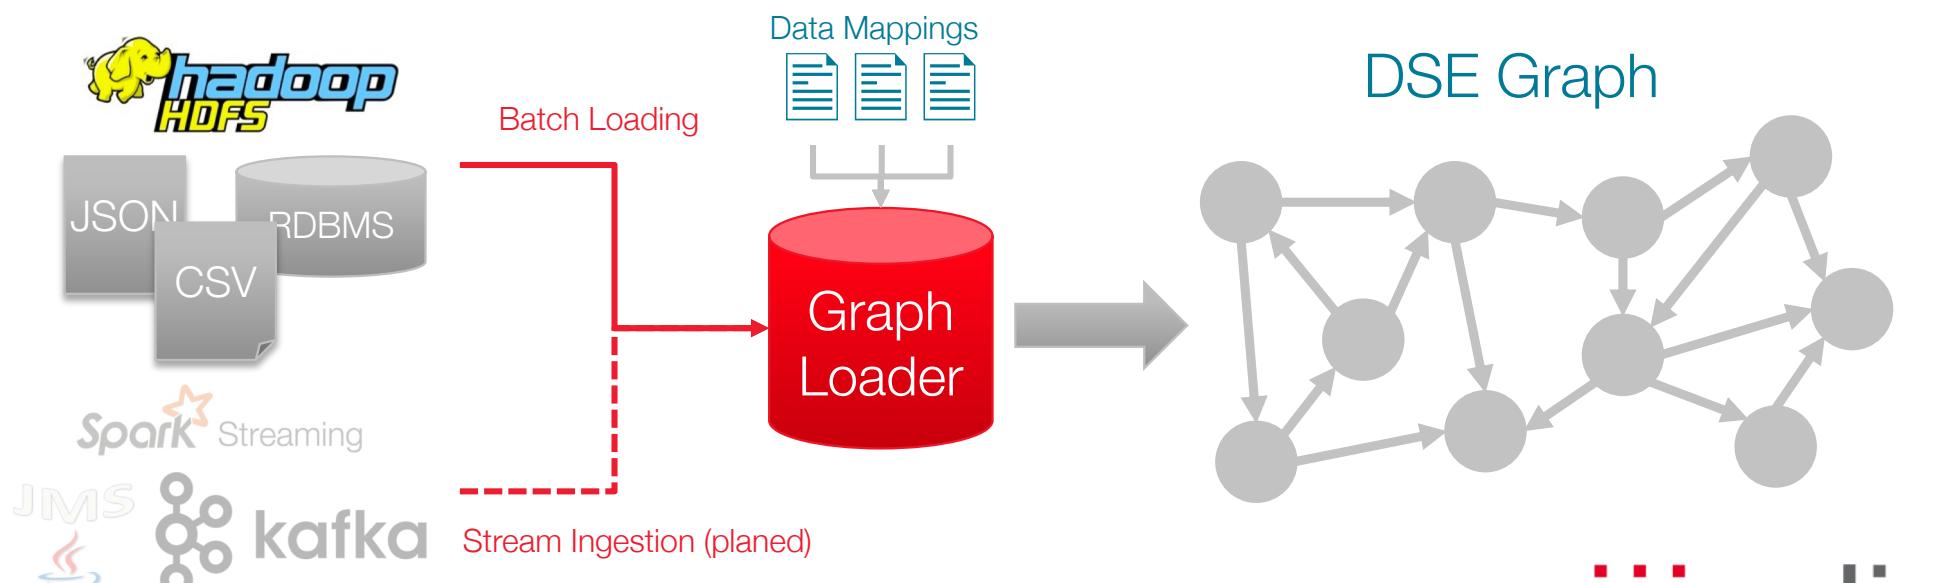
- Cell 1 (Title):** /sma/Export Tweet for Graph
- Cell 2 (Code):**

```
val tweetDF = sqlContext.sql("""  
    SELECT DISTINCT id  
        , createdAtAsLong  
        , lang  
        , LOWER(Place.fullName)  
        , CASE WHEN (coordinatesLatitude > 0.0 AND coordinatesLongitude > 0.0)  
            THEN concat(coordinatesLatitude,coordinatesLongitude)  
        ELSE null  
    END  
        , CASE WHEN (lang != 'ar' AND lang != 'ja' AND lang != 'th' AND lang != 'ko')  
            THEN TRIM(REGEXP_REPLACE(REGEXP_REPLACE(REGEXP_REPLACE(text,'\\n',''), '\\p{C}', ''))  
        ELSE null  
    END  
        , retweetedStatus.id  
        , REGEXP_REPLACE(inReplyToStatusId, '-1', '')  
        , REGEXP_REPLACE(inReplyToUserId, '-1', '')  
        , inReplyToScreenName  
        , User.id  
    FROM tweets_raw  
    """).distinct()  
tweetDF.write.format("com.databricks.spark.csv")  
    .option("header", "false").option("delimiter", DELIMITER).option("nullValue", NULL_VALUE).option("quoteMode",  
    QUOTE_MODE)  
    .save("hdfs://user/armasuisse/output/tweet.txt")  
  
tweetDF: org.apache.spark.sql.DataFrame = [tweetId: bigint, timestamp: bigint, language: string, placeId: string, geoPointId: string, text: string, retweetedId: bigint, replyToId: string, replyToUserId: string, replyToScreenName: string, publisherUserId: bigint]
```

Both cells have a status of "FINISHED". The entire job took 2 seconds to run.

Batch Ingestion with DSE Loader

- Simplifies loading of data from various sources into DSE Graph efficiently
- Inspects incoming data for schema compliance
- Uses declarative data mappings and custom transformations to handle data
- Loads all vertices first and stores in cache then edges and properties



Batch Ingestion with DSE Loader – Example

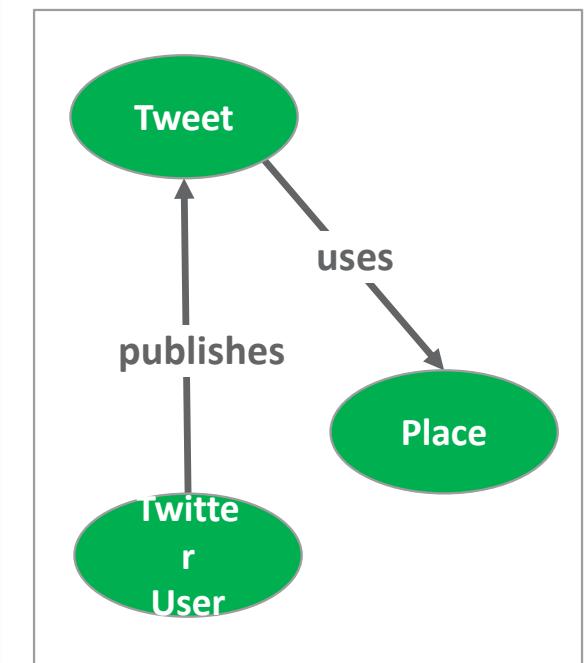
```
tweetInput = File.text("tweet.txt")  
    .delimiter(',')  
    .header('id','timestamp','language','placeId','geoPointId','text',  
        'retweetedId','replyToId','replyToUserId','replyToScreenName',  
        'publisherUserId')  
  
tweetInput = tweetInput.transform { it['replyToId'] =  
    it['replyToId'].equals('-1') ? '' : it['replyToId'];  
    it.each {key,value -> it[key] = value.replaceAll('^\"|\"$', '')};  
    it }  
  
load(tweetInput).asVertices {  
    isNew()  
    label "tweet"  
    key "id"  
    ignore "placeId"  
    ignore "publisherUserId"  
    ...  
    outV "placeId", "uses", {  
        label "place"  
        key "id"  
    }  
    ...  
    inv "publisherUserId", "publishes", {  
        label "twitterUser"  
        key "id"  
    }  
}
```

Extract

Transform

Load

```
"712181117131034624","1458632184000","en","","","",@crouchybro  
Hi, I work for @NBCNews. Are you safe? If you feel comfortable,  
would you DM me?",","","","",257242054,"crouchybro","168379717"  
...
```



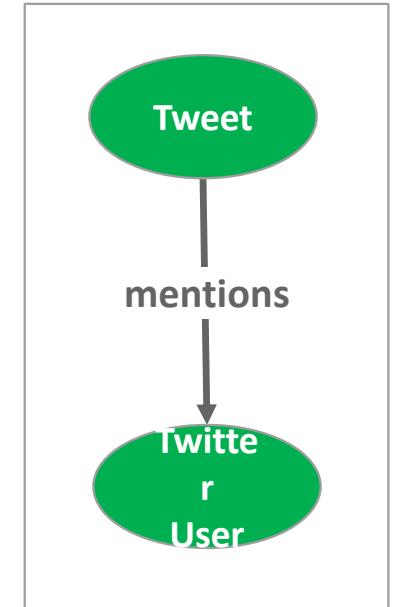
Batch Ingestion with DSE Loader - Example

```
tweetMentionsInput = File.text("tweet_mentions_user.txt")
    .delimiter(' ', '')
    .header('tweetId', 'userId', 'screenName', 'timestamp')

tweetMentionsInput = tweetMentionsInput.transform {
    it.each {key,value -> it[key] = value.replaceAll('^\"|\"$', '')}
    it
}

load(tweetMentionsInput).asEdges {
    label "mentions"
    ignore "screenName"
    outV "tweetId", {
        label "tweet"
        key "tweetId": "id"
    }
    inV "userId", {
        label "twitterUser"
        key "id"
    }
}
```

```
graphloader load-sma-graph.groovy -graph sma_graph_prod_v10 -address
192.168.69.135 -dryrun false -preparation false -abort_on_num_failures
10000 -abort_on_prep_errors false -vertex_complete false
```

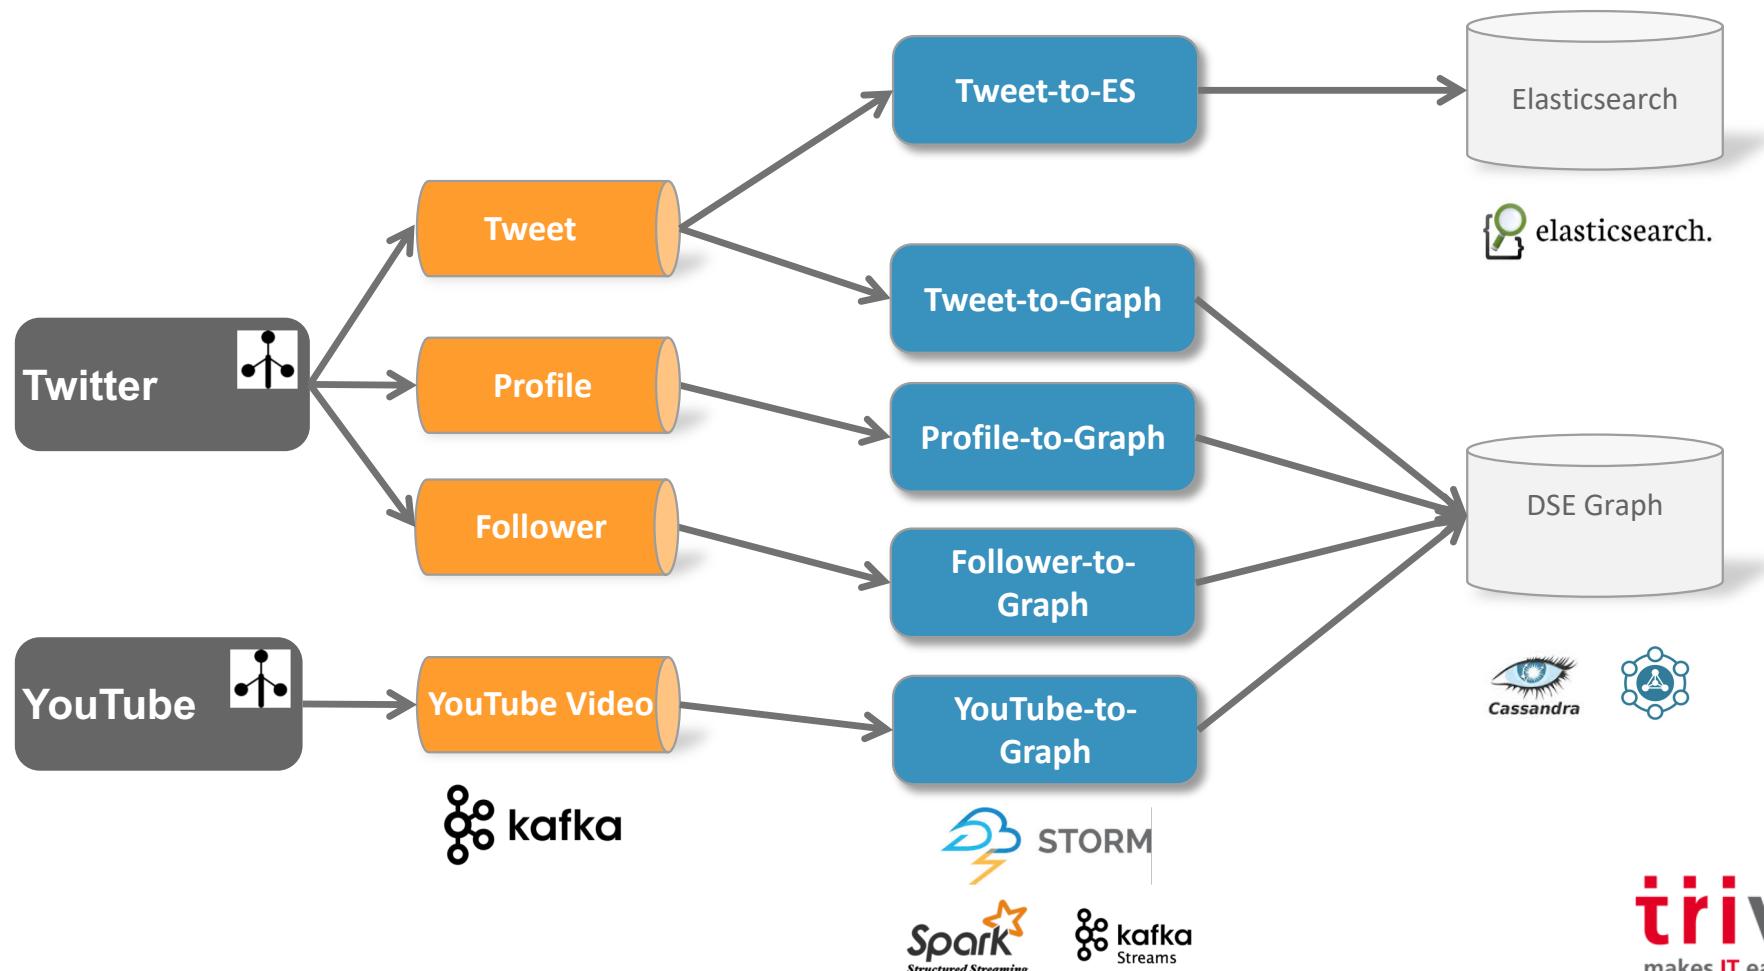


■ Performance Tests DSE Graph Loader

ID strategy	dataset	vertex threads	edge threads	number of tweets	time sec/min	tweets per sec
Gen	1-day-14	0	0	212'116	541/9.01	392
Gen	1-day-14	4	12	212'116	419/6.93	506
Gen	1-day-14	10	10	212'116	334/5.56	635
Gen	1-day-14	40	40	212'116	170/2.83	1247
Gen	1-day-14	60	80	212'116	118/1.96	1797
Cust	1-day-14	60	80	212'116	101/1.68	2100
Gen	1-mth-14	60	80	7'403'667	3888/64.8	1904
Gen	1-mth-14	60	120	7'403'667	3685/61.4	2009
Cust	1-mth-14	60	120	7'403'667	3315/55.25.	2233
Gen	1-mth-14	80	120	7'403'667	4005/66.75	1848
Cust	1-mth-14	80	120	7'403'667	3598/59.69	2057

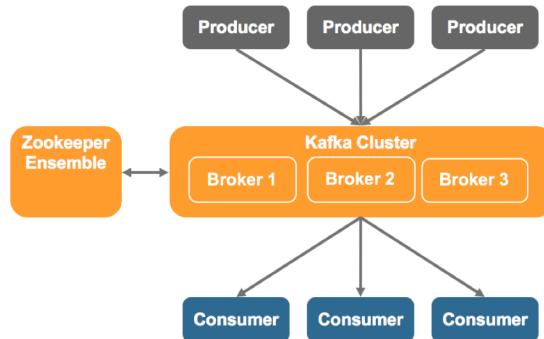
Use Case 2 – Streaming Ingestion of Social Graph

■ Streaming Ingestion of Social Network Graph

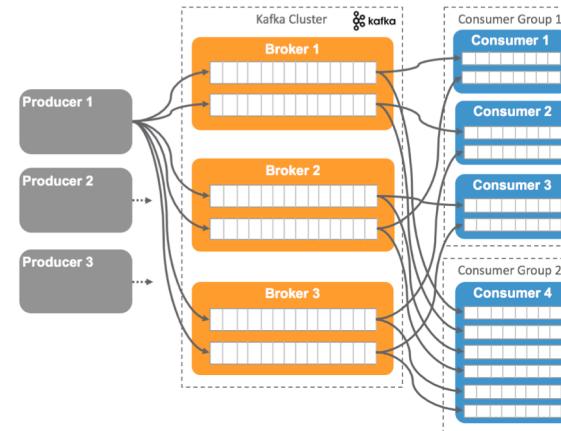


Apache Kafka – A Streaming Platform

High-Level Architecture



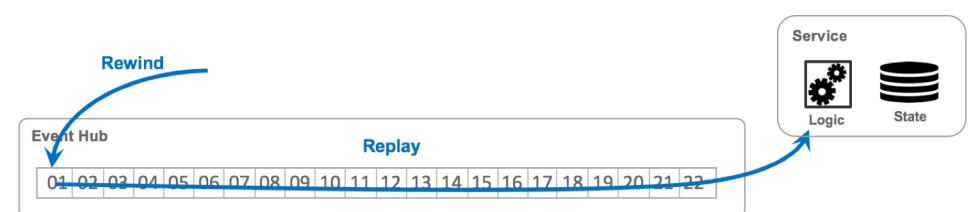
Scale-Out Architecture



Distributed Log at the Core



Logs do not (necessarily) forget



■ Graph Ingestion through DataStax Java Driver

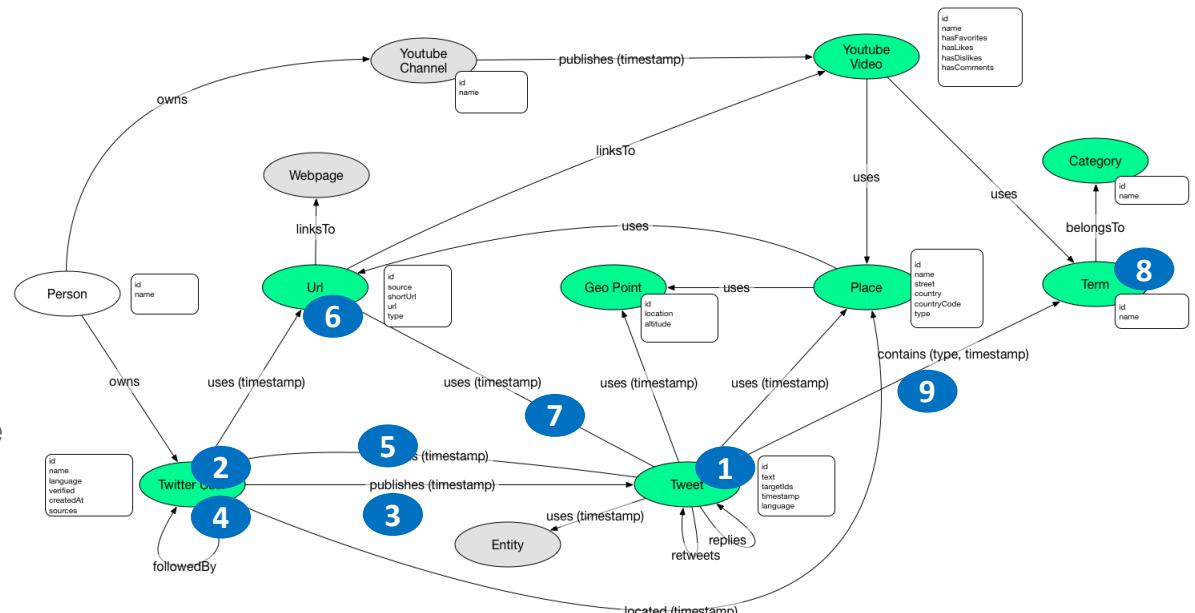
Similar to programming with JDBC: execute your gremlin statement by creating it in a String variable => also as error prone as with JDBC ;-)

fluent Java-based API available with version of DSE (5.1)

```
String stmt = ""  
    + "Vertex v" + "\n"  
    + "GraphTraversal gt = g.V().has(b.vertexLabel, b.propertyKey, b.propertyKeyValue)" + "\n"  
    + "if (!gt.hasNext()) {" + "\n"  
    + " v = graph.addVertex(label, b.vertexLabel, b.propertyKey, b.propertyKeyValue)" + "\n"  
    + "}" else {" + "\n" ...  
  
Map<String, Object> map = ImmutableMap.<String, Object>of("vertexLabel", vertexLabel, "propertyKey",  
propertyKey, "propertyKeyValue", propertyKeyValue);  
  
vertex = session.executeGraph(new SimpleGraphStatement(stmt)  
    .set("b"  
    , ImmutableMap.<String, Object>builder().putAll(map).putAll(propertyParams.getRight()).build()  
) .one() .asVertex();
```

■ Add information for new Event (tweet) to Graph – what it really means

1. Create new **Tweet** vertex and add retweets/replies edges
2. Create/Update **TwitterUser** vertex
3. Create new **publishes** edge
4. For Each Mention, Create/Update **TwitterUser** vertex
5. For each Mention, Create **mentions** edge
6. For each Link, Create/Update **Url** vertex
7. For each Link, Create **uses** edge
8. For each HashTag, Create/Update **Term** vertex
9. For each HashTag, Create **contains** edge
10.



■ Could we solve it by modeling differently?



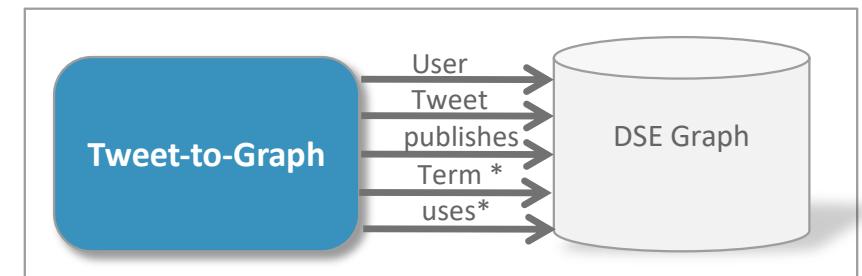
Single Insert into Social Network Graph

```
GraphTraversal gt = g.V().has(b.vertexLabel, b.propertyKey, b.propertyKeyValue)
if (!gt.hasNext()) {
    v = graph.addVertex(label, b.vertexLabel, b.propertyKey, b.propertyKeyValue)
} else {
    v = gt.next()
}
v.property("name",b.p
v.property("language"
v.property("verified"
v.property("timestamp",b.p
v.property("language",b.p
v.property("type",b.propertyParam0)
```

```
Vertex from = g.V(f).next(); Vertex to = g.V(t).next(); if
(!g.V(f).out(b.edgeLabel).V(t).hasNext()) { from.addEdge(b.edgeLabel, to) }

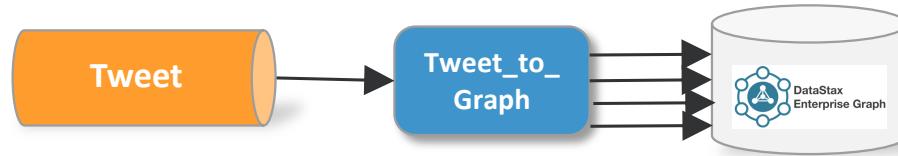
Vertex from = g.V(f).next(); Vertex to = g.V(t).next(); if
(!g.V(f).out(b.edgeLabel).V(t).hasNext()) { from.addEdge(b.edgeLabel, to) }

Vertex from = g.V(f).next(); Vertex to = g.V(t).next(); if
(!g.V(f).out(b.edgeLabel).V(t).hasNext()) { from.addEdge(b.edgeLabel, to) }
```

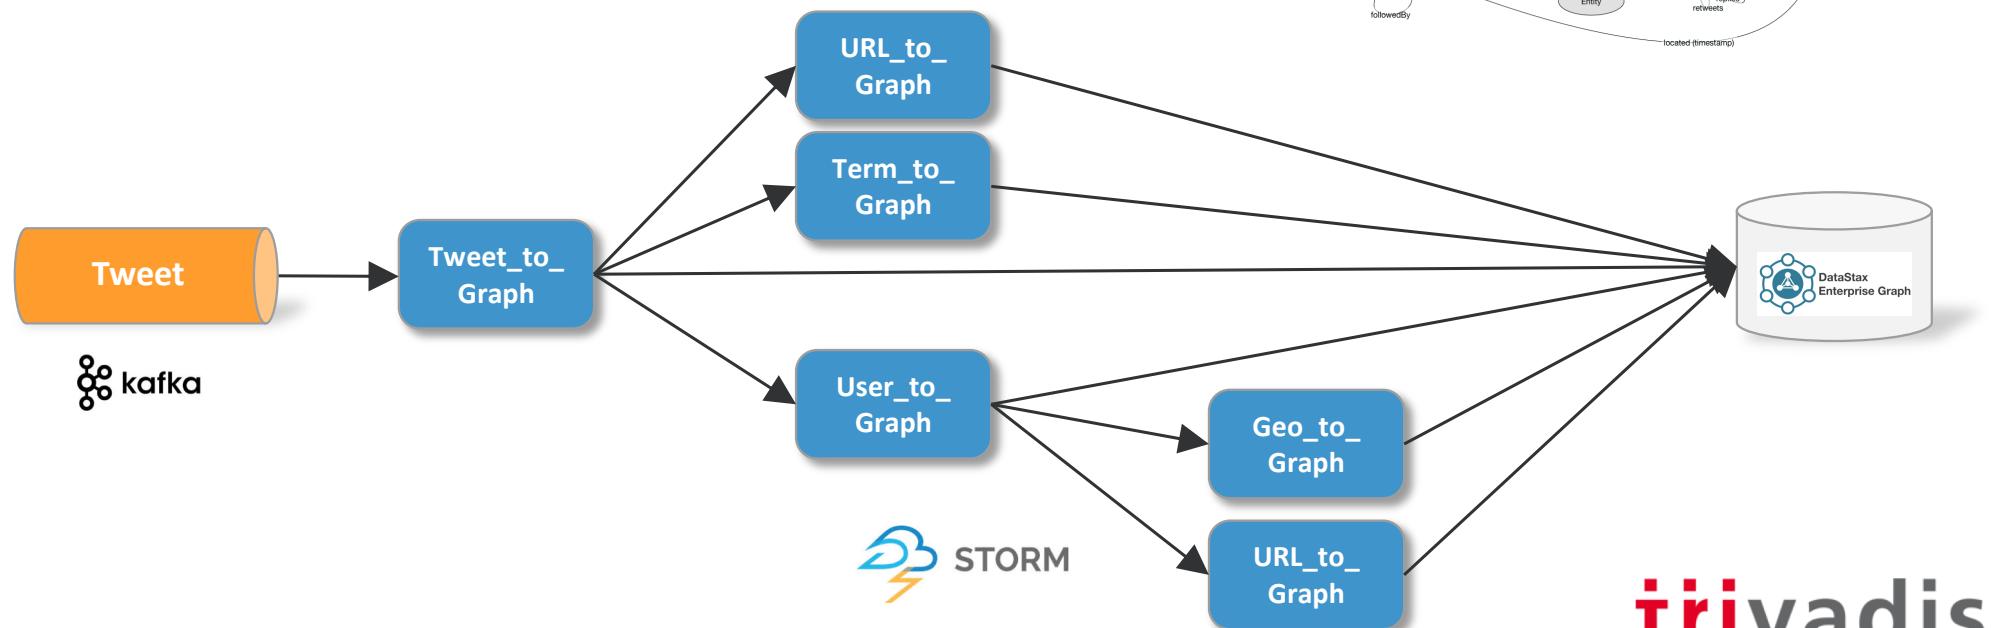


■ 2nd approach - "Fan out ingestion driven by model"

1st approach => too slow

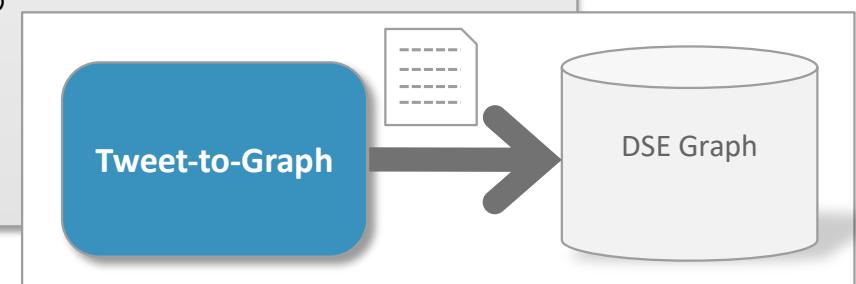


2nd approach => too difficult to implement and maintain



3rd approach - Batch Insert into Social Network Graph

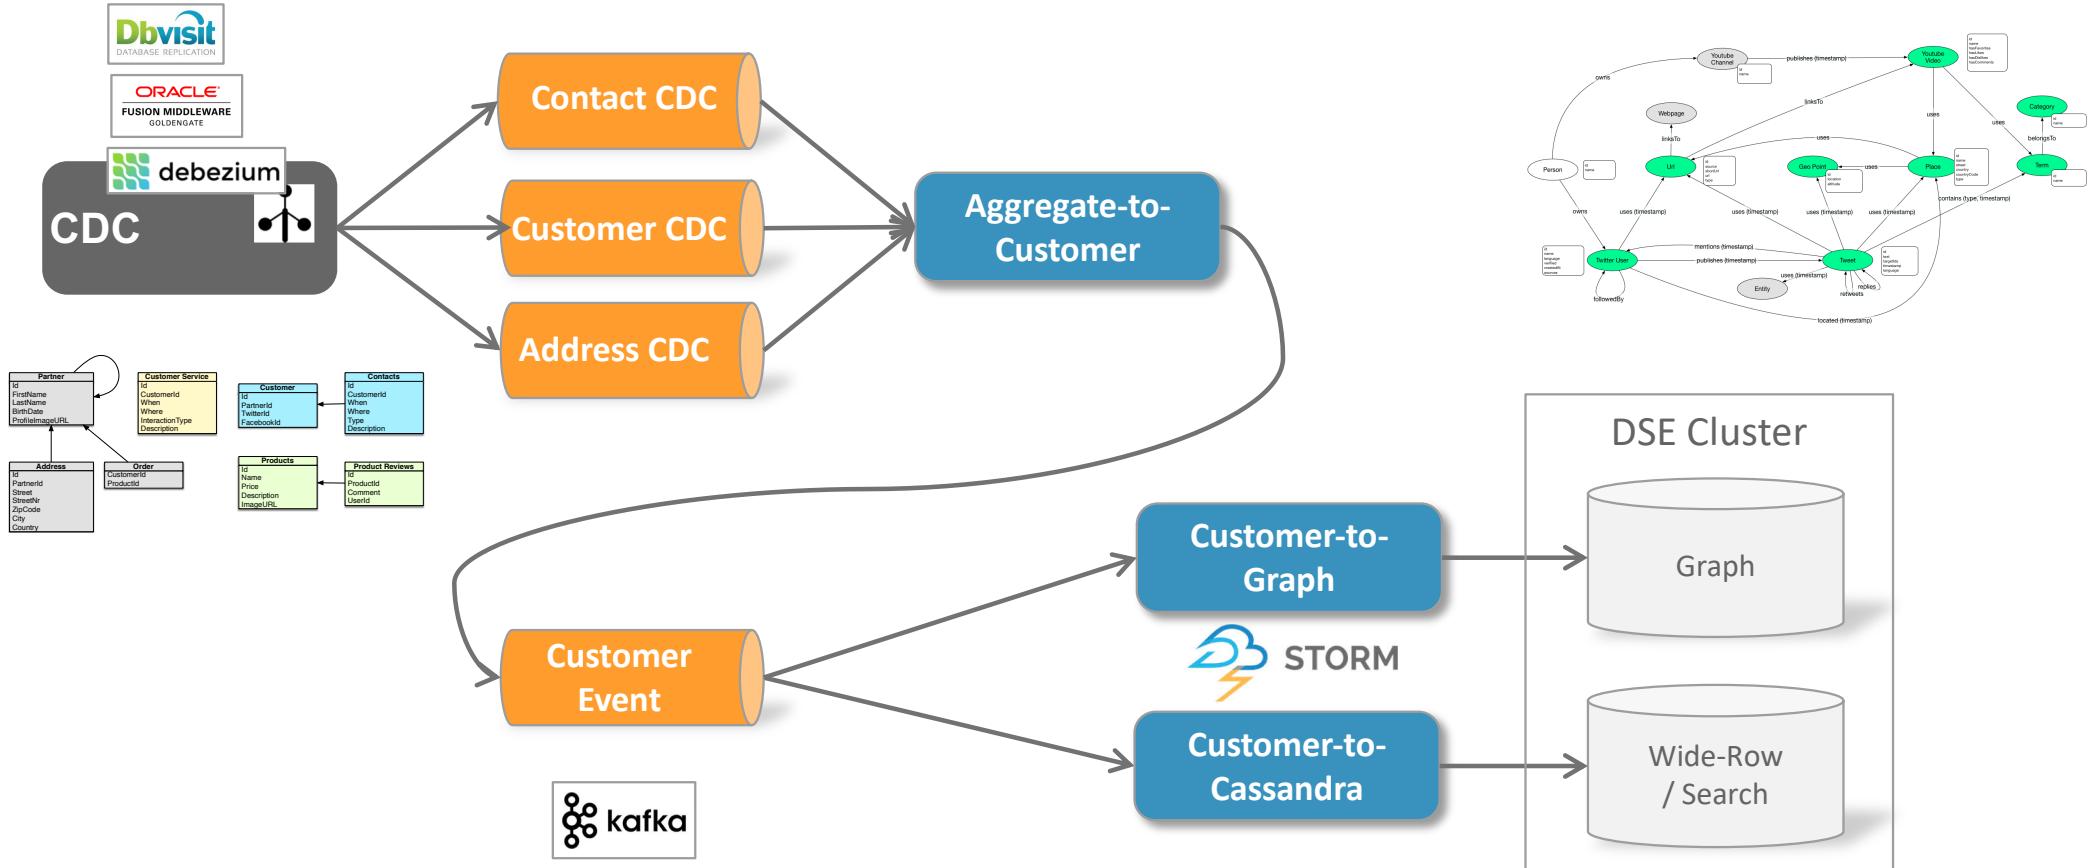
```
user = g.V().has('twitterUser', 'id', bUser.id).tryNext().orElseGet { g.addV('twitterUser').property('id', bUser.id).next() }
user.property('name', bUser.name)
user.property('language', bUser.language)
user.property('verified', bUser.verified)
tweet = g.V().has('tweet', 'id', bTweet.id).tryNext().orElseGet { g.addV('tweet').property('id', bTweet.id).next() }
tweet.property('timestamp', bTweet.timestamp)
tweet.property('language', bTweet.language)
if (!g.V(user).out('publishes').hasId(tweet.id()).hasNext()) {
    publishes = g.V(user).as('f').V(tweet).as('t').addE('publishes').from('f').next()
}
i = 0
Vertex[] term = new Vertex[10]
for (Object keyValue : bTerm.propertyKeyValues) {
    term[i] = g.V().has('term', 'name', keyValue).tryNext().orElseGet { g.addV('term').property('name', keyValue).next() }
    Map<String, Object> params = bTerm.params[i]
    if (params != null)
        for (String key : params.keySet()) {
            term[i].property(key, params.get(key))
        }
    i++
}
Edge[] usesTerm = new Edge[10]
for (i = 0; i < bUsesTerm.count; i++) {
    if (!g.V(tweet).out('uses').hasId(term[i].id()).hasNext()) {
        usesTerm[i] = g.V(tweet).as('f').V(term[i]).as('t').addE('uses').from('f').next()
    }
    Map<String, Object> params = bUsesTerm.params[i]
    if (params != null)
        for (String key : params.keySet()) {
            usesTerm[i].property(key, params.get(key))
        }
}
return nOf
```



■ Benchmark Single vs. Scripted Insert

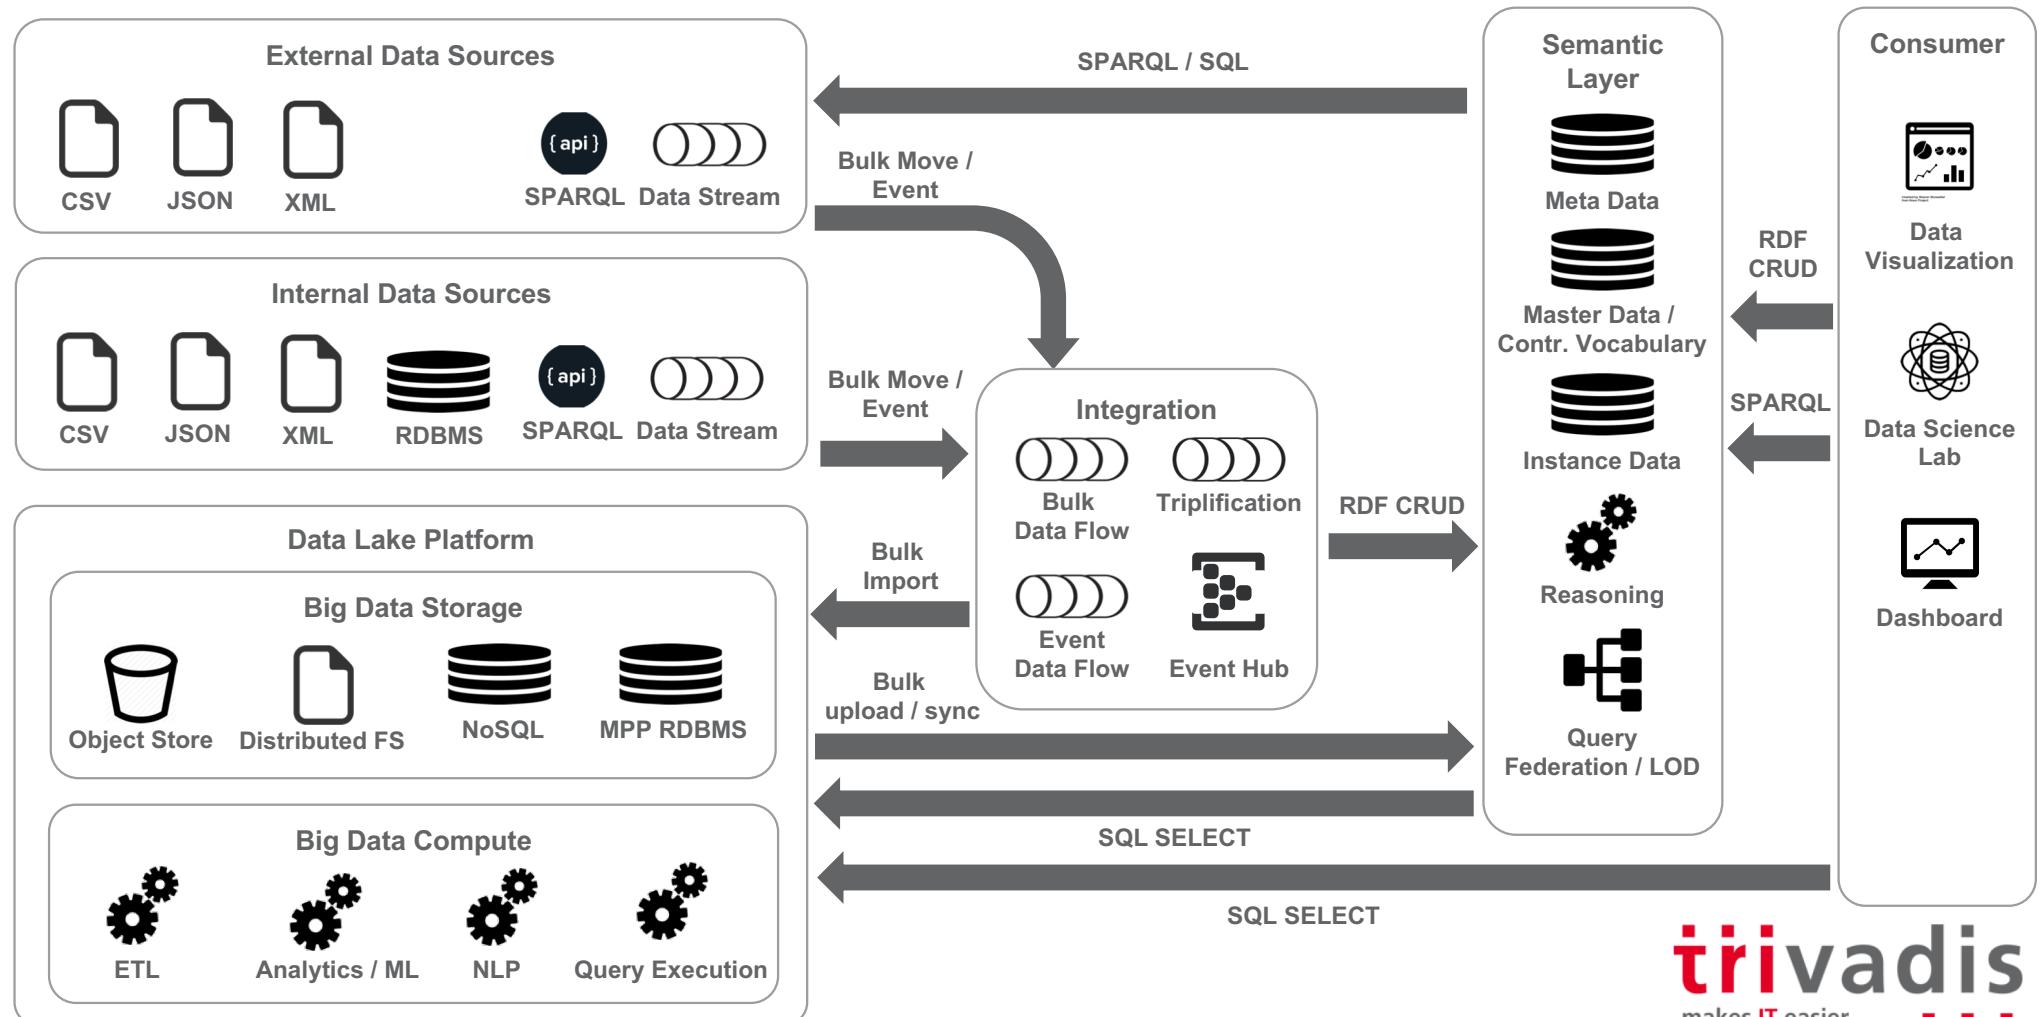
Test Case	Single Inserts	Scripted (Batch) Insert
10 tweets with 10 hashtags each	569ms - 690ms	132ms – 172ms
10 tweets with 30 hashtags each	1311ms – 1396ms	285ms – 364ms
20 tweets with 30 hashtags each	2734ms – 3385ms	640ms – 767ms
20 tweets with 40 hashtags each	3576ms – 4089ms	778ms – 992ms
100 tweets with 10 hashtags each	6340ms – 7171ms	1673ms - 2428

Process CDC (Change Data Capture) Events



Use Case 3: Streaming Ingestion into RDF Graph

Semantic Layer 3 – streaming update of Triple store



“Triplification” – 1st option

Using RML (<http://rml.io/>) and Carml (<https://github.com/carml/carml>)

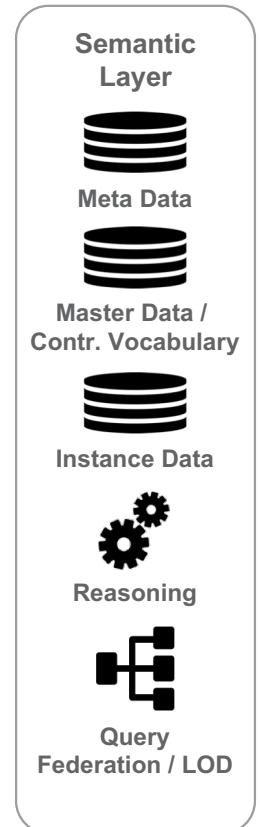
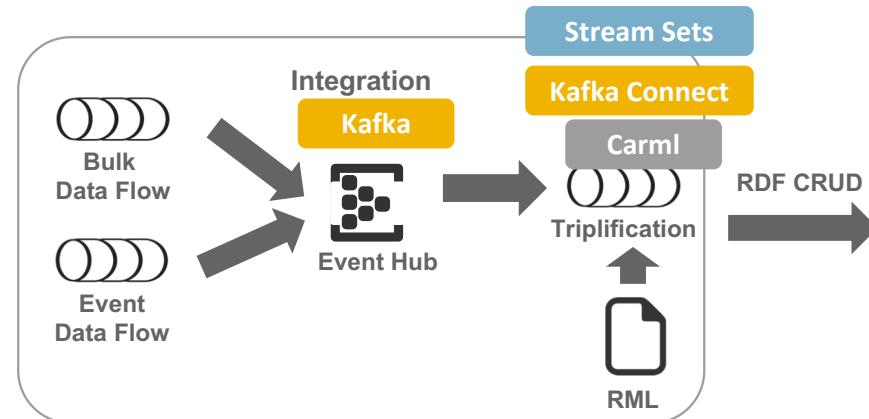
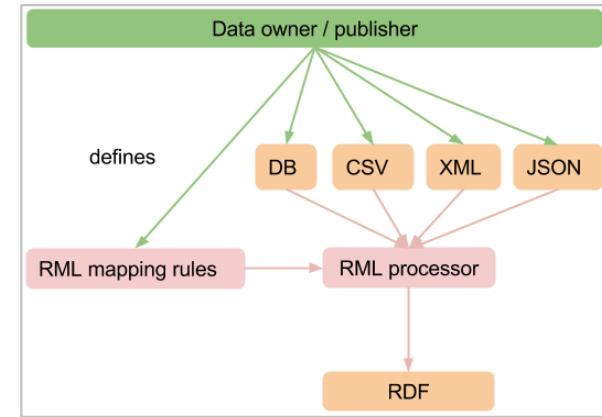
```
@prefix rml: <http://semweb.mmlab.be/ns/rml#>.
@prefix ex: <http://example.com/>.
@prefix ql: <http://semweb.mmlab.be/ns/ql#> .
@prefix carml: <http://carml.taxonic.com/carml/> .

<#SubjectMapping> a rr:TriplesMap;
  rml:logicalSource [
    rml:source [
      a carml:Stream;
      carml:streamName "input";
    ];
    rml:referenceFormulation ql:JSONPath;
    rml:iterator "$"
  ];

  rr:subjectMap [
    rr:template "http://example.com/person/{firstName}/{lastName}";
    rr:class ex:Person
  ];

  rr:predicateObjectMap [
    rr:predicate ex:likes;
    rr:objectMap [
      rml:reference "hobbies";
    ]
  ];

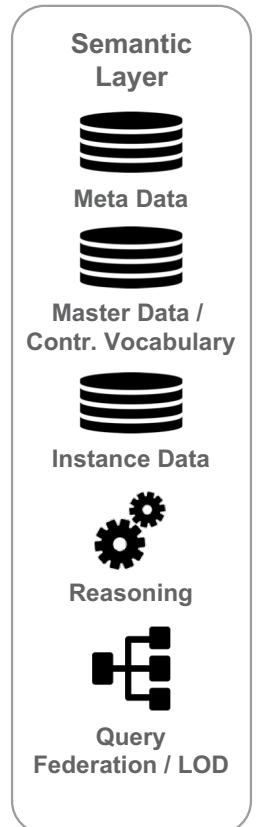
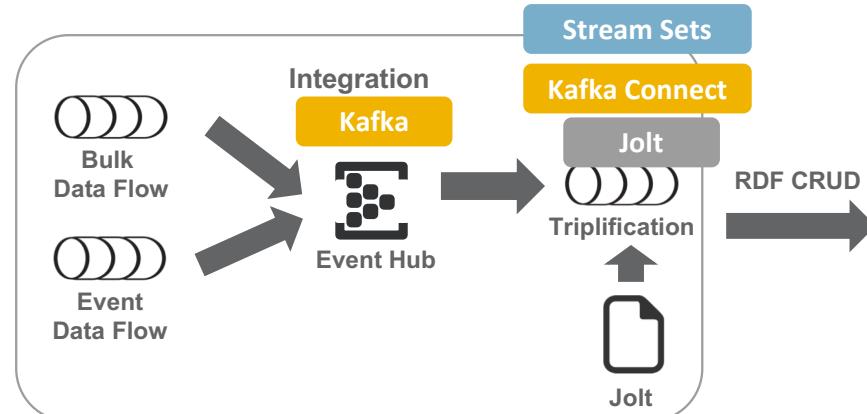
  rr:predicateObjectMap [
    rr:predicate ex:hasFirstName;
    rr:objectMap [
      rml:reference "firstName";
    ]
  ].
```



■ “Triplification” – 2nd option

Using Jolt Transformation (<https://github.com/bazaarvoice/jolt>) to create JSON-LD

Demo: <http://jolt-demo.appspot.com/#inputArrayToPrefix>



Summary - Lessons Learnt

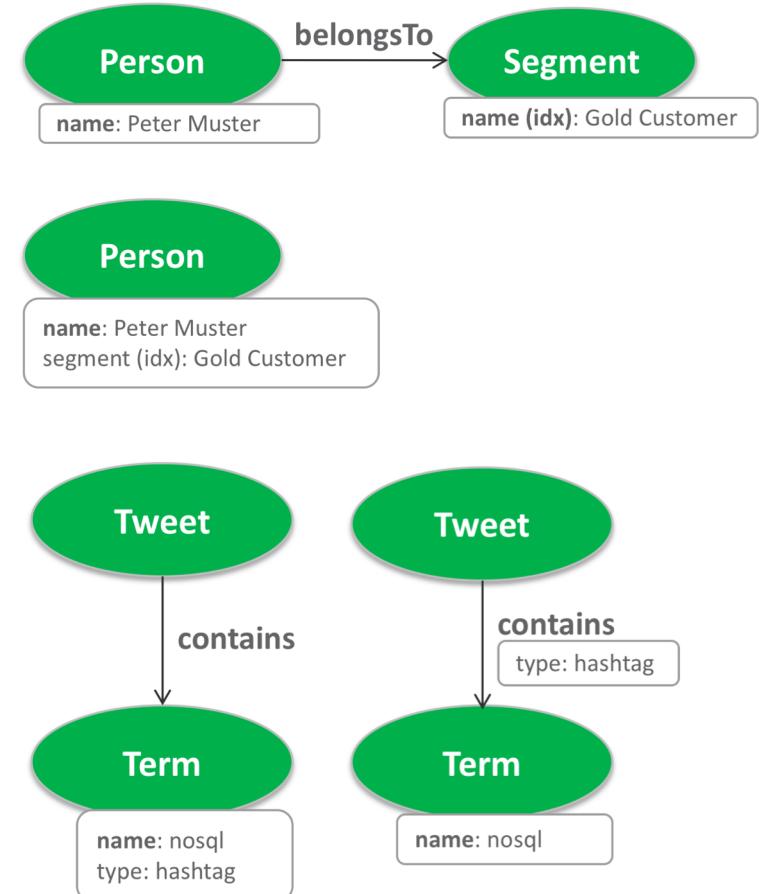
■ Summary - Lessons Learnt

Graph Schema

- Edges can not cross graph schemas
- Vertex or edge? Vertex or property?
- Custom Vertex IDs can be useful
 - idempotent
 - faster on insert (no id generation needed)
 - can lead to hotspots

Data Ingestion

- Single Nodes & Edge Insert is slow => batching by Groovy Script to speed up
- DSE graph loader works good for batch data loading



**Technology on its own won't help you.
You need to know how to use it properly.**



trivadis
makes **IT** easier. 